

# ***Program Formats***

- Programs obey specific file formats
  - CP/M and DOS: COM executables (\*.com)
  - DOS: MZ executables (\*.exe)
    - Named after Mark Zbikowski, a DOS developer
  - Windows Portable Executable (PE, PE32+) (\*.exe)
    - Modified version of Unix COFF executable format
    - PE files start with an MZ header. Why?
  - Unix/Linux: Executable and Linkable Format (ELF)
  - Mac OSX: Mach object file format (Mach-O)

# *test.c*

```
#include <stdio.h>
```

```
int big_big_array[10 * 1024 * 1024];
```

```
char *a_string = "Hello, World!";
```

```
int a_var_with_value = 100;
```

```
int main(void) {
```

```
    big_big_array[0] = 100;
```

```
    printf("%s\n", a_string);
```

```
    a_var_with_value += 20;
```

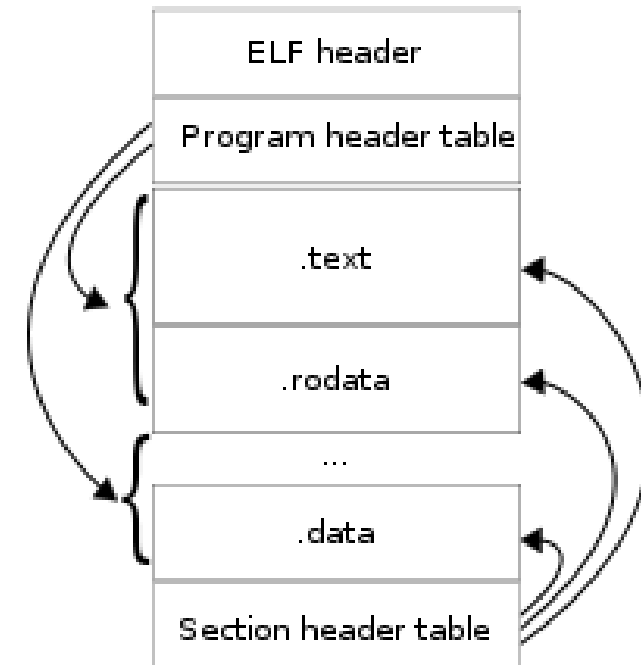
```
    printf("main is : %p\n", &main);
```

```
    return 0;
```

```
}
```

# *ELF File Format*

- ELF Header
  - Contains compatibility info
  - Entry point of the executable code
- Program header table
  - Lists all the segments in the file
  - Used to load and execute the program
- Section header table
  - Used by the linker



# ELF Header Format

```
typedef struct {  
1   unsigned char e_ident[EI_NIDENT];  
    Elf32_Half e_type;  
5   Elf32_Half e_machine;  
    Elf32_Word e_version;  
    Elf32_Addr e_entry;  
    Elf32_Off e_phoff;  
    Elf32_Off e_shoff;  
10  Elf32_Word e_flags;  
    Elf32_Half e_ehsize;  
    Elf32_Half e_phentsize;  
    Elf32_Half e_phnum;  
    Elf32_Half e_shentsize;  
15  Elf32_Half e_shnum;  
    Elf32_Half e_shstrndx;  
} Elf32_Ehdr;
```

```
$ gcc -g -o test test.c
```

```
$ readelf --header test
```

ELF Header:

Magic:	7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00
Class:	<u>ELF64</u>
Data:	2's complement, little endian
Version:	1 (current)
OS/ABI:	UNIX - System V
ABI Version:	0
Type:	EXEC (Executable file)
Machine:	<u>Advanced Micro Devices X86-64</u>
Version:	0x1
Entry point address:	<u>0x400460</u>
Start of program headers:	<u>64 (bytes into file)</u>
Start of section headers:	5216 (bytes into file)
Flags:	0x0
Size of this header:	<u>64 (bytes)</u>
Size of program headers:	56 (bytes)
Number of program headers:	<u>9</u>
Size of section headers:	64 (bytes)
Number of section headers:	<u>36</u>
Section header string table index:	33

# Investigating the Entry Point

```
int main(void) {  
    printf("main is : %p\n", &main);  
    return 0;  
}
```

- Most compilers insert extra code into compiled programs
- This code typically runs before and after main()

```
(base) kolin@mosaic:~/col7001$ objdump --disassemble -M intel hello  
hello:      file format elf64-x86-64  
  
Disassembly of section .init:  
0000000000001000 <.init>:  
1000:    f3 0f 1e fa                endbr64  
1004:    48 83 ec 08                sub     rsp,0x8  
1008:    48 8b 05 d9 2f 00 00      mov     rax,QWORD PTR [rip+0x2fd9]    # 3fe8 <__gmon_start__@Base>  
100f:    48 85 c0                  test    rax,rax  
1012:    74 02                      je      1016 <_init+0x16>  
1014:    ff d0                      call    rax  
1016:    48 83 c4 08                add     rsp,0x8  
101a:    c3                        ret  
  
Disassembly of section .plt:  
0000000000001020 <.plt>:  
1020:    ff 35 9a 2f 00 00      push    QWORD PTR [rip+0x2f9a]    # 3fc0 <_GLOBAL_OFFSET_TABLE_+0x8>  
1026:    f2 ff 25 9b 2f 00 00    bnd jmp QWORD PTR [rip+0x2f9b]    # 3fc8 <_GLOBAL_OFFSET_TABLE_+0x10>  
102d:    0f 1f 00                nop     DWORD PTR [rax]  
1030:    f3 0f 1e fa                endbr64  
1034:    68 00 00 00 00          push    0x0  
1039:    f2 e9 e1 ff ff ff      bnd jmp 1020 <_init+0x20>  
103f:    90                        nop  
  
Disassembly of section .plt.got:  
t>:    endbr64  
0      bnd jmp QWORD PTR [rip+0x2fad]    # 3ff8 <__cxa_finalize@GLIBC_2.2.5>  
      nop     DWORD PTR [rax+rax*1+0x0]  
  
0      endbr64  
      bnd jmp QWORD PTR [rip+0x2f75]    # 3fd0 <printf@GLIBC_2.2.5>  
      nop     DWORD PTR [rax+rax*1+0x0]  
  
1004:    31 ed                endbr64  
      xor     ebp,ebp  
      mov     r9,rdx
```

```
(base) kolin@mosaic:~/col7001$ readelf --headers hello |grep Entry  
Entry point address:          0x1060  
(base) kolin@mosaic:~/col7001$ ./hello  
main is : 0x555555555149  
(base) kolin@mosaic:~/col7001$
```

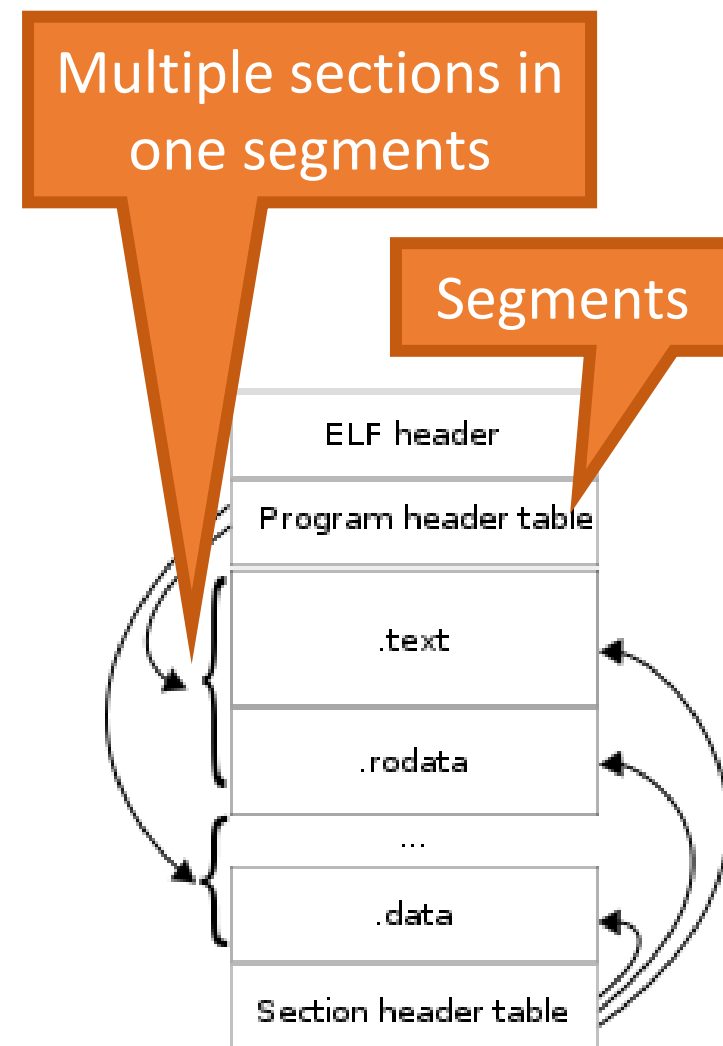
```
(base) kolin@mosaic:~/col7001$ readelf --headers hello |grep Entry  
Entry point address:          0x1060  
(base) kolin@mosaic:~/col7001$
```

```
(base) kolin@mosaic:~/col7001$ objdump --disassemble -M intel hello |less
```

```
0000000000001060 <_start>:  
1060:    f3 0f 1e fa                endbr64  
1064:    31 ed                xor     ebp,ebp  
1066:    48 8b 05 d9 2f 00 00      mov     rax,QWORD PTR [rip+0x2fd9]    # 3fe8 <__gmon_start__@Base>  
106d:    48 85 c0                test    rax,rax  
1072:    74 02                      je      1076 <_start+0x16>  
1074:    ff d0                call    rax  
1076:    48 83 c4 08                add     rsp,0x8  
107a:    c3                        ret
```

# Sections and Segments

- Sections are the various pieces of code and data that get linked together by the compiler
- Each segment contains one or more sections
  - Each segment contains sections that are related
    - E.g. all code sections
- Segments are the basic units for the loader
- Key sections:
  - .text – Executable code
  - .bss – Global variables initialized to zero
  - .data, .rodata – Initialized data and strings
  - .strtab – Names of functions and variables
  - .symtab – Debug symbols



# Section Example

String variable → .data

```
int big_big_array[10*1024];  
char *a_string = "Hello, World!";  
int a_var_with_value = 0x100;
```

Empty 10 MB  
array → .bss

```
int main(void) {  
    big_big_array[0] = 100;  
    printf("%s\n", a_string);  
    a_var_with_value += 20;  
    ...  
}
```

Initialized global  
variable → .data

String constant → .rodata

Code → .text

```
$ readelf --headers ./test
```

...

Section to Segment mapping:

Segment Sections...

00

01 .interp

02 .interp .note.ABI-tag .note.gnu.build-id .gnu.hash .dynsym .dynstr .gnu.version .gnu.version\_r

.rela.dyn .rela.plt .init .plt .text fin .rodata eh\_frame\_hdr .eh\_frame

03 .ctors .dtors .jcr .dynamic .got .got.plt .data .bss

04 .dynamic

05 .note.ABI-tag .note.gnu.build-id

06 .eh\_frame\_hdr

07

08 .ctors .dtors .jcr .dynamic .got

...

There are 36 section headers, starting at offset 0x1460:

Section Headers:

[Nr]	Name	Type	Address	Offset	Size	ES	Flags	Link	Info	Align
[ 0]		NULL	00000000	00000000	00000000	00		0	0	0
[ 1]	.interp	PROGBITS	00400238	00000238	0000001c	00	A	0	0	1
[ 2]	.note.ABI-tag	NOTE	00400254	00000254	00000020	00	A	0	0	4
[ 3]	.note.gnu.build-id	NOTE	00400274	00000274	00000024	00	A	0	0	4
[ 4]	.gnu.hash	GNU_HASH	00400298	00000298	0000001c	00	A	5	0	8
[ 5]	.dynsym	DYNSYM	004002b8	000002b8	00000078	18	A	6	1	8
[ 6]	.dynstr	STRTAB	00400330	00000330	00000044	00	A	0	0	1
[ 7]	.gnu.version	VERSYM	00400374	00000374	0000000a	02	A	5	0	2

...



# *.text Example Header*

```
typedef struct {  
    Elf32_Word p_type;  
    Elf32_Off  p_offset;  
5   Elf32_Addr p_vaddr;  
    Elf32_Addr p_paddr;  
    Elf32_Word p_filesz;  
    Elf32_Word p_memsz;  
    Elf32_Word p_flags;  
10   Elf32_Word p_align;
```

Address to load  
section in memory

Offset of data in the file

How many bytes (in hex)  
are in the section

Data for the  
program

Executable

```
... ons ./test  
...  
Section Headers:  
...  
[Nr] Name      Type      Address    Offset     Size       ES  Flags  Link  Info  Align  
[13] .text     PROGBITS  00400460   00000460   00000218   00  AX    0     0     16  
...
```

# *.bss Example Header*

```
int big_big_array[10*1024*1024];
```

```
typedef struct {
```

```
    Elf32_Word p_type;
```

```
    Elf32_Off  p_offset;
```

```
    5 Elf32_Addr p_vaddr;
```

```
    Elf32_Addr p_paddr;
```

```
    Elf32_Word p_filesz;
```

```
    Elf32_Word p_memsz;
```

```
    Elf32_Word p_flags;
```

```
    10 Elf32_Word p_align;
```

Offset of data in the file  
(Notice the length = 0)

Address to load  
section in memory

Contains  
no data

hex(4\*10\*1024\*1024) =  
0x2800020

Writable

```
$ readelf -S sections ./
```

```
...
```

```
Section Headers:
```

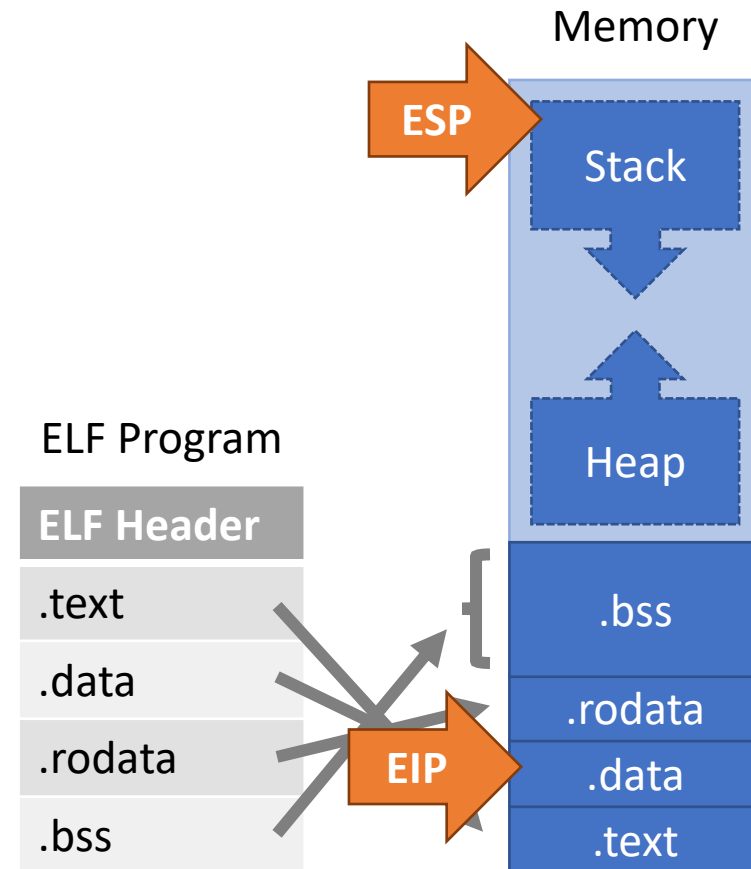
```
...
```

[Nr]	Name	Type	Address	Offset	Size	ES	Flags	Link	Info	Align
[25]	.bss	NOBITS	00601040	00001034	02800020	00	WA	0	0	32
[26]	.comment	PROGBITS	00000000	00001034	000002a	01	MS	0	0	1

```
...
```

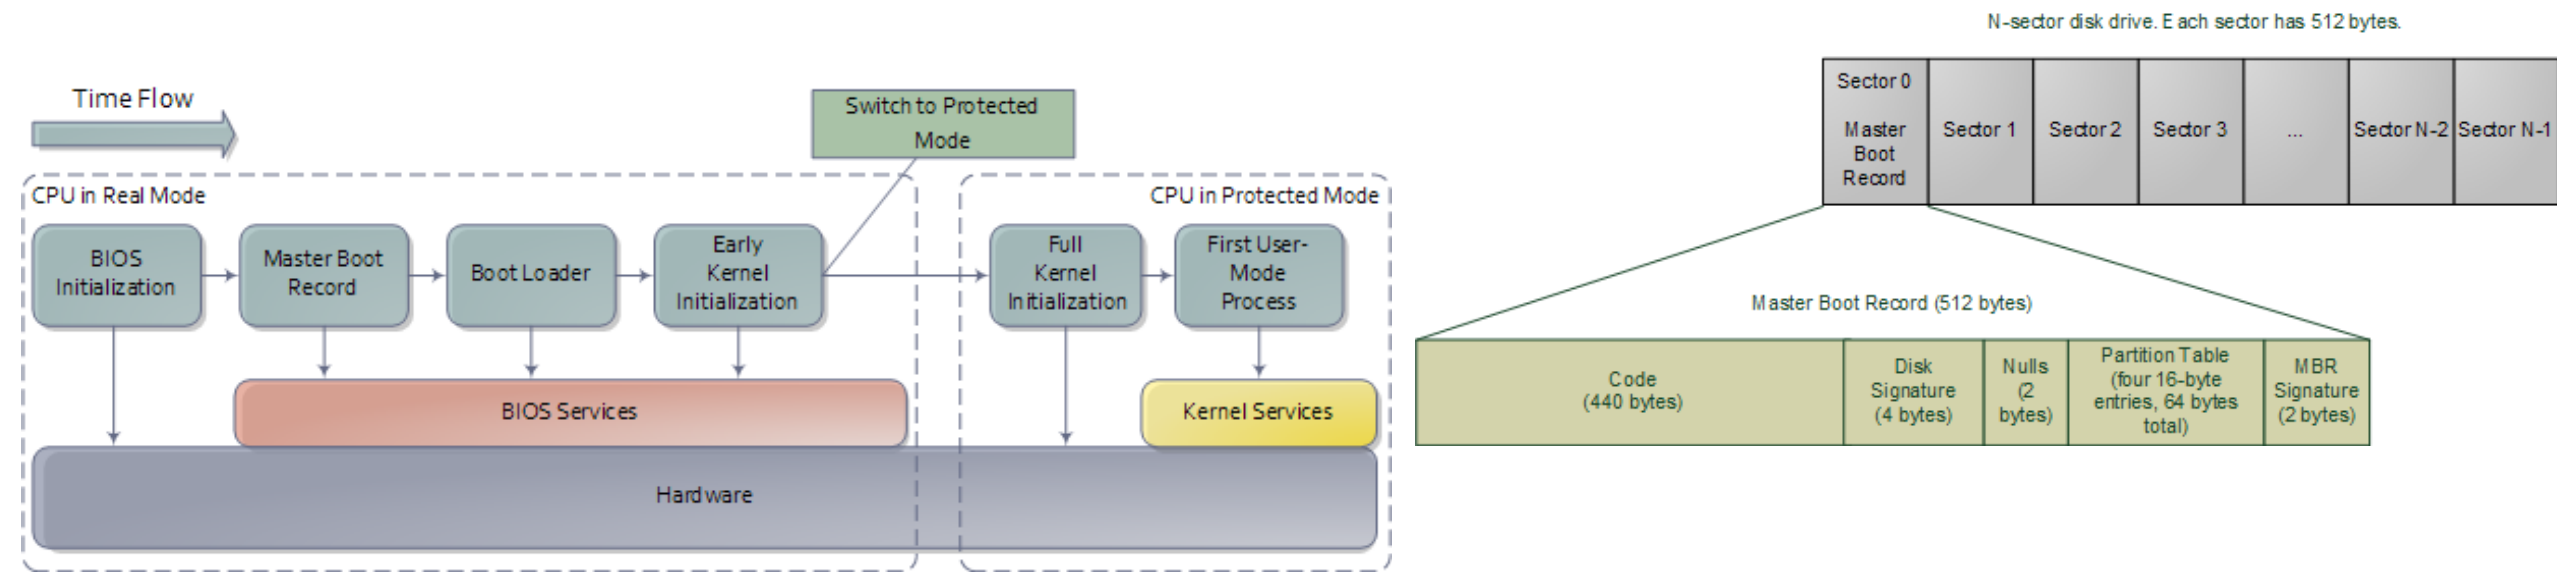
# *The Program Loader*

- OS functionality that loads programs into memory, creates processes
  - Places segments into memory
    - Expands segments like .bss
  - Loads necessary dynamic libraries
  - Performs relocation
  - Allocated the initial stack frame
  - Sets EIP to the programs entry point



# Boot Sequence

- The BIOS
  - Basic I/O System
  - Initial interface between the hardware and the operating system
  - Responsible for allowing you to control your computer's hardware settings for booting up
- Bootable devices
  - MBR
  - Loaded at location 0x7c00 in RAM and control is given to this code





**POWER**



**FIRMWARE**

BIOS/UEFI  
0xFFFFFFFF0

- Initialize CPU, memory
- Power-On Self-Test (POST)
- Initialize drivers



**BOOTLOADER**

MBR/UEFI

- Select OS/kernel
- Load kernel
- Set up memory, pass parameters



**KERNEL**

Linux / Windows

- Initialize CPU, memory, devices
- Mount root filesystem
- Initialize scheduler



**INIT PROCESS**

init / systemd

- Launch services, daemons
- Initialize network, logging

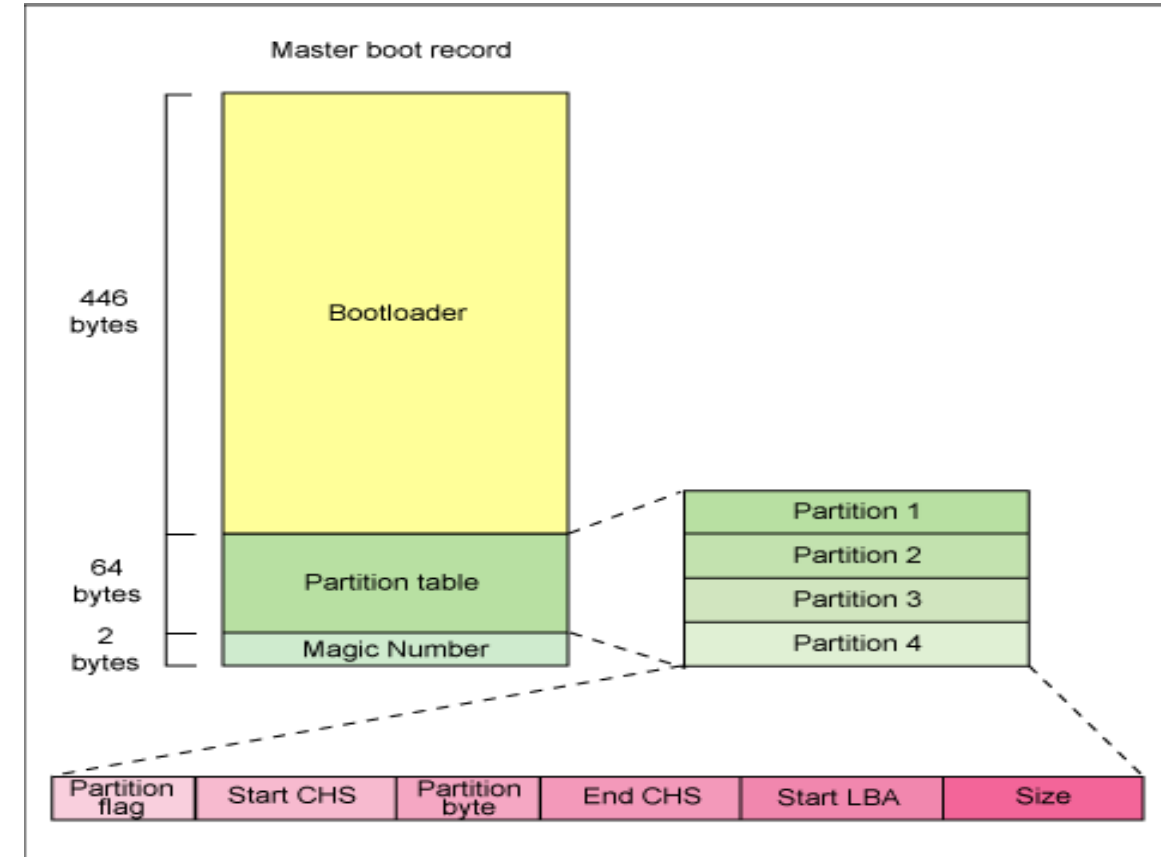


**USER ENVIRONMENT**

Shell / Desktop

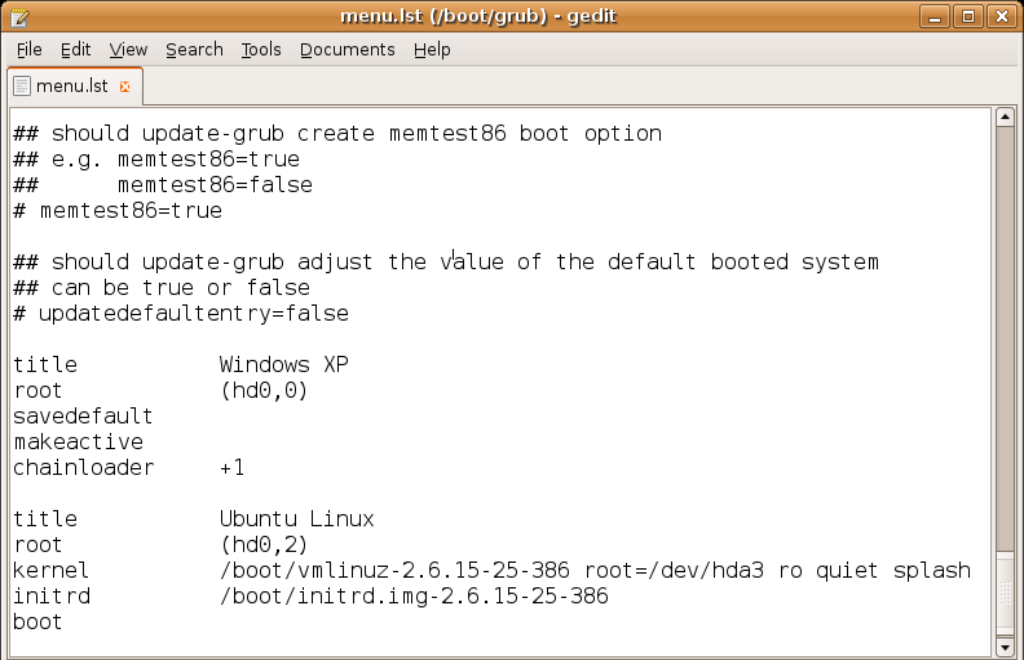
# Boot Loaders - Linux

- A multi-stage program which eventually loads the kernel image and initial RAM Disk(initrd)
  - Stage-1 Boot Loader is less than 512 bytes (why?)
    - Just does enough to load next stage
  - Next stage can reside in boot sector or the partition or area in the disk which is hardcoded in MBR
  - Stage-1.5 is a crucial feature
    - Makes grub file-system aware



# GRUB Stage-2

- Since file-system aware, it can display the boot options to user from
  - */boot/grub/grub.cfg*
- GRUB command-line
  - *grub> kernel /bzImage-2.6.14.2*
    - *[Linux-bzImage,setup=0x1400, size=0x29672e]*
  - *grub> initrd /initrd-2.6.14.2.img*
    - *[Linux-initrd @ 0x5f13000, 0xcc199 bytes]*
- So now kernel time
  - But where are we going to place kernel image in the memory??



```
menu.lst (/boot/grub) - gedit
File Edit View Search Tools Documents Help
menu.lst
## should update-grub create memtest86 boot option
## e.g. memtest86=true
##      memtest86=false
# memtest86=true

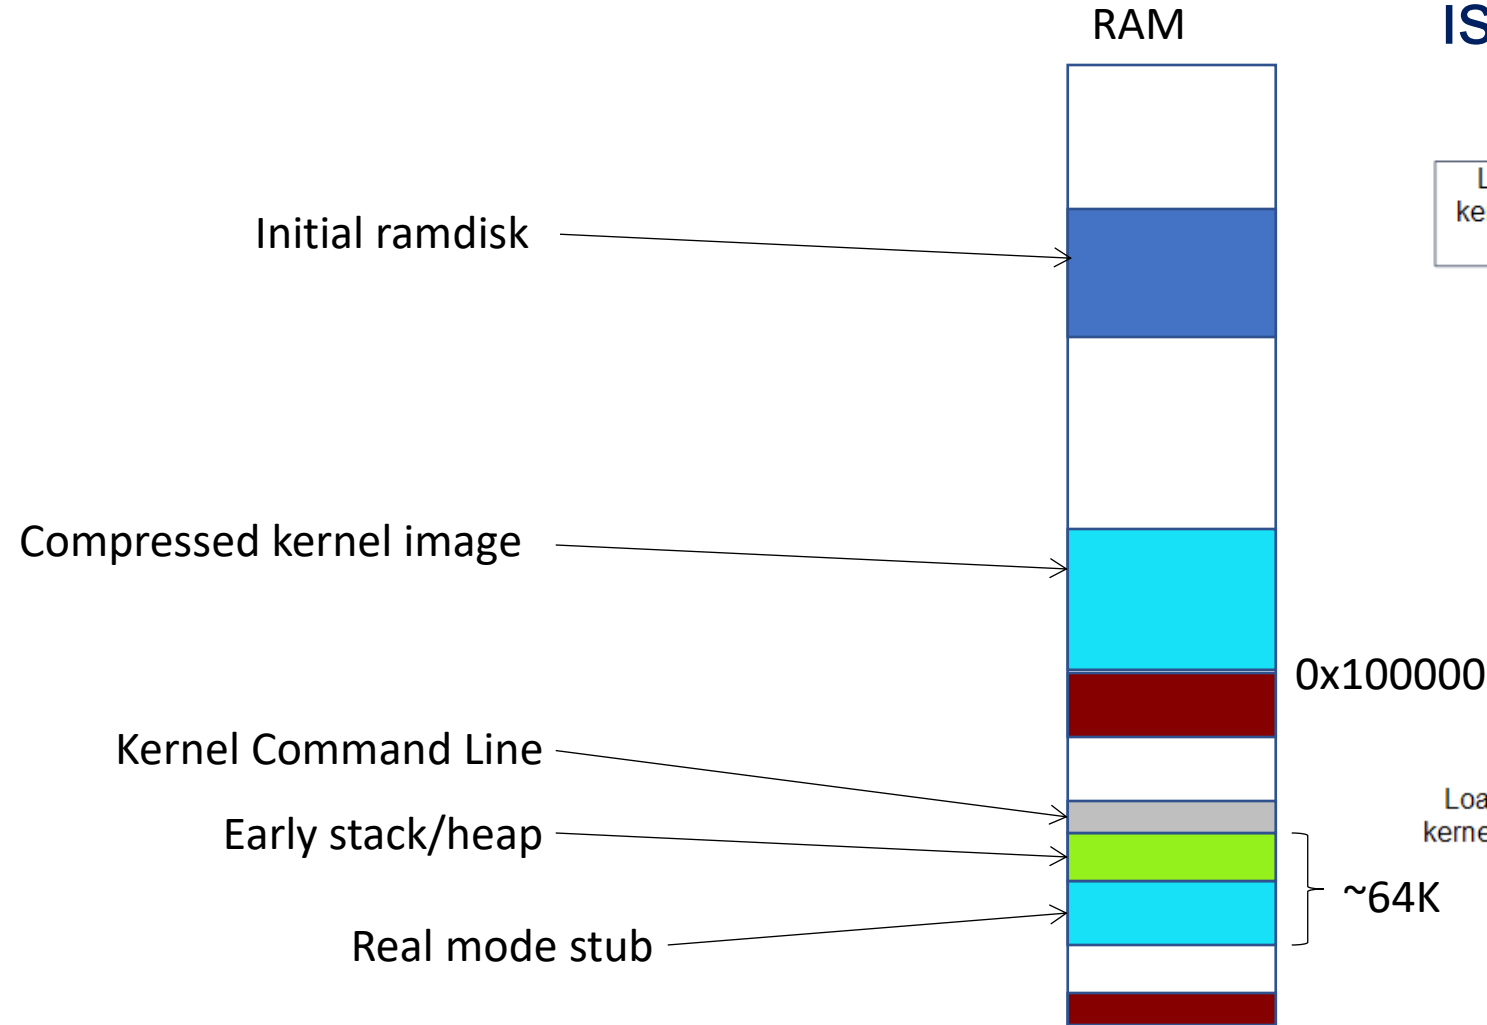
## should update-grub adjust the value of the default booted system
## can be true or false
# updatedefaultentry=false

title      Windows XP
root       (hd0,0)
savedefault
makeactive
chainloader +1

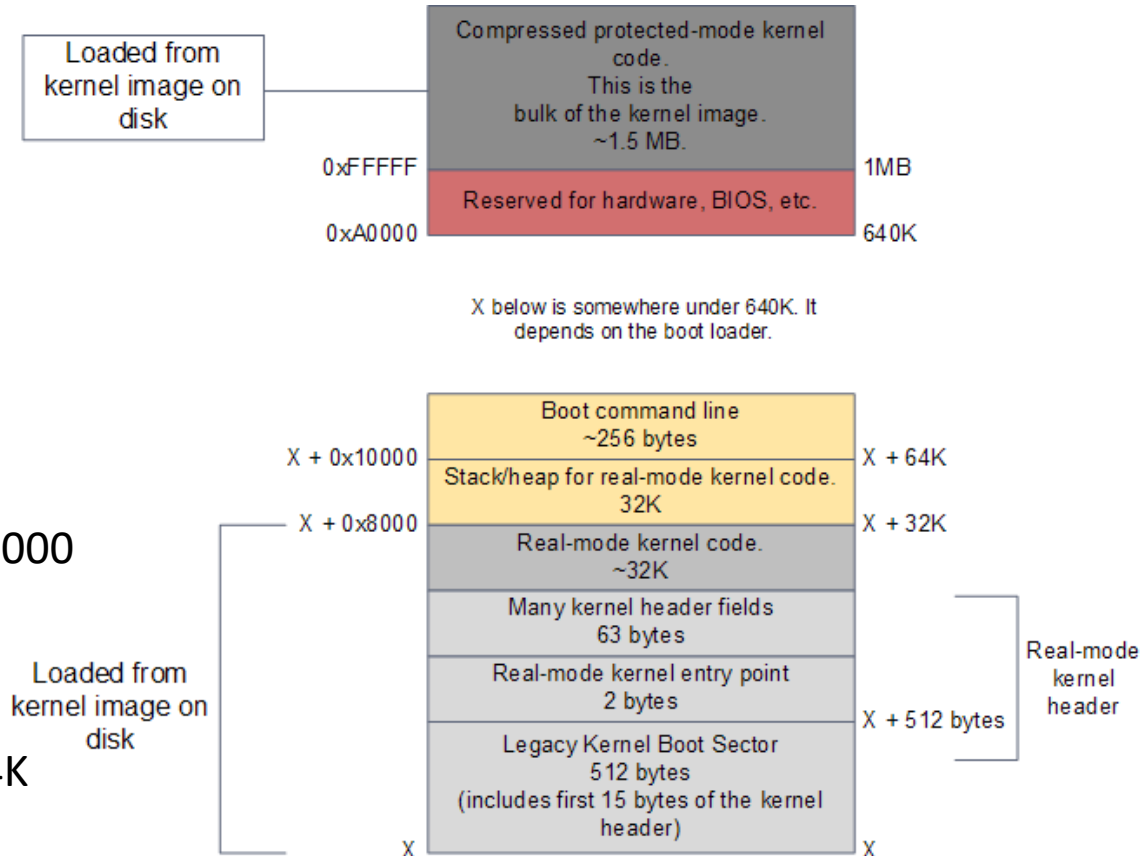
title      Ubuntu Linux
root       (hd0,2)
kernel    /boot/vmlinuz-2.6.15-25-386 root=/dev/hda3 ro quiet splash
initrd    /boot/initrd.img-2.6.15-25-386
boot
```

# So Kernel Time

- In the beginning...



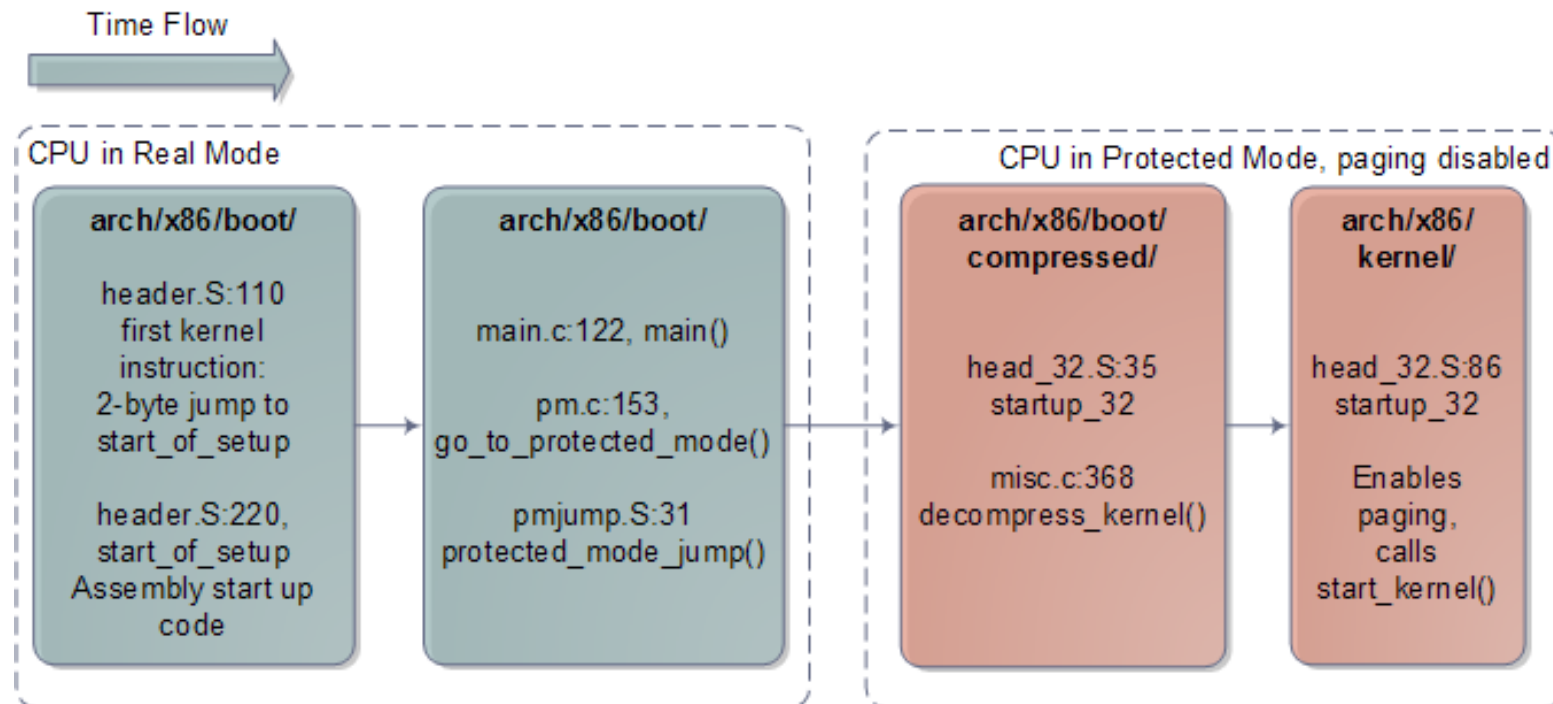
- RAM contents after boot loader is done





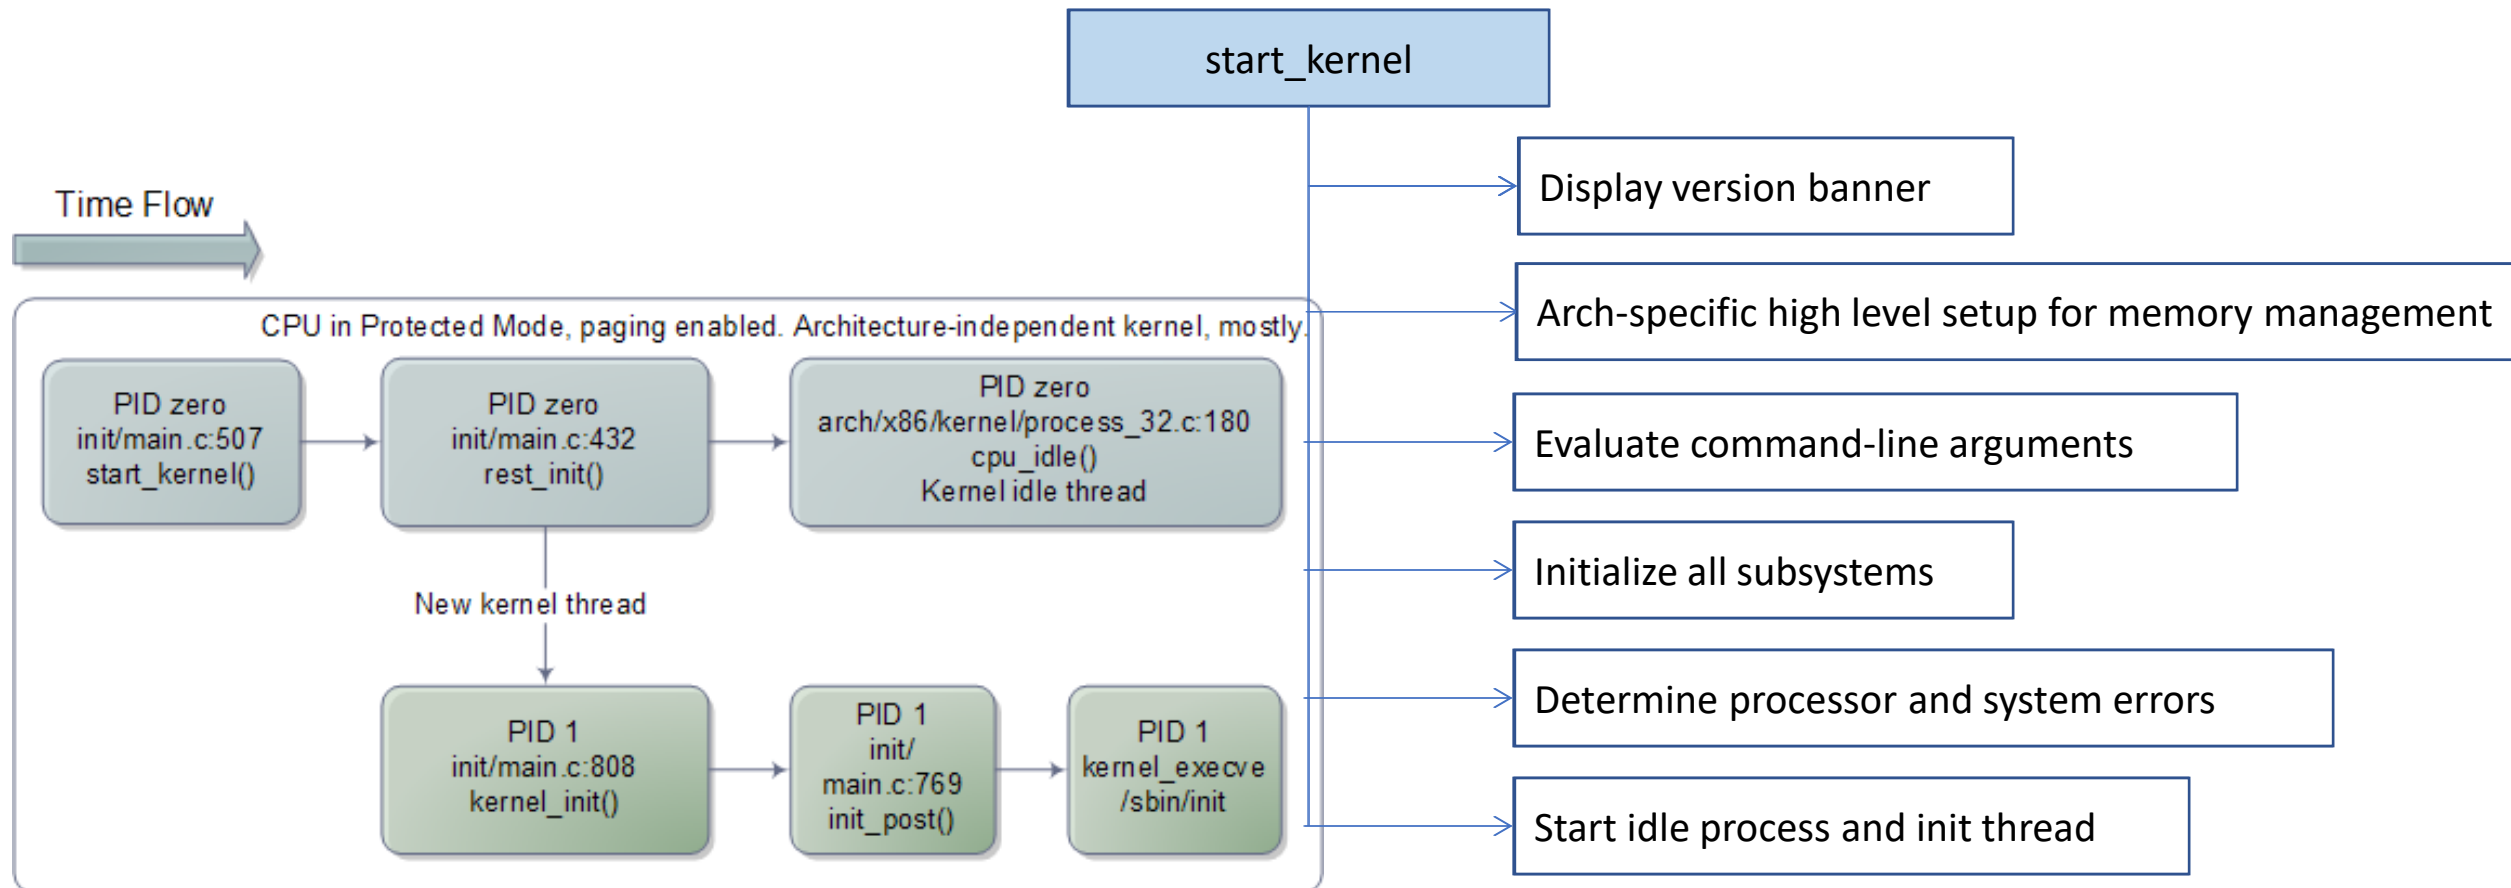
# *Major Steps in Kernel Initialization*

- 1) Platform-specific initialization  
(In assembly language)
- 2) Platform-independent initialization  
(In high-level language C)

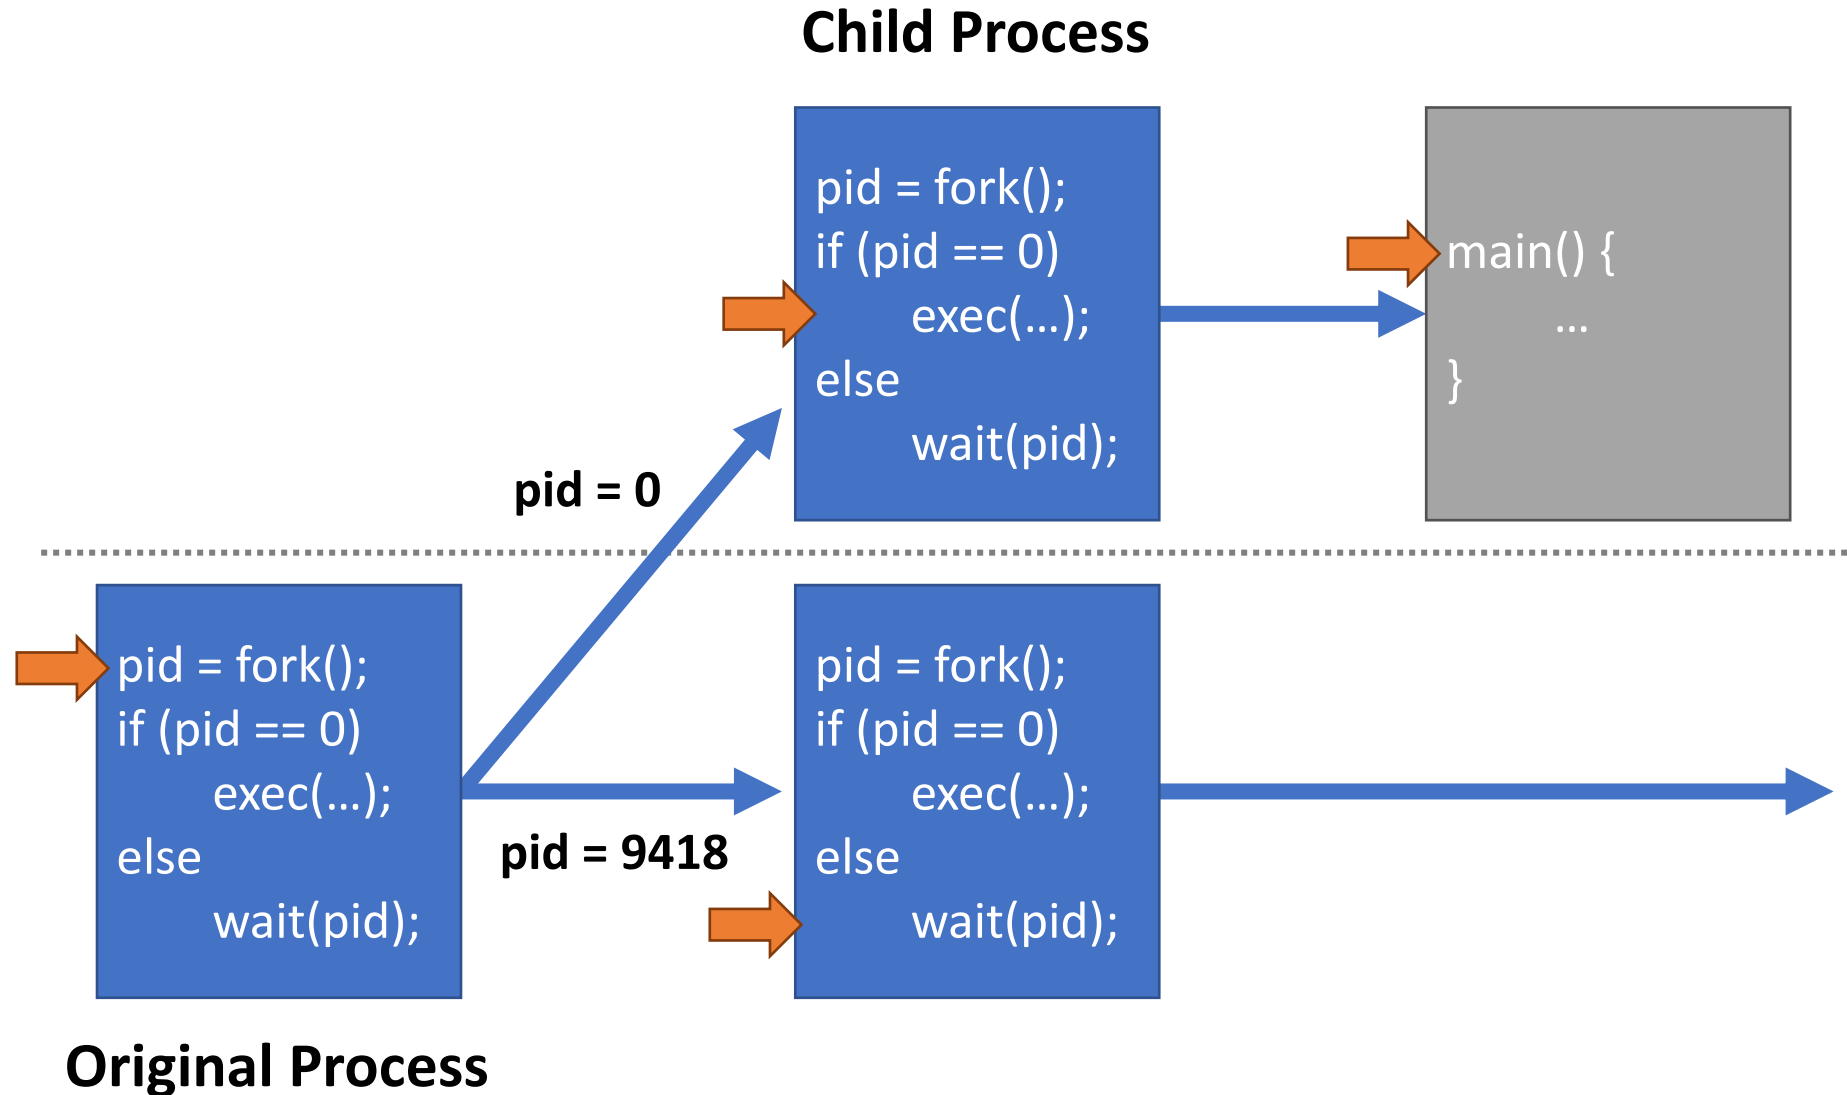


# Major Steps in Kernel Initialization

- Platform-independent initialization  
(In high-level language C)



# UNIX Process Management



# ***Question: What does this code print?***

```
int child_pid = fork();  
if (child_pid == 0) {           // I'm the child process  
    printf("I am process # %d\n", getpid());  
    return 0;  
} else {                       // I'm the parent process  
    printf("I am parent of process # %d\n", child_pid);  
    return 0;  
}
```