



## Reorder Buffer (RoB)

- Register values and memory values are not written until an instruction commits
- RoB effectively renames the destination registers
  - every destination gets a new entry in the RoB
- On misprediction:
  - Speculated entries in RoB are cleared
- Exceptions:
  - Not recognized/taken until it is ready to commit
  - Precise exceptions require that ‘later’ entries in RoB are cleared

# *Flynn\* Taxonomy, 1966*

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345 (Clovertown)

# Flynn\* Taxonomy, 1966

- In 2013, SIMD and MIMD most common parallelism in architectures – usually both in same system!
- Most common parallel processing programming style: Single Program Multiple Data (“SPMD”)
  - Single program that runs on all processors of a MIMD
  - Cross-processor execution coordination using synchronization primitives

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345 (Clovertown)

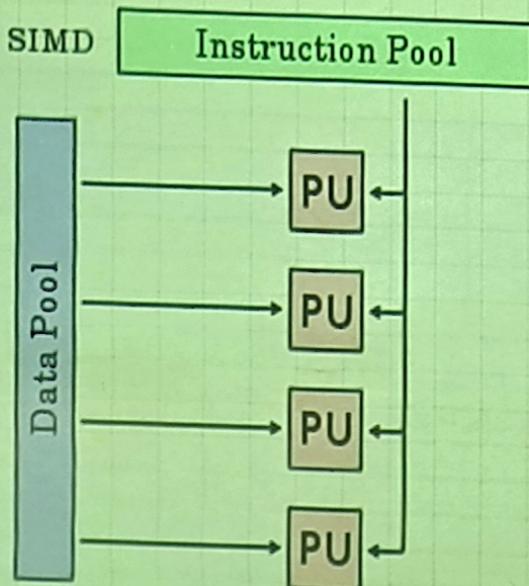
# Flynn\* Taxonomy, 1966

- In 2013, SIMD and MIMD most common parallelism in architectures – usually both in same system!
- Most common parallel processing programming style: Single Program Multiple Data (“SPMD”)
  - Single program that runs on all processors of a MIMD
  - Cross-processor execution coordination using synchronization primitives
- SIMD (aka hw-level *data parallelism*): specialized function units, for handling lock-step calculations involving arrays
  - Scientific computing, signal processing, multimedia (audio/video processing)

		Data Streams	
		Single	Multiple
Instruction Streams	Single	SISD: Intel Pentium 4	SIMD: SSE instructions of x86
	Multiple	MISD: No examples today	MIMD: Intel Xeon e5345 (Clovertown)

# Single-Instruction/Multiple-Data Stream (SIMD or “sim-dee”)

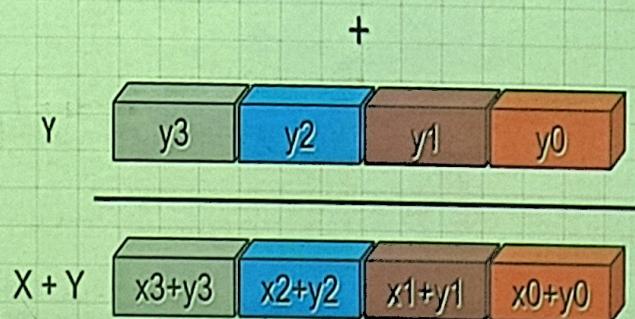
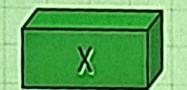
- SIMD computer exploits multiple data streams against a single instruction stream to operations that may be naturally parallelized, e.g., Intel SIMD instruction extensions or NVIDIA Graphics Processing Unit (GPU)





# SIMD: Single Instruction, Multiple Data

- Scalar processing
  - traditional mode
  - one operation produces one result





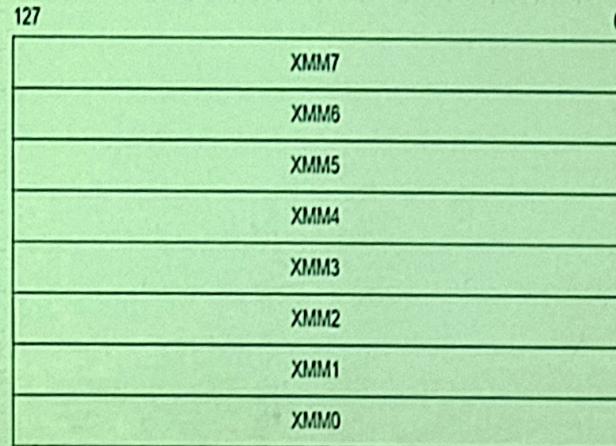
# What does this mean to you?

- In addition to SIMD extensions, the processor may have other special instructions
  - Fused Multiply-Add (FMA) instructions:
$$x = y + c * z$$
is so common some processor execute the multiply/add as a single instruction, at the same rate (bandwidth) as + or \* alone
- In theory, the compiler understands all of this
  - When compiling, it will rearrange instructions to get a good “schedule” that maximizes pipelining, uses FMAs and SIMD
  - It works with the mix of instructions inside an inner loop or other block of code
- But in practice the compiler may need your help
  - Choose a different compiler, optimization flags, etc.
  - Rearrange your code to make things more obvious
  - Using special functions (“intrinsics”) or write in assembly ☺



# Intel SIMD Extensions

- MMX 64-bit registers, reusing floating-point registers [1992]
- SSE2/3/4, new 8 128-bit registers [1999]

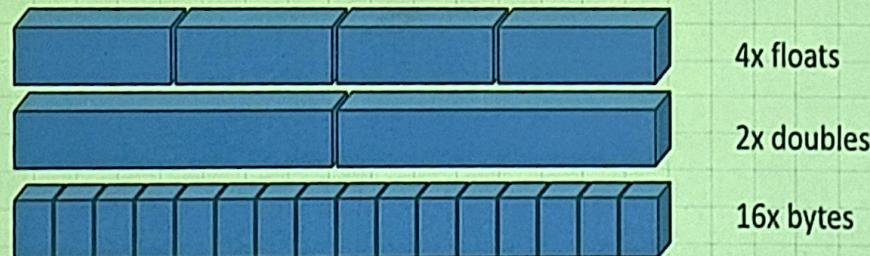


- AVX, new 256-bit registers [2011]
  - Space for expansion to 1024-bit registers

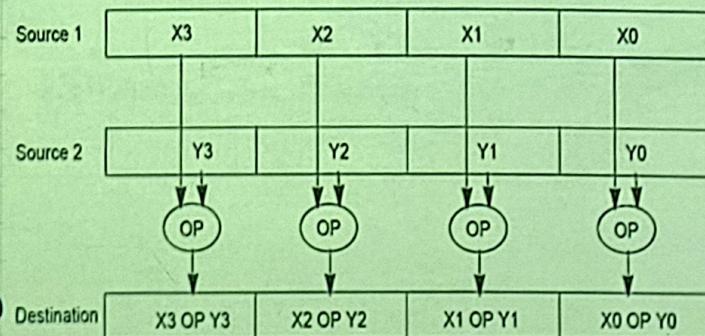


# SSE / SSE2 SIMD on Intel

- SSE2 data types: anything that fits into 16 bytes, e.g.,



- Instructions perform add, multiply etc. on all the data in parallel



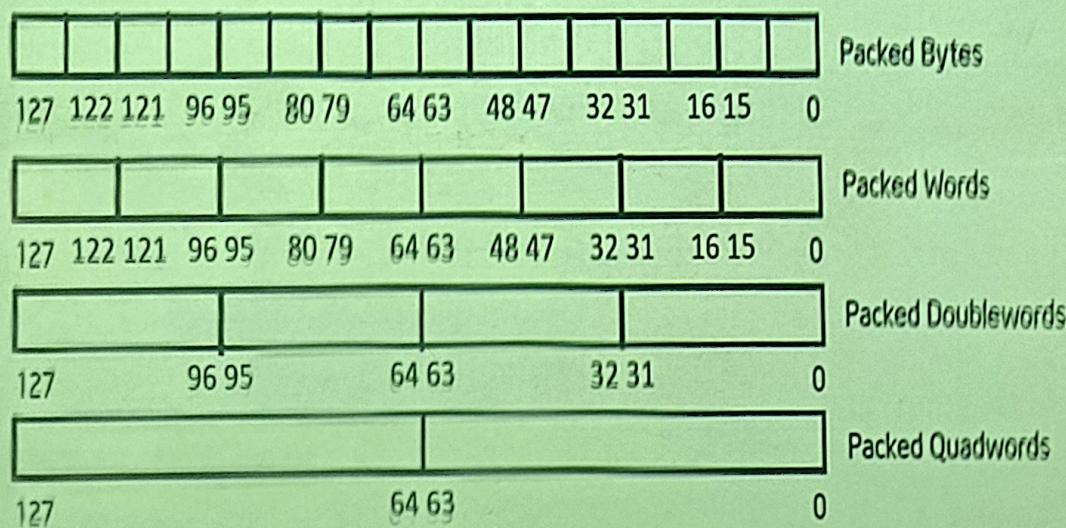
- Similar on GPUs, vector p erations)



# Intel Architecture SSE2+ 128-Bit SIMD Data Types

- Note: in Intel Architecture (unlike MIPS) a word is 16 bits
  - Single-precision FP: Double word (32 bits)
  - Double-precision FP: Quad word (64 bits)

Fundamental 128-Bit Packed SIMD Data Types





# Example: SIMD Array Processing

for each f in array  
f = sqrt(f)

for each f in array  
{  
load f to floating-point register  
calculate the square root  
write the result from the  
register to memory  
}

{ for each 4 members in array  
{  
load 4 members to the SSE register  
calculate 4 square roots in one operation  
store the 4 results from the register to memory  
}

SIMD style



## Data-Level Parallelism and SIMD

- SIMD wants adjacent values in memory that can be operated in parallel
- Usually specified in programs as loops

```
for (i=1000; i>0; i=i-1)
```

```
    x[i] = x[i] + s;
```

- How can reveal more data-level parallelism than available in a single iteration of a loop?
- *Unroll loop* and adjust iteration rate



## *Loop Unrolling in C*

- Instead of compiler doing loop unrolling, could do it yourself in C

```
for(i=1000; i>0; i=i-1)
```

```
    x[i] = x[i] + s;
```

- Could be rewritten

```
for(i=1000; i>0; i=i-4) {
```

```
    x[i] = x[i] + s;
```

```
    x[i-1] = x[i-1] + s;
```

```
    x[i-2] = x[i-2] + s;
```

```
    x[i-3] = x[i-3] + s;
```

```
}
```



# *Building a Python-C extension*

- Write the C function
  - PyObject
  - PyArg\_ParseTuple()
  - PyLong\_FromLong()
  - First import calls this



# Building a Python-C extension

- Write the C function
  - PyObject
  - PyArg\_ParseTuple()
  - PyLong\_FromLong()

- Write the init function
  - PyMethodDef
  - PyModuleDef

- First import calls this

```
#include <Python.h>

static PyObject *method_fputs(PyObject *self, PyObject *args) {
    char *str, *filename = NULL;
    int bytes_copied = -1;

    /* Parse arguments */
    if(!PyArg_ParseTuple(args, "ss", &str, &filename)) {
        return NULL;
    }

    FILE *fp = fopen(filename, "w");
    bytes_copied = fputs(str, fp);
    fclose(fp);

    return PyLong_FromLong(bytes_copied);
}

static PyMethodDef FputsMethods[] = {
    {"fputs", method_fputs, METH_VARARGS, "Python interface for fputs
    (NULL, NULL, 0, NULL)}
};

static struct PyModuleDef fputsmodule = {
    PyModuleDef_HEAD_INIT,
    "fputs",
    "Python interface for the fputs C library function",
    -1,
    FputsMethods
};
```

# Python-C extension

```
#include <Python.h>
```

```
static PyObject *method_fputs(PyObject *self, PyObject *args) {
    char *str, *filename = NULL;
    int bytes_copied = -1;
```

```
/* Parse arguments */
```

```
if(!PyArg_ParseTuple(args, "ss", &str, &filename)) {
    return NULL;
}
```

```
FILE *fp = fopen(filename, "w");
bytes_copied = fputs(str, fp);
fclose(fp);
```

```
return PyLong_FromLong(bytes_copied);
}
```

```
static PyMethodDef FputsMethods[] = {
    {"fputs", method_fputs, METH_VARARGS, "..."},
```

```
FILE *fp = fopen(filename, "w");
bytes_copied = fputs(str, fp);
fclose(fp);

return PyLong_FromLong(bytes_copied);

static PyMethodDef FputsMethods[] = {
    {"fputs", method_fputs, METH_VARARGS, "Python interface for fputs
NULL, NULL, 0, NULL}
};

static struct PyModuleDef fputsmodule = {
PyModuleDef_HEAD_INIT,
"fputs",
"Python interface for the fputs C library function",
-1,
FputsMethods
};
```

- Write the init function

- PyMethodDef
- PyModuleDef

- PyMODINIT\_FUNC

- First import calls this

```
PyMODINIT_FUNC PyInit_fputs(void) {
    return PyModule_Create(&fputsmodule);
}
```

```
FILE *fp = fopen(filename, "w");
bytes_copied = fputs(str, fp);
fclose(fp);

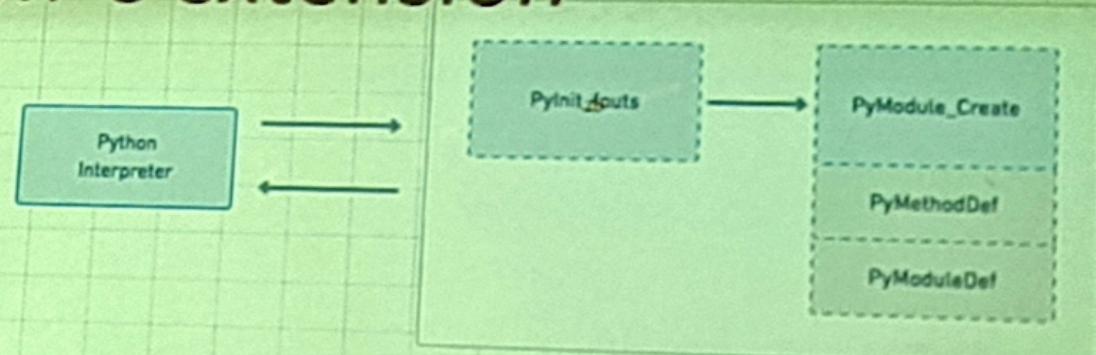
return PyLong_FromLong(bytes_copied);
}

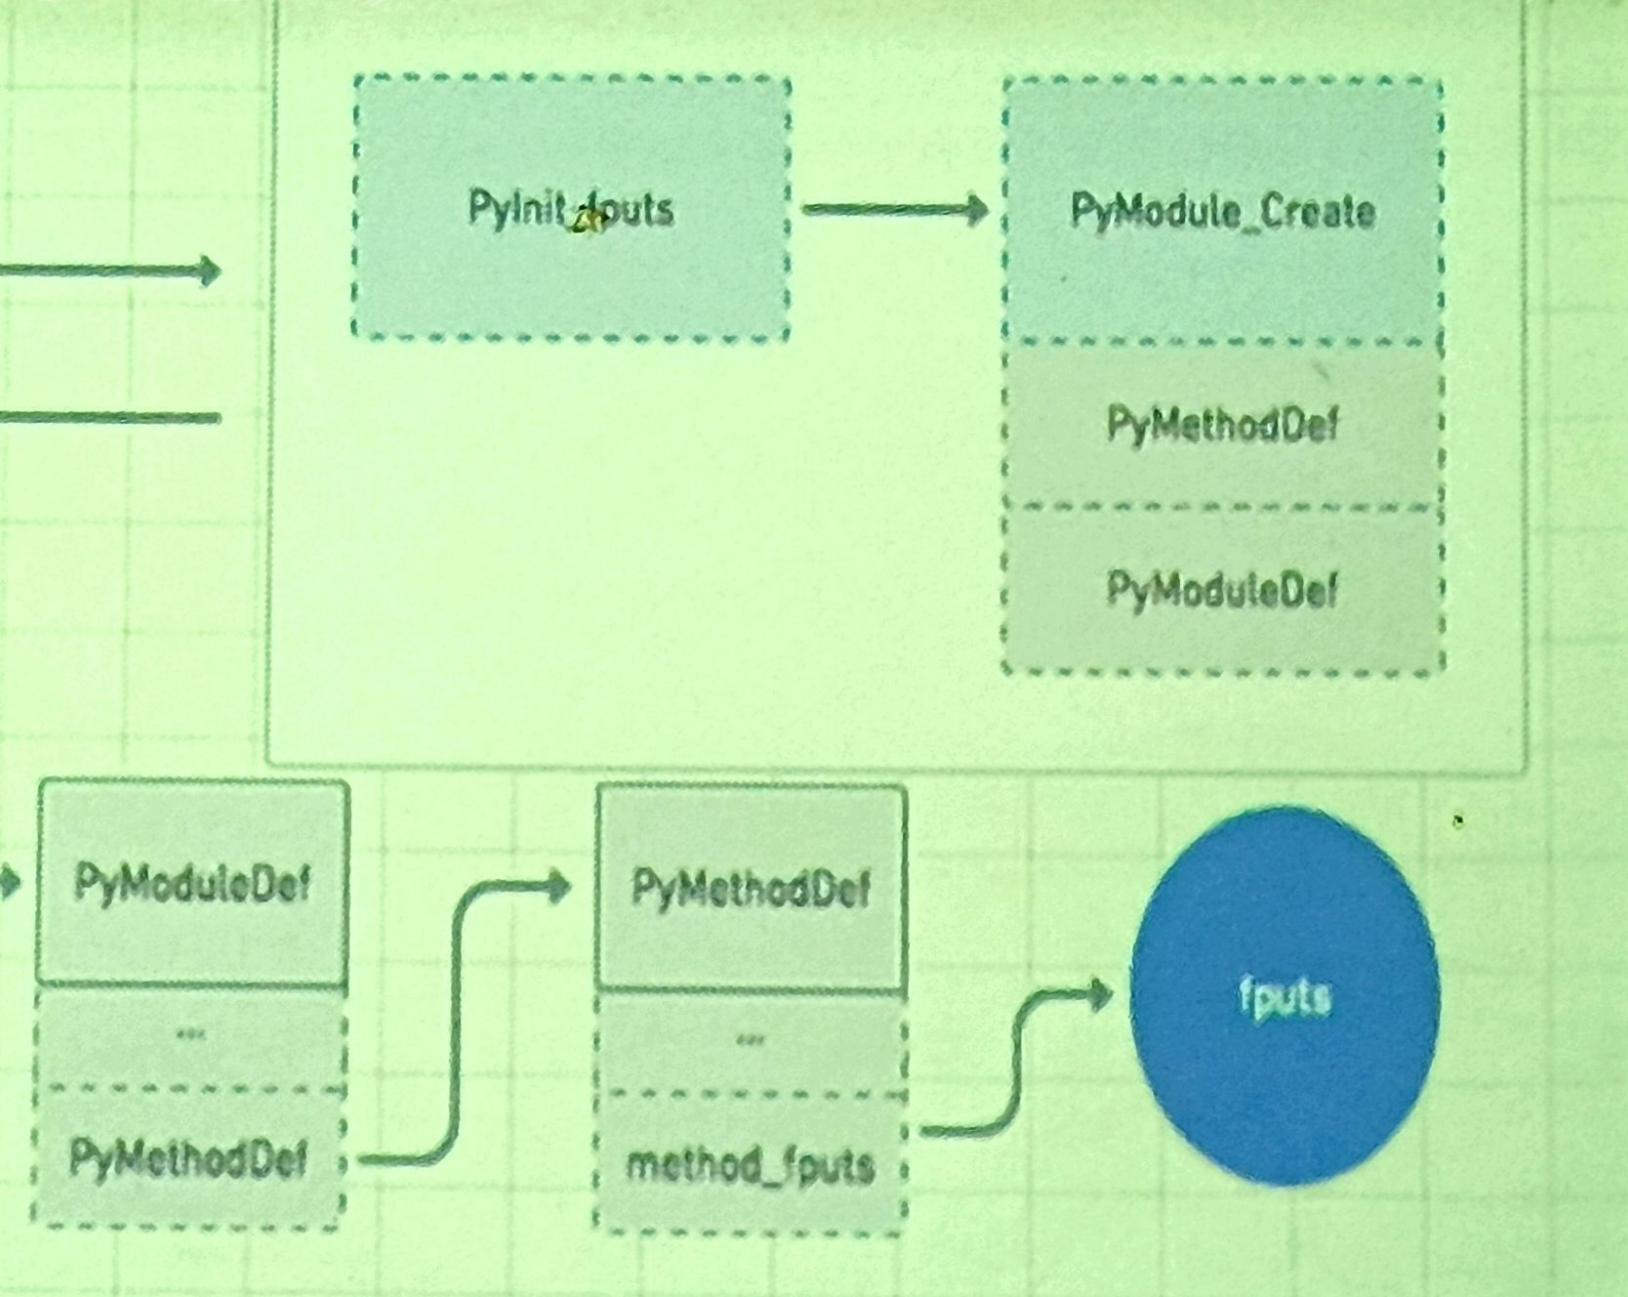
static PyMethodDef FputsMethods[] = {
    {"fputs", method_fputs, METH_VARARGS, "Python interface for fputs
    (NULL, NULL, 0, NULL)
};

static struct PyModuleDef fputsmodule = {
    PyModuleDef_HEAD_INIT,
    "fputs",
    "Python interface for the fputs C library function",
    -1,
    FputsMethods
};
```

# *Building a Python-C extension*

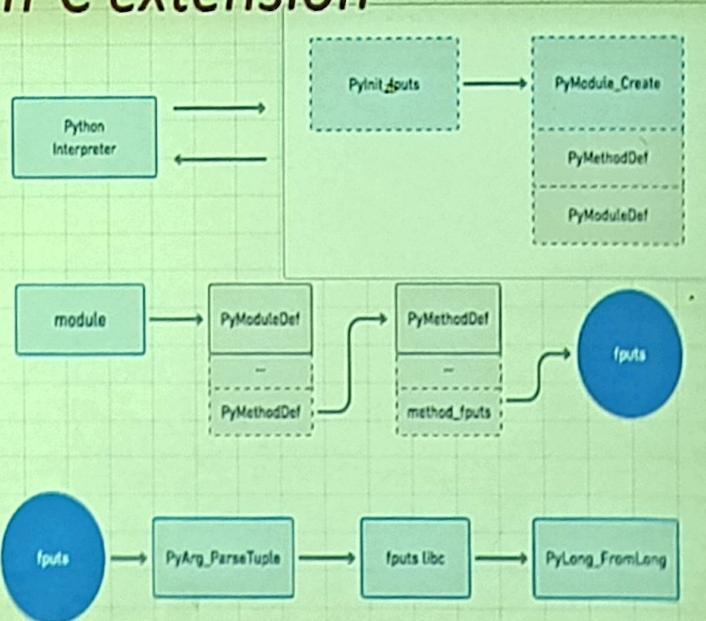
- This is how it appears





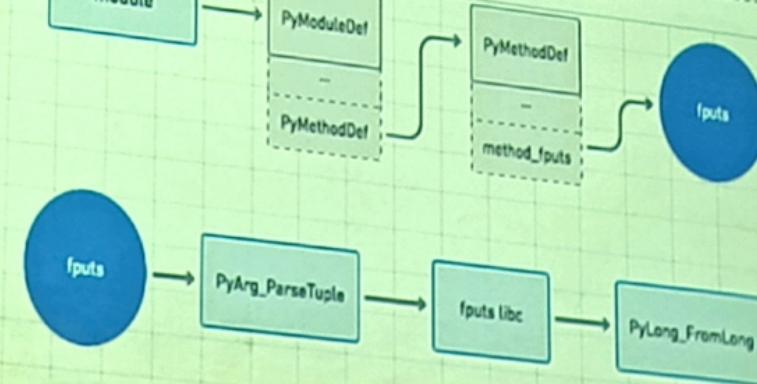
# *Building a Python-C extension*

- This is how it appears
  - Create setup.py
  - Compile and Install
    - python3 setup.py install



- python3\_setup.py install

Use in your code



```
from distutils.core import setup, Extension
def main():
    setup(name="fputs",
          version="1.0.0",
          description="Python interface for the fputs C library function",
          author=<your name>,
          author_email=<your_email@gmail.com>,
          ext_modules=[Extension("fputs", ["fputsmodule.c"])])
if __name__ == "__main__":
    main()
```

# • Use in your code

```
>>> import fput
>>> fput.__doc__
'Python interface for the fput C library function'
>>> fput.__name__
'fput'
>>> # Write to an empty file named 'write.txt'
>>> fput.fput("Real Python!", "write.txt")
13
>>> with open("write.txt", "r") as f:
>>>     print(f.read())
'Real Python!'
```

```

graph LR
    fput((fput)) --> PyArg_ParseTuple[PyArg_ParseTuple]
    PyArg_ParseTuple --> fputLibc[fput.libc]
    fputLibc --> PyLong_FromLong[PyLong_FromLong]
  
```

```

from distutils.core import setup, Extension
def main():
    setup(name="fput",
          version="1.0.0",
          description="Python interface for the fput C library function",
          author=<your name>,
          author_email=<your_email@gmail.com>,
          ext_modules=[Extension("fput", ["fputmodule.c"])])
if __name__ == "__main__":
    main()
  
```



# Accelerate matrixMultiply

```
#include <x86intrin.h>
#include <sys/time.h>
#include <stdlib.h>
#include <stdint.h>
#include <stdio.h>
#include <math.h>
#include <mm_malloc.h>

unsigned long long get_timestamp ()
{
    struct timeval now;
    gettimeofday (&now, NULL);
    return now.tv_usec + (unsigned long long)now.tv_sec * 1000000;
}
```

```
void dgemm_avx256(const uint32_t n, const double* A, const double* B,
{
    for( uint32_t i = 0; i < n; i += 4 )
    {
        for( uint32_t j = 0; j < n; j++ )
        {
            m256d c0 = _mm256_load_pd(C + i + j * n); /* c0 = C[i][j] */
            for( uint32_t k = 0; k < n; k++ )
            {
                c0 = _mm256_add_pd(c0, /* c0 += A[i][k]*B[k][j] */ _mm256_mul_pd(_mm256_load_pd(A + i + k * n), B + k + j * n));
            }
            _mm256_store_pd(C + i + j * n, c0); /* C[i][j] = c0 */
        }
    }
}
```

# Accelerate matrixMultiply

```
n.h>
.h>
.h>
    void dgemm_avx256(const uint32_t n, const double* A, const double* B, double* C)
    {
        for( uint32_t i = 0; i < n; i += 4 )
        {
            for( uint32_t j = 0; j < n; j++ )
            {
                m256d c0 = _mm256_load_pd(C + i + j * n); /* c0 = C[i][j] */
                for( uint32_t k = 0; k < n; k++ )
                {
                    c0 = _mm256_add_pd(c0, /* c0 += A[i][k]*B[k][j] */
                        _mm256_mul_pd( _mm256_load_pd(A + i + k * n), _mm256_broadcast_sd(B + k + j * n) )
                }
                _mm256_store_pd(C+i+j*n, c0); /* C[i][j] = c0 */
            }
        }
    }
};
```

```
long long now.tv_sec * 1000000; } } _mm256_store_pd(C + i + j * n, c0); /* C[i][j] = c0 */ }

void dgemm_avx512(const uint32_t n, const double* A, const double* B, double* C)
{
    for( uint32_t i = 0; i < n; i += 8)
    {
        for( uint32_t j = 0; j < n; ++j)
        {
            m512d c0 = _mm512_load_pd(C + i + j * n); // c0 = C[i][j]
            for( uint32_t k = 0; k < n; k++)
            {
                // c0 += A[i][k] * B[k][j]
                m512d bb = _mm512_broadcast_sd_pd(_mm_load_sd(B + j * n + k));
                c0 = _mm512_fmaddd_pd(_mm512_load_pd(A + n * k + i), bb, c0);
            }
            _mm512_store_pd(C + i + j * n, c0); // C[i][j] = c0
        }
    }
}
```

# Accelerate matrixMultiply

```
#include <x86intrin.h>
#include <sys/time.h>
#include <stdlib.h>
#include <stdint.h>
#include <stdio.h>
#include <math.h>
#include <mm_malloc.h>

unsigned long long get_timestamp ()
{
    struct timeval now;
    gettimeofday (&now, NULL);
    return now.tv_usec + (unsigned long long)now.tv_sec * 1000000;
}

void dgemm_avx256(const uint32_t n, const double* A, const double* B, double* C)
{
    for( uint32_t i = 0; i < n; i += 4 )
    {
        for( uint32_t j = 0; j < n; j++ )
        {
            m256d c0 = _mm256_load_pd(C + i + j * n); /* c0 = C[i][j] */
            for( uint32_t k = 0; k < n; k++ )
            {
                c0 = _mm256_add_pd(c0, /* c0 += A[i][k]*B[k][j] */
                    _mm256_mad_pd(_mm256_load_pd(A + i + k * n), _mm256_broadcast_sd(B + k + j * n) ) );
            }
            _mm256_store_pd(C+i+j*n, c0); /* C[i][j] = c0 */
        }
    }
}

void dgemm_avx512(const uint32_t n, const double* A, const double* B, double* C)
{
    for( uint32_t i = 0; i < n; i += 8 )
    {
        for( uint32_t j = 0; j < n; ++j)
        {
            m512d c0 = _mm512_load_pd(C + i + j * n); // c0 = C[i][j]
            for( uint32_t k = 0; k < n; k++)
            {
                // c0 += A[i][k] * B[k][j]
                m512d bb = _mm512_broadcastsd_pd(_mm_load_sd(B + j * n + k));
                c0 = _mm512_fmaddd_pd(_mm512_load_pd(A + n * k + i), bb, c0);
            }
            _mm512_store_pd(C + i + j * n, c0); // C[i][j] = c0
        }
    }
}
```

```
void dgemm_basic(const uint32_t n, const double* A, const double* B, double* C)
{
    for(uint32_t i = 0; i < n; ++i)
    {
        for(uint32_t j = 0; j < n; ++j)
        {
            double cij = C[i + j * n]; /* cij = C[i][j] */
            for(uint32_t k = 0; k < n; k++)
            {
                cij += A[i + k * n] * B[k + j * n]; /* cij += A[i][k]*B[k][j] */
            }
            C[i + j * n] = cij; /* C[i][j] = cij */
        }
    }
}
```

```
void dgemm_avx512(const uint32_t n, const double* A, const double* B, double* C)
{
    for( uint32_t i = 0; i < n; i += 8)
    {
        for( uint32_t j = 0; j < n; ++j)
        {
            m512d c0 = _mm512_load_pd(C + i + j * 8);
            for( uint32_t k = 0; k < n; k++)
            {
                // c0 += A[i][k] * B[k][j]
                m512d bb = _mm512_broadcastsd(C0);
                c0 = _mm512_fmadd_pd(_mm512_loa
            }
            _mm512_store_pd(C + i + j * n, c0);
        }
    }
}
```

# Accelerate matrixMultiply

```
int main(){
    uint32_t trial_no = 11;
    uint32_t n = 32 * 20;

    double *a;
    double *b;
    double *c;
    unsigned long long t0;
    unsigned long long t1;
    unsigned long long t;
    struct timeval tv1, tv2;

    double result;

    a = (double*) _mm_malloc (n * n * sizeof(double), 64);
    b = (double*) _mm_malloc (n * n * sizeof(double), 64);
    c = (double*) _mm_malloc (n * n * sizeof(double), 64);

    for(uint32_t i = 0; i < n * n; ++i)
    {
        a[i] = rand();
    }

    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_basic(n, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }

    result = result / (double)(trial_no);
    printf("Time Taken: Ref:\t%lf\n",result);

    result=0;
    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_avx256(n, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }

    result = result / (double)(trial_no);
    printf("Time Taken: AVX256:\t%lf\n",result);
}
```

# Accelerate matrixMultiply

```
int main()
{
    uint32_t trial_no = 11;
    uint32_t n = 32 * 20;

    double *a;
    double *b;
    double *c;
    unsigned long long t0;
    unsigned long long t1;
    unsigned long long t;
    struct timeval tv1, tv2;

    double result;

    double* _mm_malloc (n * n * sizeof(double), 64);
    double* _mm_malloc (n * n * sizeof(double), 64);
    double* _mm_malloc (n * n * sizeof(double), 64);

    (uint32_t i = 0; i < n * n; ++i)
        a[i] = rand();

    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_basic(n, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }

    result = result / (double)(trial_no);
    printf("Time Taken: Ref:\t%f\n", result);

    result=0;
    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_avx256(n, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }

    result = result / (double)(trial_no);
    printf("Time Taken: AVX256:\t%f\n", result);
}
```

```
double* _mm_malloc (n * n * sizeof(double), 64);
double* _mm_malloc (n * n * sizeof(double), 64);
double* _mm_malloc (n * n * sizeof(double), 64);

(uint32_t i = 0; i < n * n; ++i)
    a[i] = rand();

result=0;
for(uint32_t i = 0; i < trial_no; i++)
{
    gettimeofday(&tv1, NULL);
    dgemm_avx256(n, a, b, c);
    gettimeofday(&tv2, NULL);
    result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
}
result = result / (double)(trial_no);
printf("Time Taken: AVX256:\t%f\n",result);

result=0;
for(uint32_t i = 0; i < trial_no; i++)
{
    gettimeofday(&tv1, NULL);
    dgemm_avx512(n, a, b, c);
    gettimeofday(&tv2, NULL);
    result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
}
result = result / (double)(trial_no);
printf("Time Taken: AVX512:\t%f\n",result);
```

```
for(uint32_t i = 0; i < n * n; ++i)
{
    a[i] = rand();
}
```

```
Kollegmtinerva: /asset/extes/CONAVS$ ./a.out
Time Taken: Ref:      1.027594
Time Taken: AVX256:   0.436642
Time Taken: AVX512:   0.412659
```

```
gcc -Wall -pedantic matmul-3.c -march=icelake-server
```

```
gettimeofday(&tv1, NULL);
dgemm_avx256(n, a, b, c);
gettimeofday(&tv2, NULL);
result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 +
}
result = result / (double)(trial_no);
printf("Time Taken: AVX256:\t%f\n",result);

result=0;
for(uint32_t i = 0; i < trial_no; i++)
{
    gettimeofday(&tv1, NULL);
    dgemm_avx512(n, a, b, c);
    gettimeofday(&tv2, NULL);
    result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 +
}
result = result / (double)(trial_no);
printf("Time Taken: AVX512:\t%f\n",result);
```

# Accelerate matrixMultiply

```
int main()
{
    uint32_t trial_no = 11;
    uint32_t n = 32 * 20;

    double *a;
    double *b;
    double *c;
    unsigned long long t0;
    unsigned long long t1;
    unsigned long long t;
    struct timeval tv1, tv2;

    double result;

    a = (double*) _mm_malloc (n * n * sizeof(double), 64);
    b = (double*) _mm_malloc (n * n * sizeof(double), 64);
    c = (double*) _mm_malloc (n * n * sizeof(double), 64);

    for(uint32_t i = 0; i < n * n; ++i)
    {
        a[i] = rand();
    }

    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_basic(n, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }

    result = result / (double)(trial_no);
    printf("Time Taken: Ref:\t%f\n", result);

    result=0;
    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_avx256(n, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }

    result = result / (double)(trial_no);
    printf("Time Taken: AVX256:\t%f\n", result);

    result=0;
    for(uint32_t i = 0; i < trial_no; i++)
    {
        gettimeofday(&tv1, NULL);
        dgemm_avx512(n, a, b, c);
        gettimeofday(&tv2, NULL);
        result += (double) (tv2.tv_usec - tv1.tv_usec) / 1000000 + (double) (tv2.tv_sec - tv1.tv_sec);
    }

    result = result / (double)(trial_no);
    printf("Time Taken: AVX512:\t%f\n", result);
}
```

```
kolin@minerva:~/acceleration/tw$ ./a.out
Time Taken: Ref:      1.027594
Time Taken: AVX256:   0.436642
Time Taken: AVX512:   0.412659
```

# Using shared Objects

```
from ctypes import CDLL, POINTER
from ctypes import c_size_t, c_double, c_uint
import numpy as np
from numpy import random
import time

# load the library
mylib = CDLL("/home/kolin/accel/extension/two/mylib.so")

# C-type corresponding to numpy array
ND_POINTER_1 = np.ctypeslib.ndpointer(dtype=np.float64,
                                       ndim=1,
                                       flags="C")

# define prototypes
mylib.print_array.argtypes = [ND_POINTER_1 , c_size_t]
mylib.print_array.restype = None
mylib.dgemm_basic.argtypes = [c_uint, ND_POINTER_1 , ND_POINTER_1 , ND_POINTER_1 ]
mylib.dgemm_basic.restype = None
mylib.dgemm_avx512.argtypes = [c_uint, ND_POINTER_1 , ND_POINTER_1 , ND_POINTER_1]
mylib.dgemm_avx512.restype = None
mylib.dgemm_avx256.argtypes = [c_uint, ND_POINTER_1 , ND_POINTER_1 , ND_POINTER_1]
mylib.dgemm_avx256.restype = None
```

```
gcc -Wall -pedantic -shared -fPIC -o mylib.so try.c -march=icelake-server
```

```
cd /opt/kutin/accel/extension/two/mylib.so
g++ -fPIC -c try.c -o mylib.so -march=icelake-server
[INTER_1 = np.ctypeslib.ndpointer(dtype=np.float64,
                                    ndim=1,
                                    flags="C")]
The prototypes
print_array.argtypes = [ND_POINTER_1 , c_size_t]
print_array.restype = None
gemm_basic.argtypes = [c_uint, ND_POINTER_1 , ND_POINTER_1 , ND_POINTER_1 ]
gemm_basic.restype = None
emm_avx512.argtypes = [c_uint, ND_POINTER_1 , ND_POINTER_1 , ND_POINTER_1 ]
emm_avx512.restype = None
emm_avx256.argtypes = [c_uint, ND_POINTER_1 , ND_POINTER_1 , ND_POINTER_1 ]
emm_avx256.restype = None
n=32*20
#s = (2,2)
s=n*n
a = random.rand(s)
b =random.rand(s)
c =np.zeros(s)
X=np.ones(s)
# call function
mylib.print_array(X, X.size)
start=time.time()
mylib.dgemm_basic(n,a,b,c)
end=time.time()
print(c)
print("Elapsed Time Ref:", end-start)
c =np.zeros(s)
start1=time.time()
mylib.dgemm_avx256(n,a,b,c)
end1=time.time()
print("Elapsed Time AVX256:", end1-start1)
```



## *Graphical Pipeline:*

- Set of steps need to be taken in order to turn 3D scene into 2D image
  - The graphical pipeline is conceptual model
  - Highly dependent on the underlying available Software and Hardware accelerators

# Graphical Pipeline:

- Set of steps need to be taken in order to turn 3D scene into 2D image
  - The graphical pipeline is conceptual model
  - Highly dependent on the underlying available Software and Hardware accelerators
- The model of graphical pipeline is usually used in real-time rendering
  - Each step is backed with efficient algorithms that are usually hardware accelerated (e.g., using GPUs)

- The graphical pipeline is conceptual model
- Highly dependent on the underlying available Software and Hardware accelerators
- The model of graphical pipeline is usually used in real-time rendering
  - Each step is backed with efficient algorithms that are usually hardware accelerated (e.g., using GPUs)
- How do you represent 3D surfaces ?
- A 3D mesh (usually triangles):
  - Simple objects can be composed of thousands of triangles!



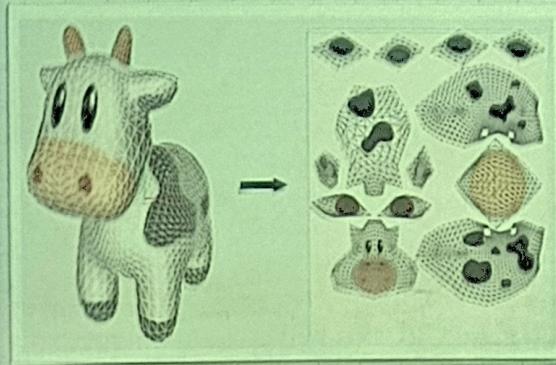
## *Graphical Pipeline:*

- Set of steps need to be taken in order to turn 3D scene into 2D image
  - The graphical pipeline is conceptual model
  - Highly dependent on the underlying available Software and Hardware accelerators
- The model of graphical pipeline is usually used in real-time rendering
  - Each step is backed with efficient algorithms that are usually hardware accelerated (e.g., using GPUs)
- How do you represent 3D surfaces ?
- A 3D mesh (usually triangles):
  - Simple objects can be composed of thousands of triangles!

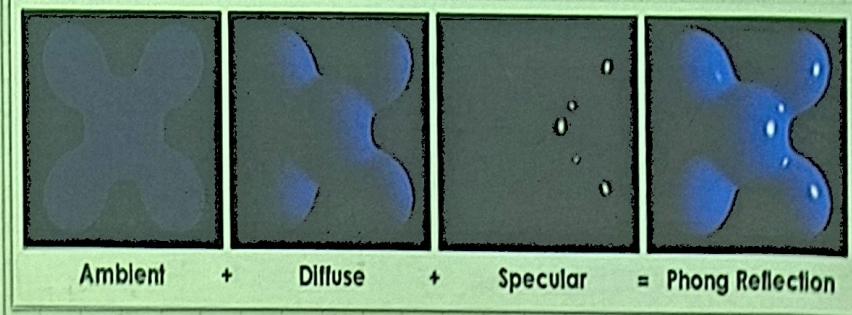


# Graphic Processing – Some History

- 1990s: Real-time 3D rendering for video games were becoming common
  - Doom, Quake, Descent, ... (Nostalgia!)
- 3D graphics processing is immensely computation-intensive

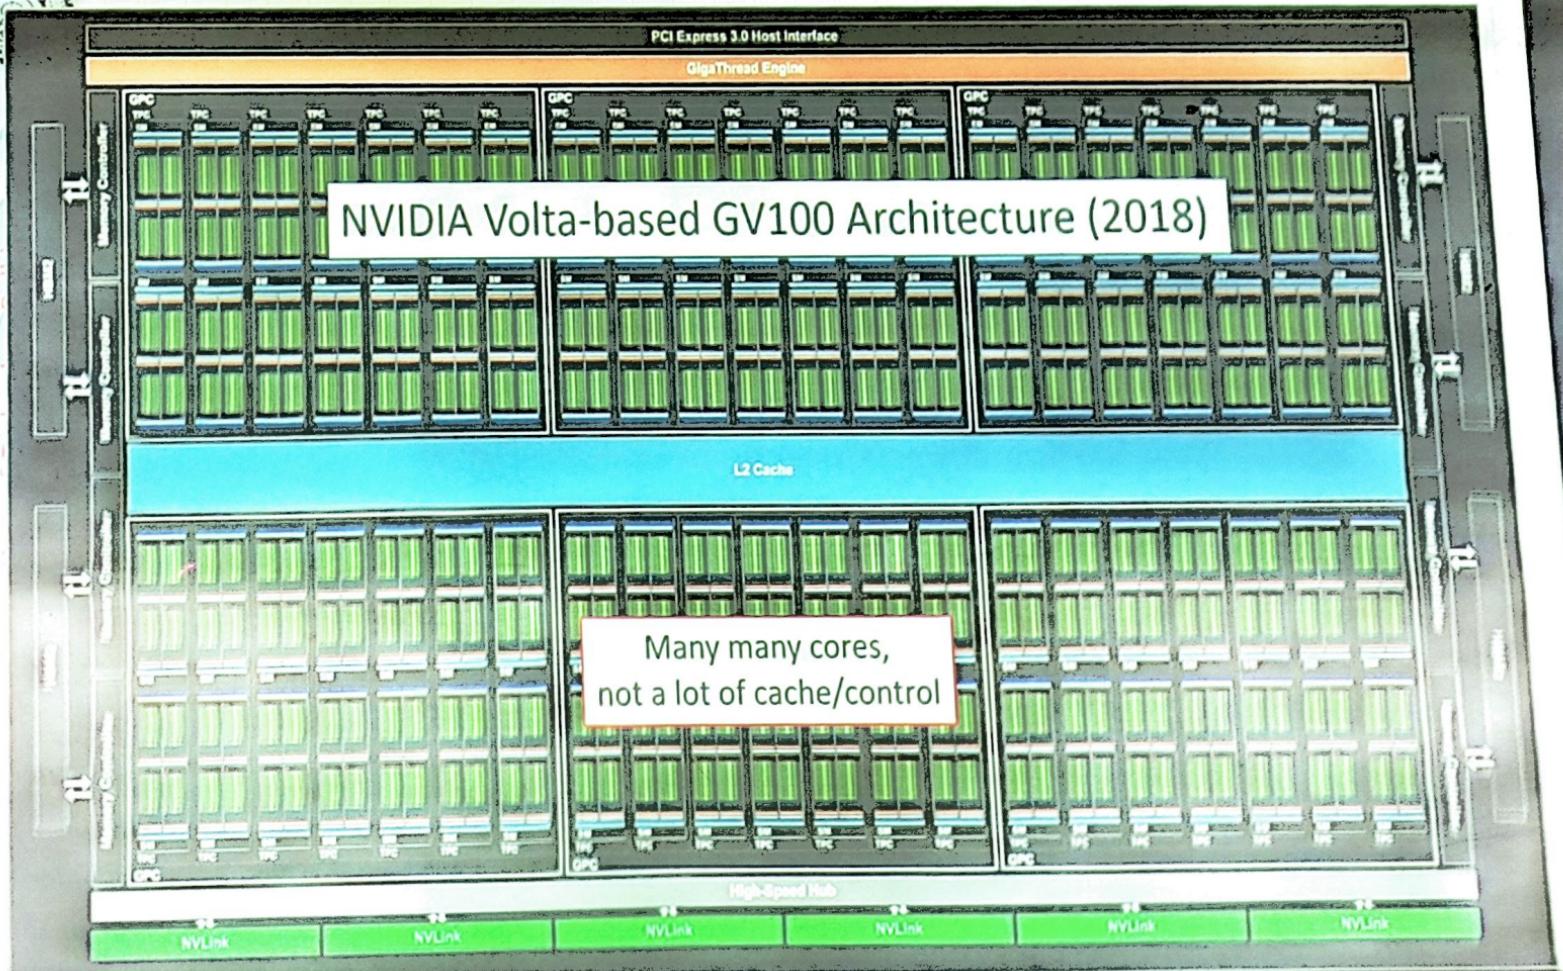


Texture mapping



Shading

Warren Moore, "Textures and Samplers in Metal," Metal by Example, 2014

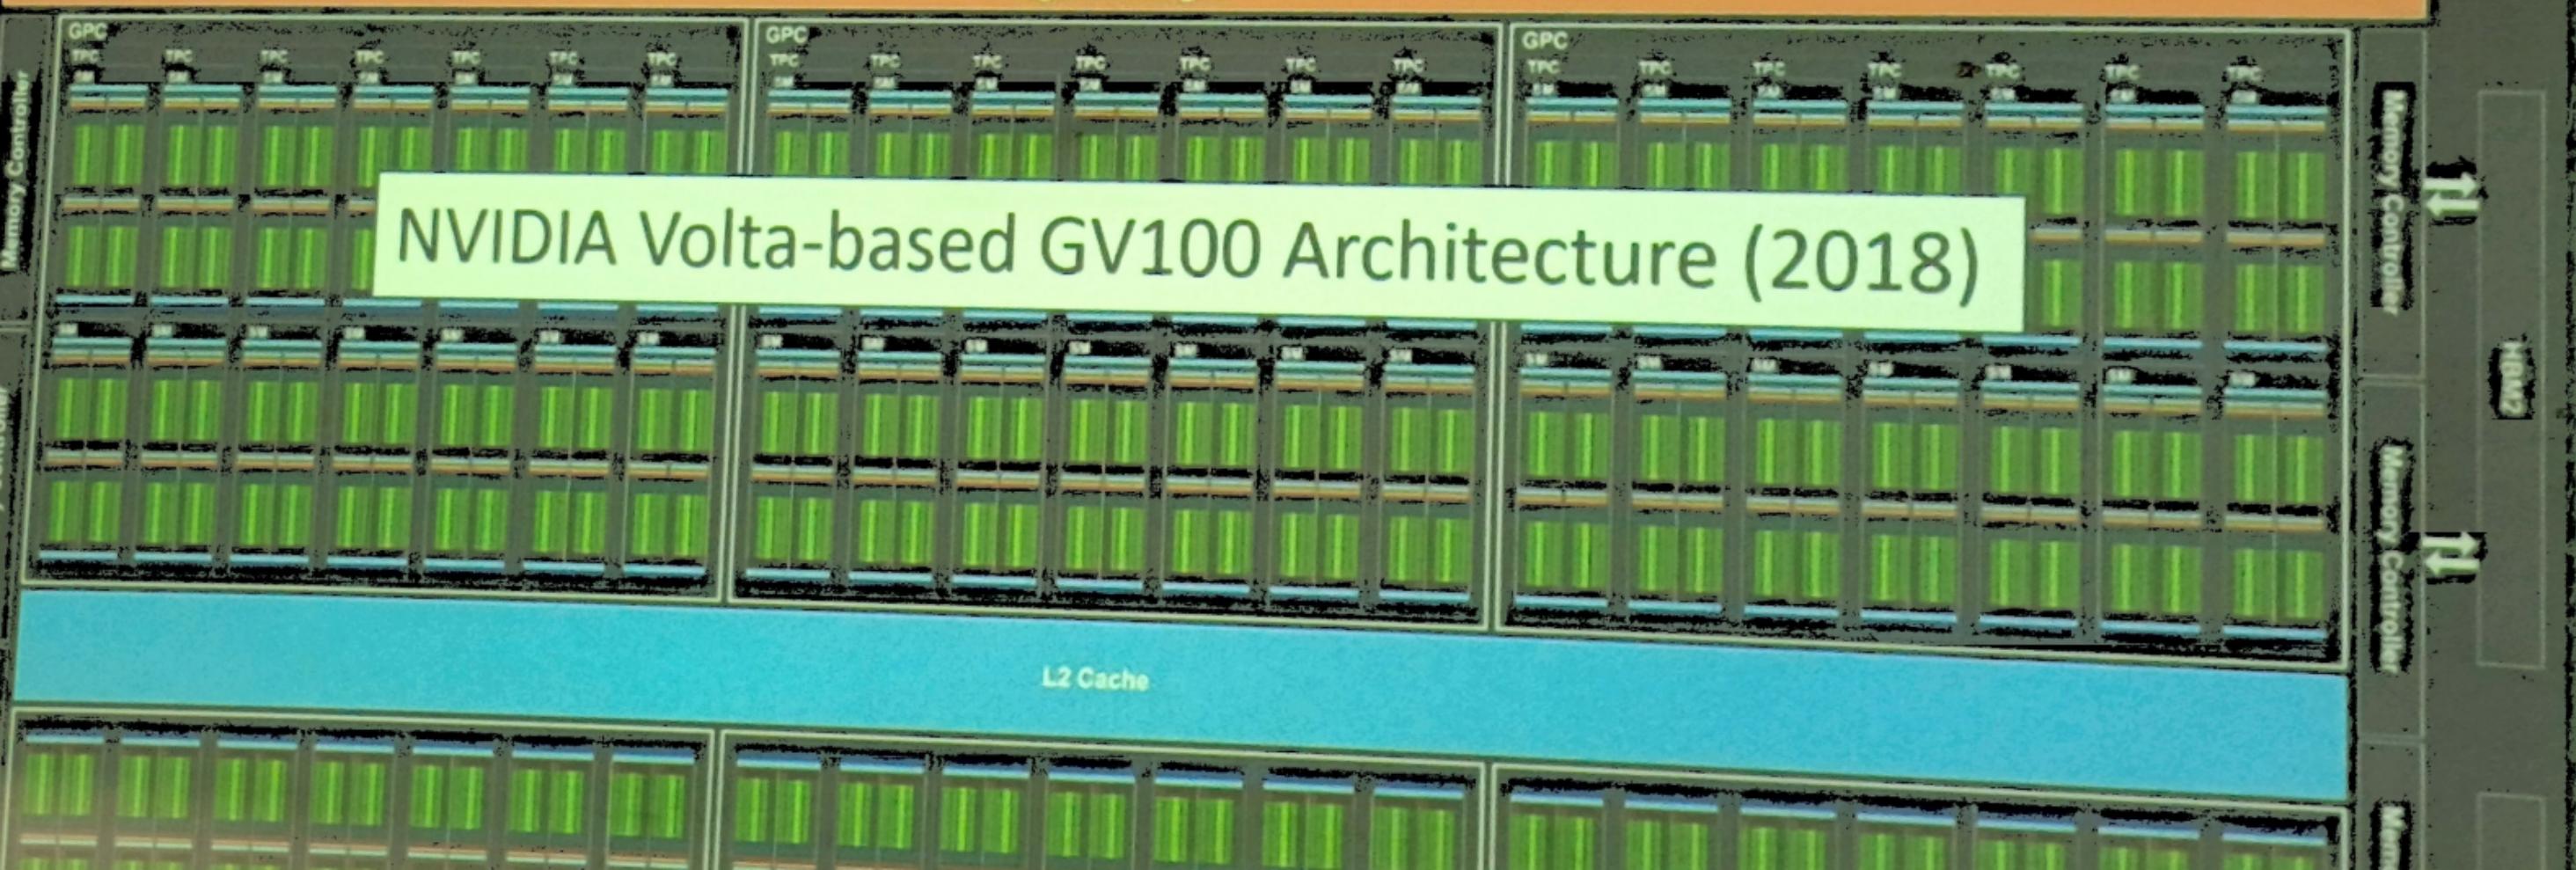


NVIDIA Volta-based GV100 Architecture (2018)

Many many cores,  
not a lot of cache/control

### MultiThread Engine

# NVIDIA Volta-based GV100 Architecture (2018)



L2 Cache

Many many cores,  
not a lot of cache/control

Memory Controller

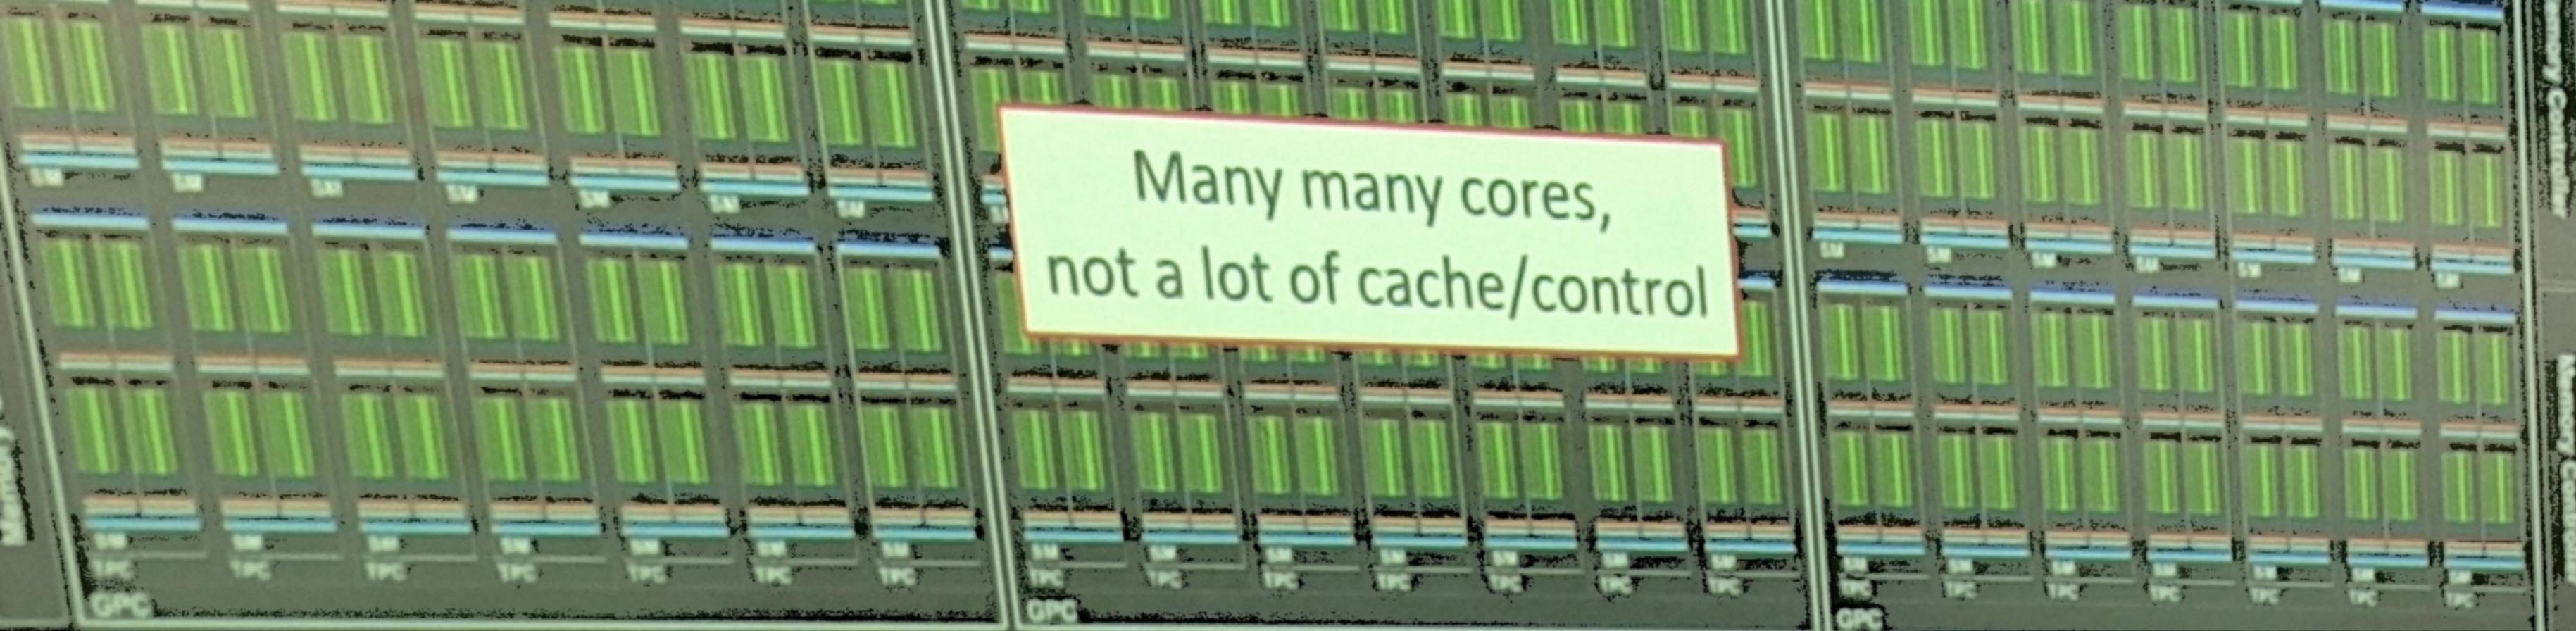
GPC

High-Speed Hub

Memory Controller

Memory Controller

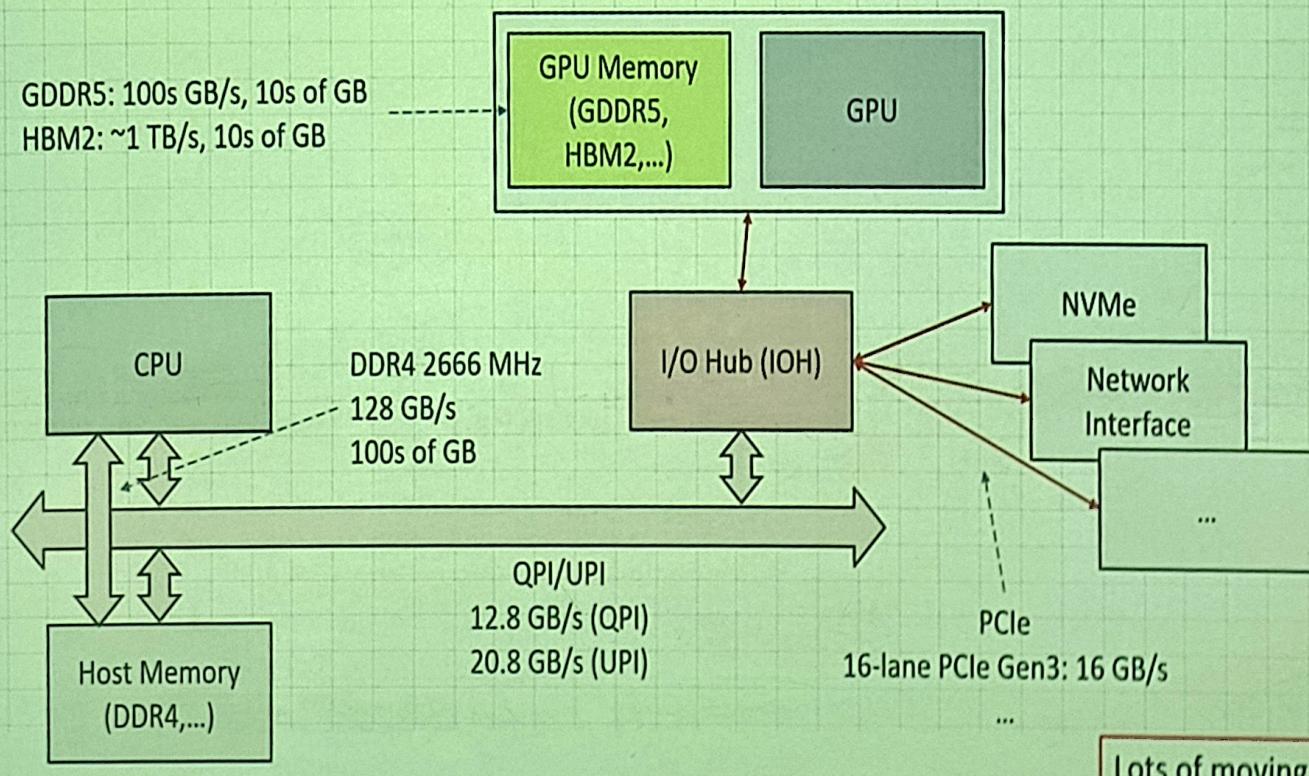
Memory Controller



Many many cores,  
not a lot of cache/control



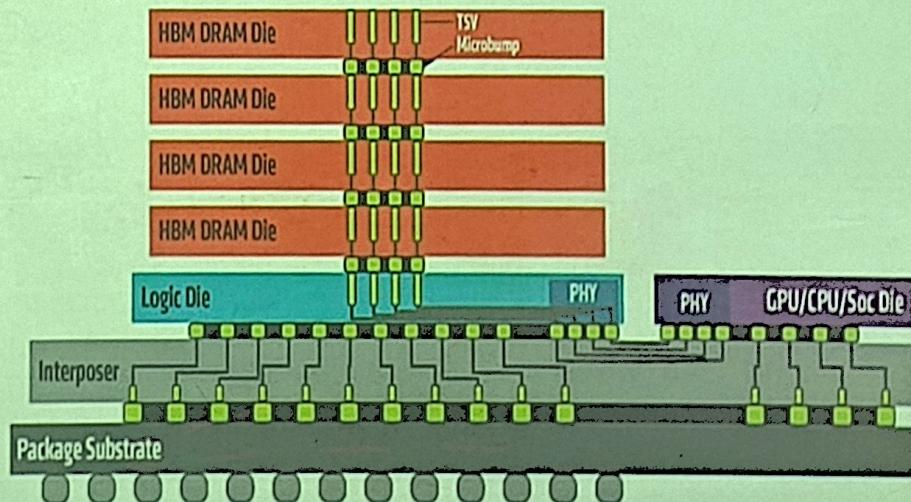
# System Architecture Snapshot With a GPU





# High-Performance Graphics Memory

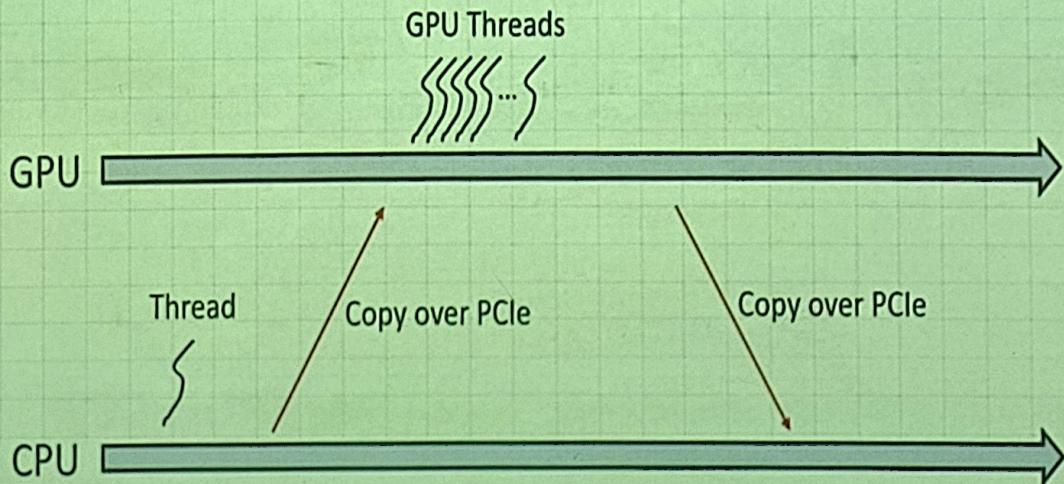
- Modern GPUs even employing 3D-stacked memory via silicon interposer
  - Very wide bus, very high bandwidth
  - e.g., HBM2 in Volta





# Massively Parallel Architecture For Massively Parallel Workloads!

- NVIDIA CUDA (Compute Uniform Device Architecture) – 2007
  - A way to run custom programs on the massively parallel architecture!
- OpenCL specification released – 2008
- Both platforms expose synchronous execution of a massive number of threads



# Device

Multiprocessor N

⋮

Multiprocessor 2

Multiprocessor 1

Shared Memory

Registers

Registers

Registers

Instruction Unit

Processor 1

Processor 2

⋮

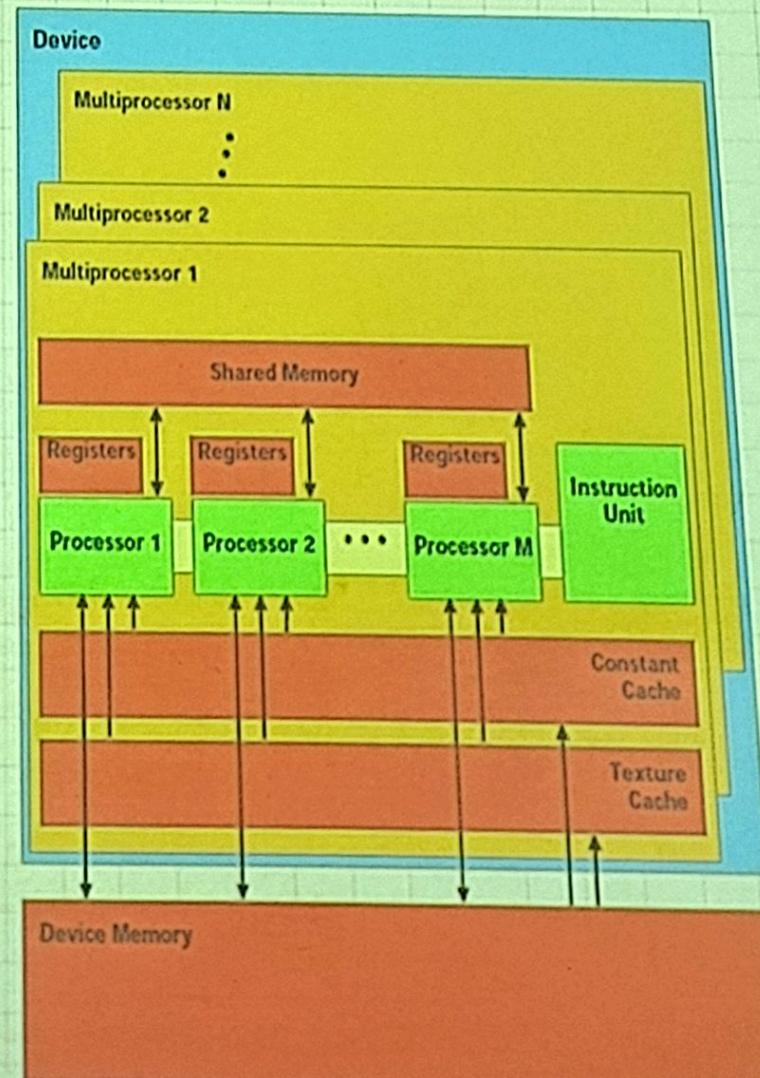
Processor M

Constant Cache

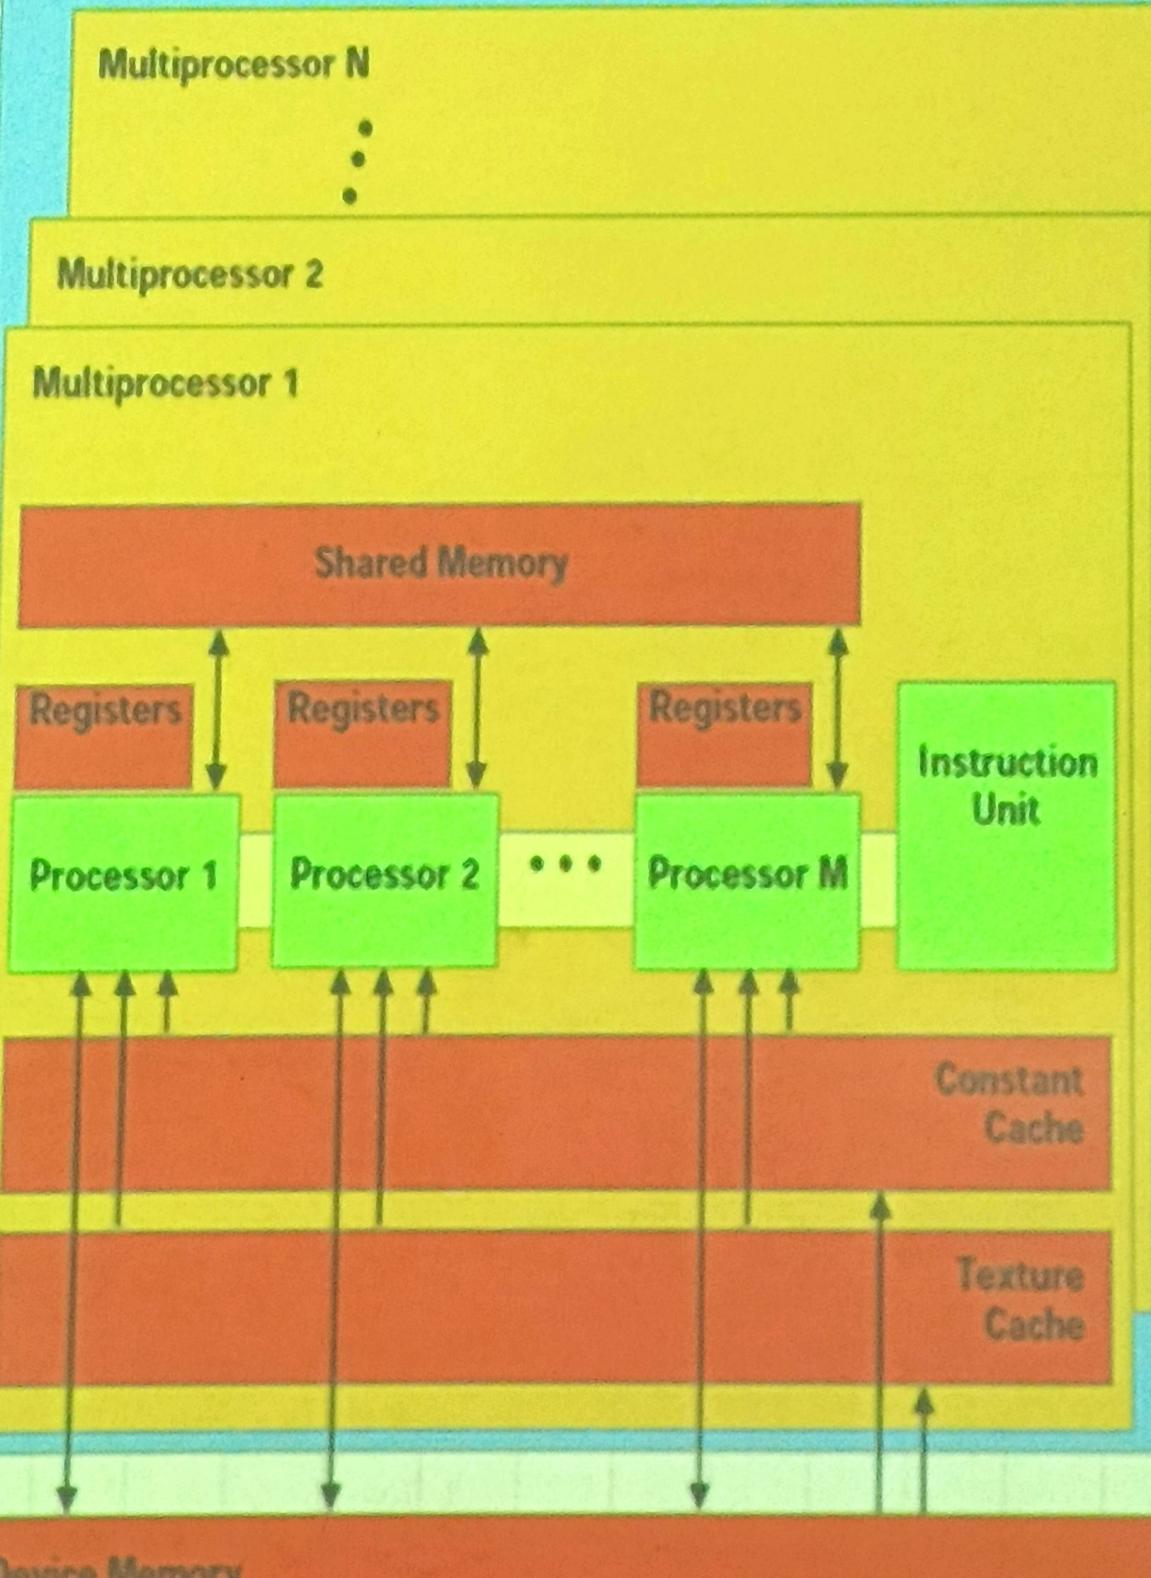
Texture Cache

Device Memory

# NVIDIA GPU Architecture



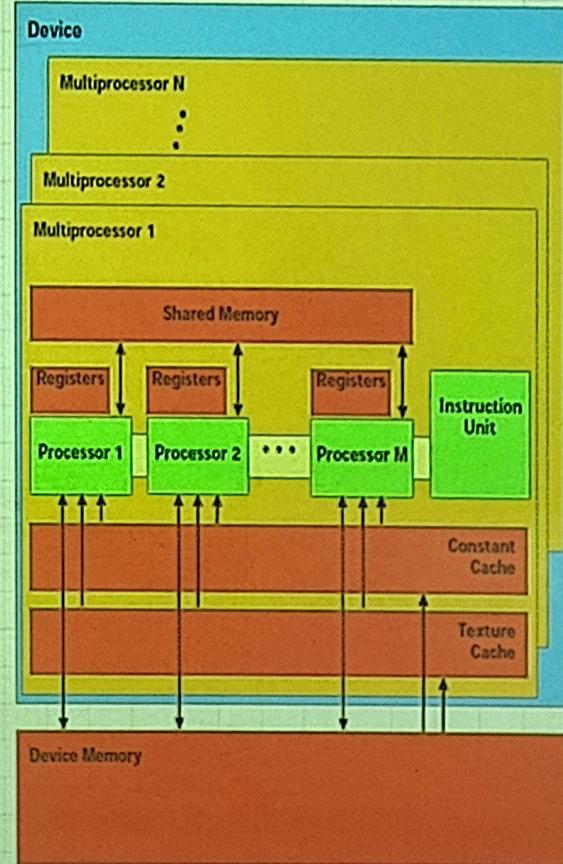
# Device





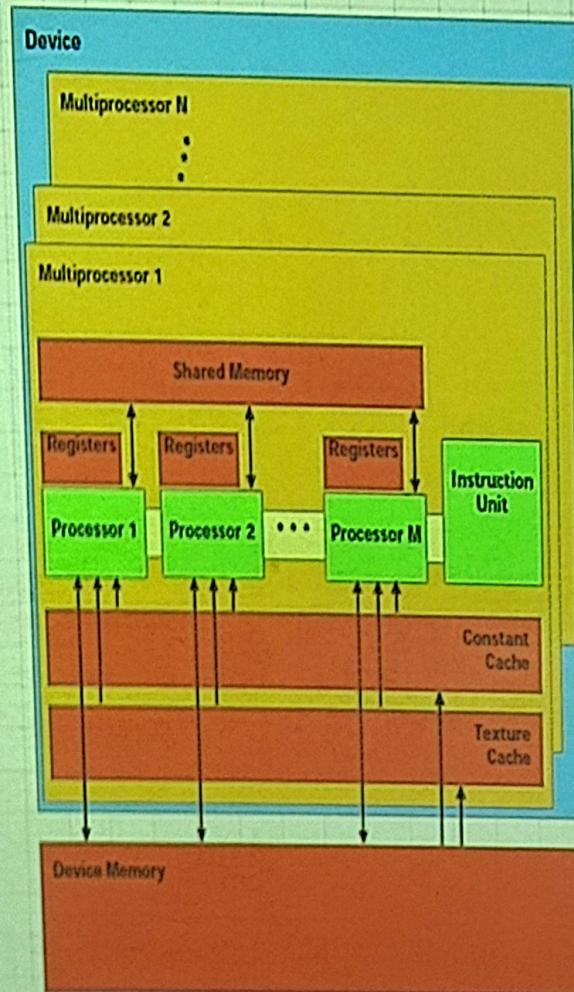
# NVIDIA GPU Architecture

- A scalable array of multithreaded Streaming Multiprocessors (SMs), each SM consists of
  - 8 Scalar Processor (SP) cores
  - 2 special function units for transcendentals



# NVIDIA GPU Architecture

- A scalable array of multithreaded Streaming Multiprocessors (SMs), each SM consists of
  - 8 Scalar Processor (SP) cores
  - 2 special function units for transcendentals
  - A multithreaded instruction unit
  - On-chip shared memory





# Volta Execution Architecture

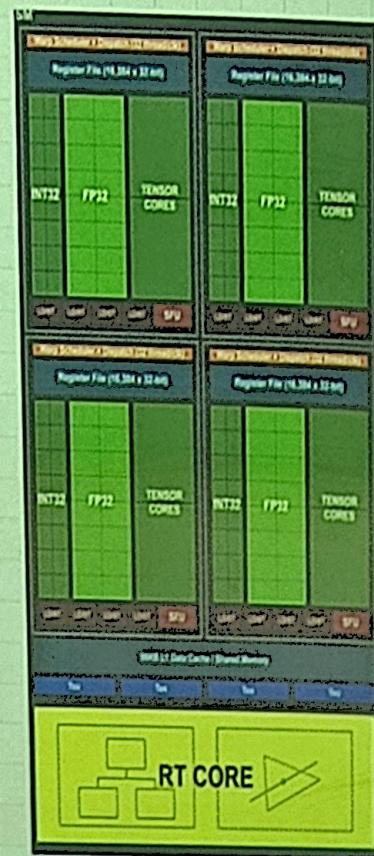
- 64 INT32 Cores, 64 FP32 Cores, 4 Tensor Cores, Ray-tracing cores..
  - Specialization to make use of chip space...?
- Not much on-chip memory per thread
  - 96 KB Shared memory
  - 1024 Registers per FP32 core
- Hard limit on compute management
  - 32 blocks AND 2048 threads AND 1024 threads/block
  - e.g., 2 blocks with 1024 threads, or 4 blocks with 512 threads
  - Enough registers/shared memory for all threads must be available (all context is resident during execution)



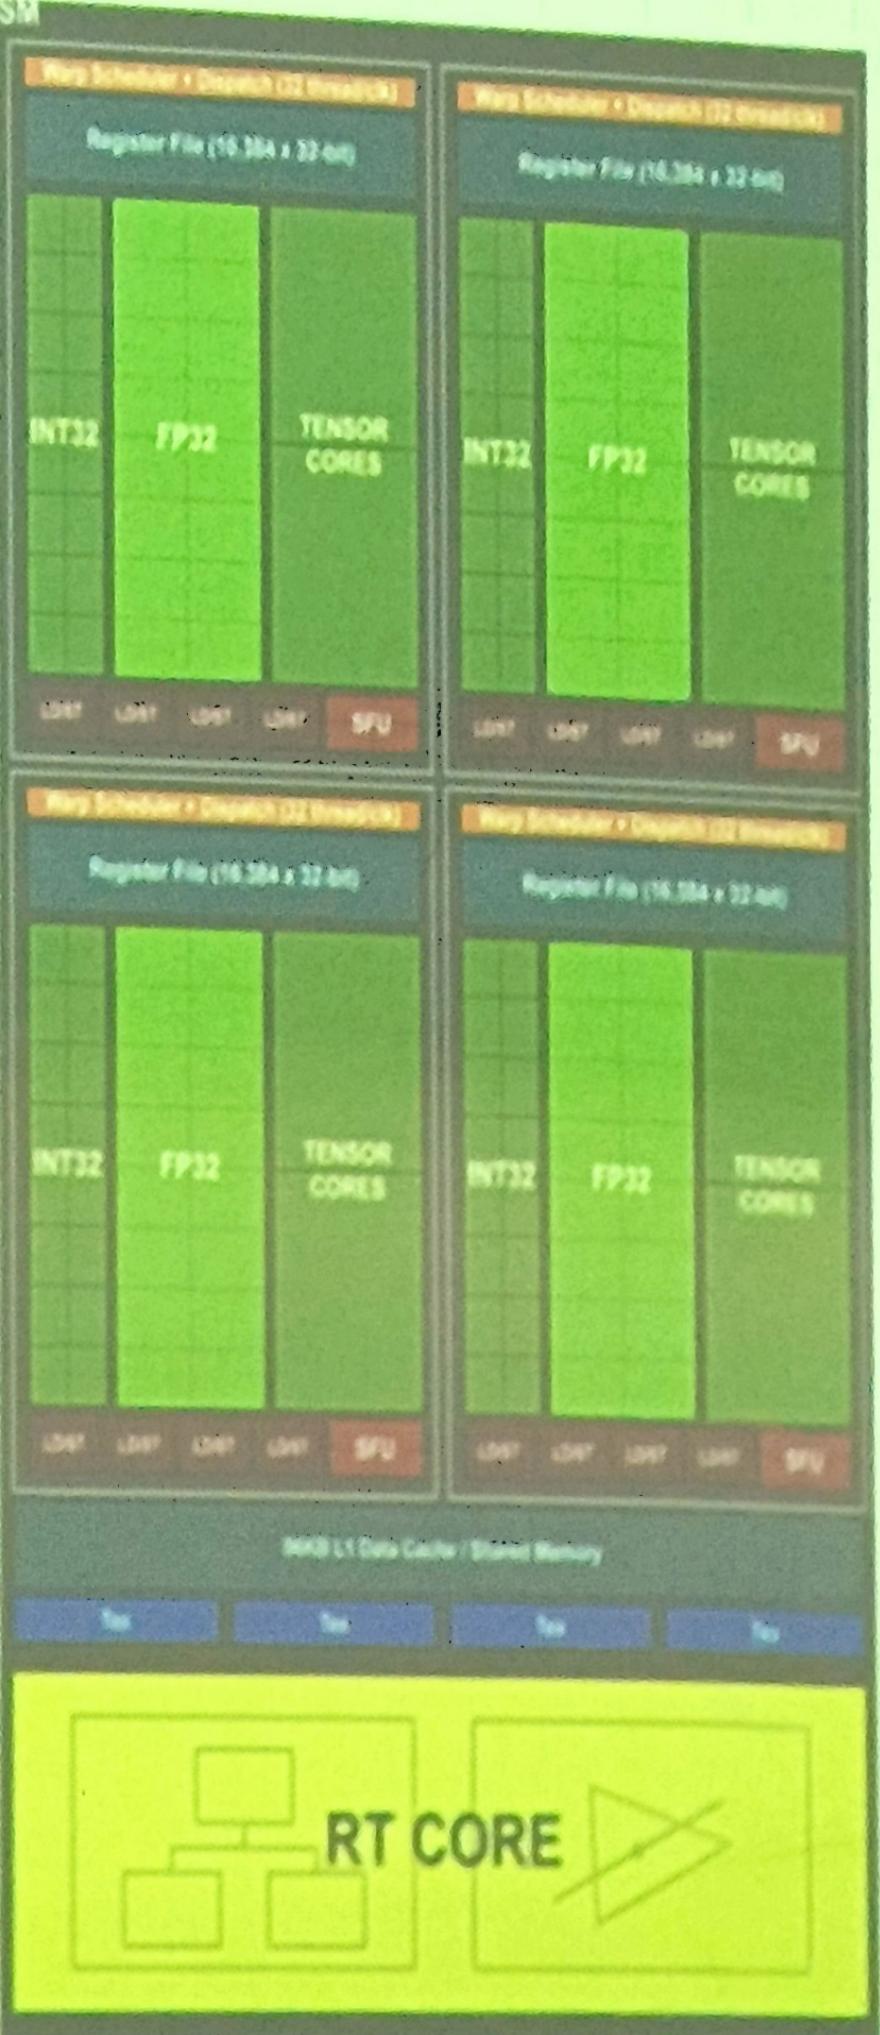
More threads than cores = Threads interleaved to hide memory latency

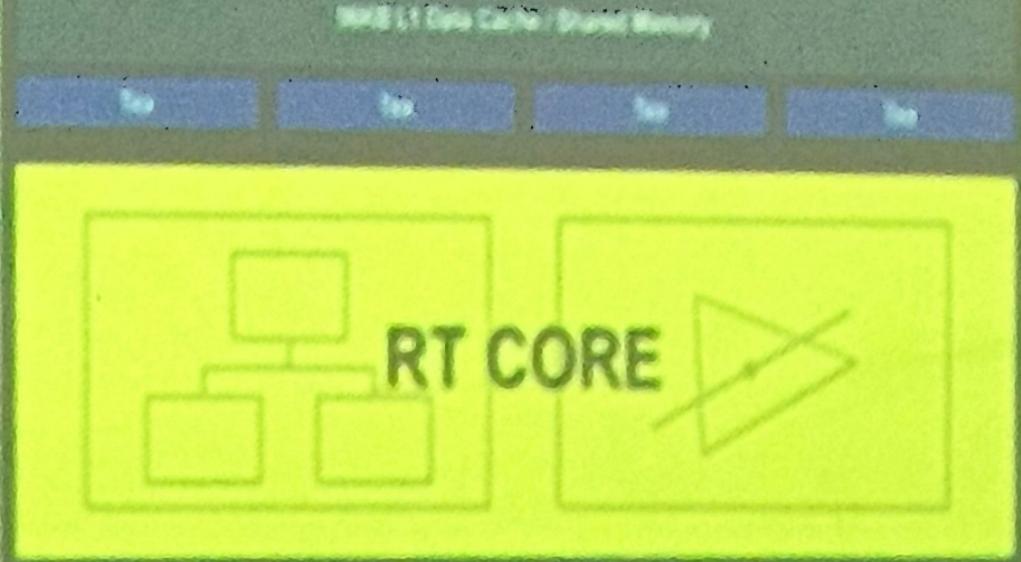
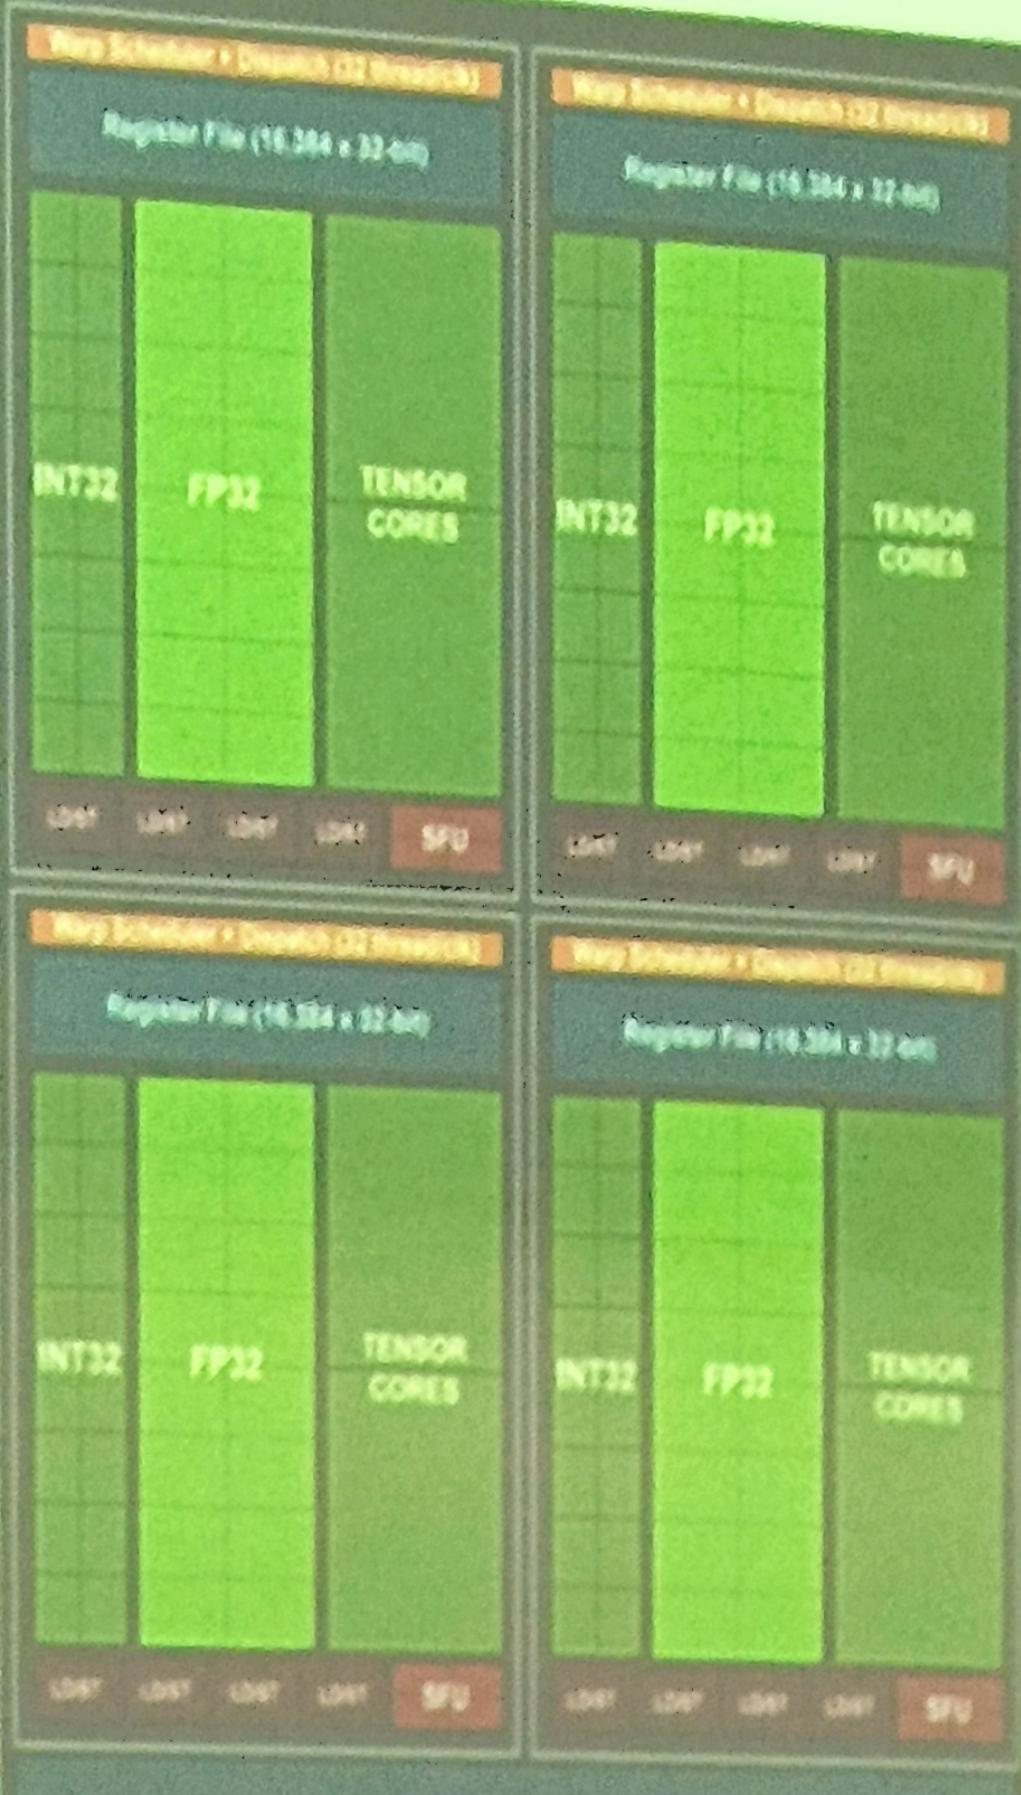
# Volta Execution Architecture

- 64 INT32 Cores, 64 FP32 Cores, 4 Tensor Cores, Ray-tracing cores..
  - Specialization to make use of chip space...?
- Not much on-chip memory per thread
  - 96 KB Shared memory
  - 1024 Registers per FP32 core
- Hard limit on compute management
  - 32 blocks AND 2048 threads AND 1024 threads/block
  - e.g., 2 blocks with 1024 threads, or 4 blocks with 512 threads
  - Enough registers/shared memory for all threads must be available (all context is resident during execution)



More threads than cores – Threads interleaved to hide memory latency







## Resource Balancing Details

- How many threads in a block?
- Too small:  $4 \times 4$  window == 16 threads
  - 128 blocks to fill 2048 thread/SM
  - SM only supports 32 blocks -> only 512 threads used
    - SM has only 64 cores... does it matter? Sometimes!
- Too large:  $32 \times 48$  window == 1536 threads
  - Threads do not fit in a block!
- Too large: 1024 threads using more than 64 registers
- Limitations vary across platforms (Fermi, Pascal, Volta, ...)



## *Resource Balancing Details*

- How many threads in a block?
- Too small:  $4 \times 4$  window == 16 threads
  - 128 blocks to fill 2048 thread/SM
  - SM only supports 32 blocks -> only 512 threads used
    - SM has only 64 cores... does it matter? Sometimes!
- Too large:  $32 \times 48$  window == 1536 threads
  - Threads do not fit in a block!
- Too large: 1024 threads using more than 64 registers
- Limitations vary across platforms (Fermi, Pascal, Volta, ...)

# Volta Execution Architecture

- 64 INT32 Cores, 64 FP32 Cores, 4 Tensor Cores, Ray-tracing cores..
  - Specialization to make use of chip space...?
- Not much on-chip memory per thread
  - 96 KB Shared memory
  - 1024 Registers per FP32 core
- Hard limit on compute management
  - 32 blocks AND 2048 threads AND 1024 threads/block
    - e.g., 2 blocks with 1024 threads, or 4 blocks with 512 threads
  - Enough registers/shared memory for all threads must be available (all context is resident during execution)

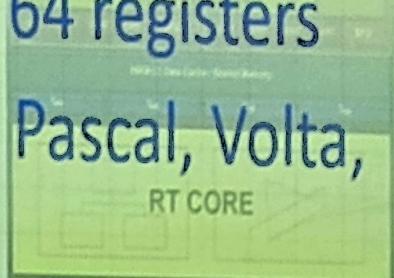


More threads than cores – Threads interleaved to hide memory latency

## Resource Balancing Details

- How many threads in a block?
- Too small: 4x4 window == 16 threads
  - 128 blocks to fill 2048 thread/SM
  - SM only supports 32 blocks -> only 512 threads used
    - SM has only 64 cores... does it matter? Sometimes!
- Too large: 32x48 window == 1536 threads
  - Threads do not fit in a block!
- Too large: 1024 threads using more than 64 registers  
available (all context is resident during execution)
- Limitations vary across platforms (Fermi, Pascal, Volta,...)

More threads than cores – Threads interleaved to hide memory latency

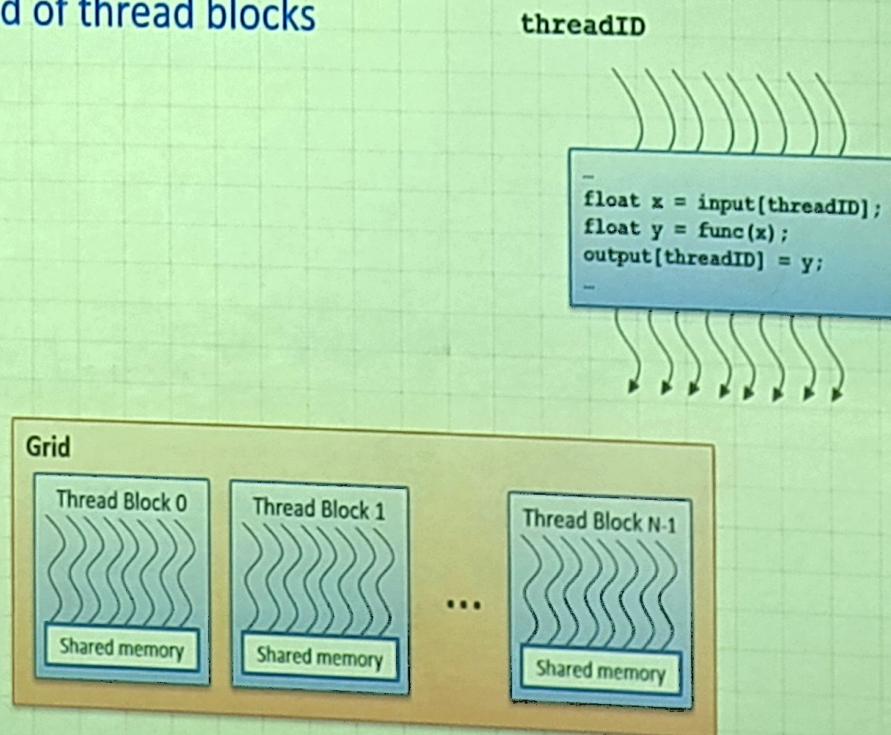


# Warp Scheduling Unit

- Threads in a block are executed in 32-thread “warp” unit
  - Not part of language specs, just architecture specifics
  - A warp is SIMD – Same PC, same instructions executed on every core
- What happens when there is a conditional statement?
- Too large: 32x48 window == 1536 threads
  - Prefix operations, or control divergence
  - Threads do not fit in a block!
  - More on this later!
- Too large: 1024 threads using more than 64 registers
- Warps have been 32-threads so far, but may change in the future
- Limitations vary across platforms (Fermi, Pascal, Volta, ...)

# CUDA Programming Model

- A CUDA kernel is executed by an array of threads
  - All threads run the same code (SPMD)
  - Each thread has an ID that it uses to compute memory addresses and make control decisions
- Threads are arranged as a grid of thread blocks
  - Threads within a block have access to a segment of shared memory

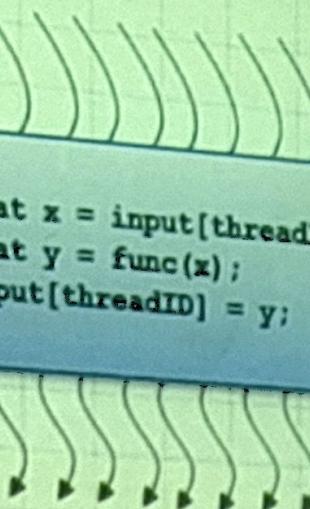


# of thread blocks

make control decisions

threadID

```
-  
float x = input[threadID];  
float y = func(x);  
output[threadID] = y;  
-
```

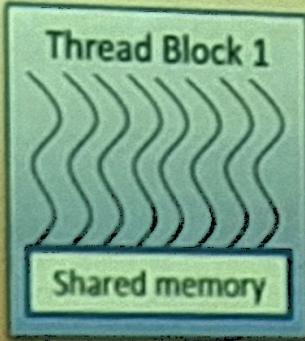


## Grid

Thread Block 0

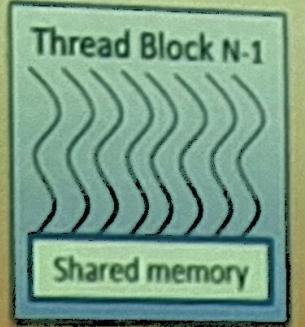


Thread Block 1



...

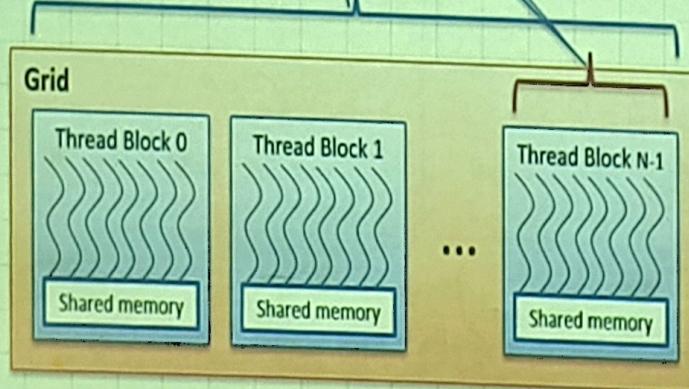
Thread Block N-1



# Kernel Invocation Syntax

grid & thread block dimensionality

```
vecAdd<<<32, 512>>>(devPtrA, devPtrB, devPtrC);
```



```
int i = blockIdx.x * blockDim.x + threadIdx.x;
```

block ID within a grid

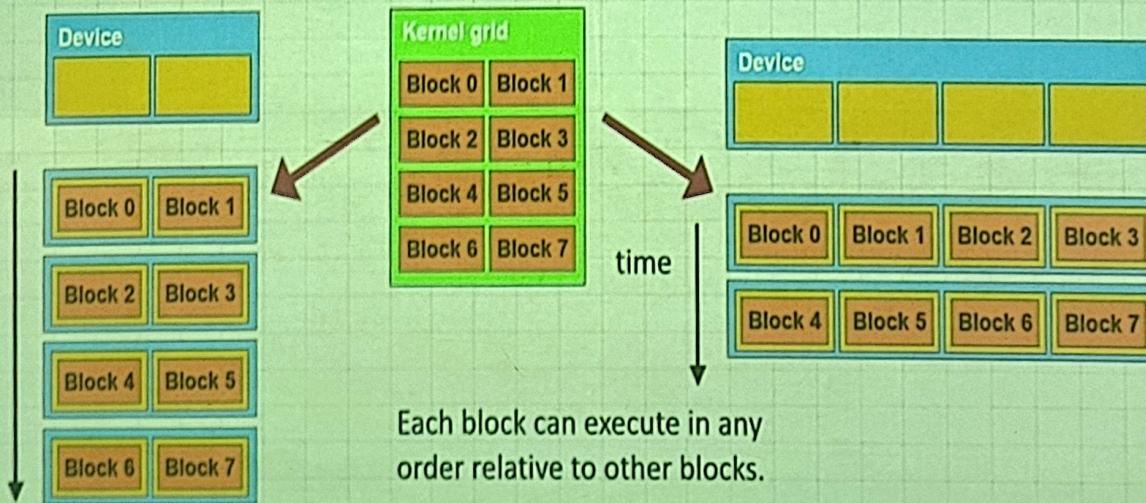
number of threads per block

thread ID within a thread block



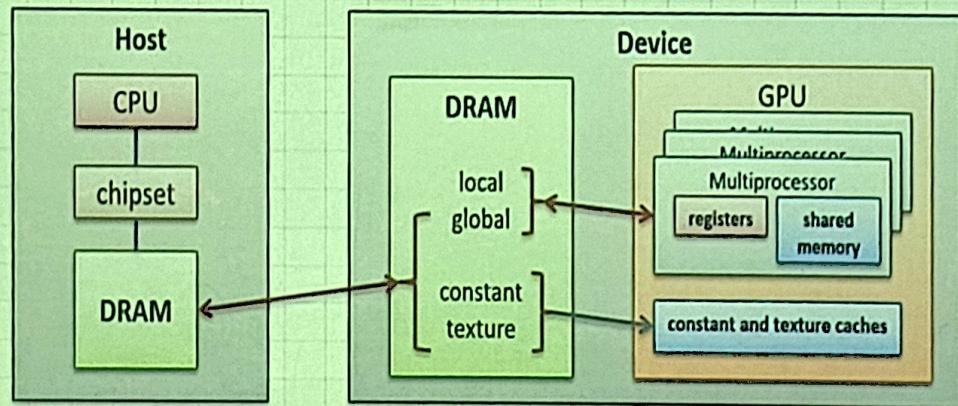
# Mapping Threads to the Hardware

- Blocks of threads are transparently assigned to SMs
  - A block of threads executes on one SM & does not migrate
  - Several blocks can reside concurrently on one SM



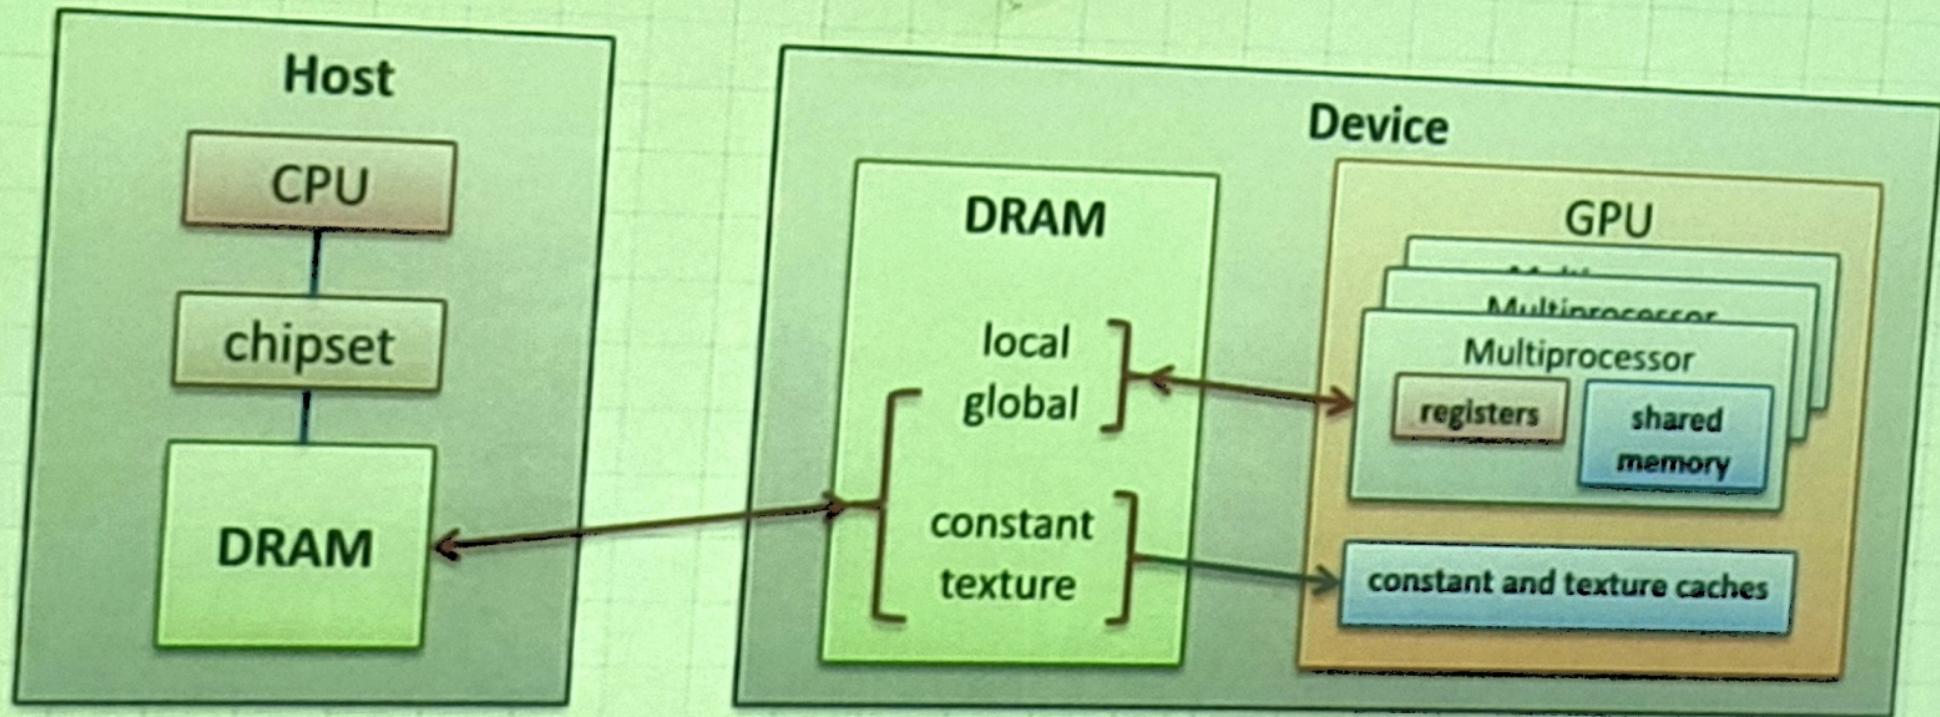


# GPU Memory Hierarchy



Memory	Location	Cached	Access	Scope	Lifetime
Register	On-chip	N/A	R/W	One thread	Thread
Local	Off-chip	No	R/W	One thread	Thread
Shared	On-chip	N/A	R/W	All threads in a block	Block
Global	Off-chip	No	R/W	All threads + host	Application
Constant	Off-chip	Yes	R	All threads + host	Application
Texture	Off-chip	Yes	R	All threads + host	Application

# 3. Memory Hierarchy



Memory	Location	Cached	Access	Scope	Lifetime
Register	On-chip	N/A	R/W	One thread	Thread
Local	Off-chip	No	R/W	One thread	Thread
Shared	On-chip	N/A	R/W	All threads in a block	Block
Global	Off-chip	No	R/W	All threads + host	Application
Constant	Off-chip	Yes	R	All threads + host	Application
Texture	Off-chip	Yes	R	All threads + host	Application

## CPU side

```
int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>(A, B, C);
    ...
}
```

## Asynchronous call

## GPU side

```
// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
}
```

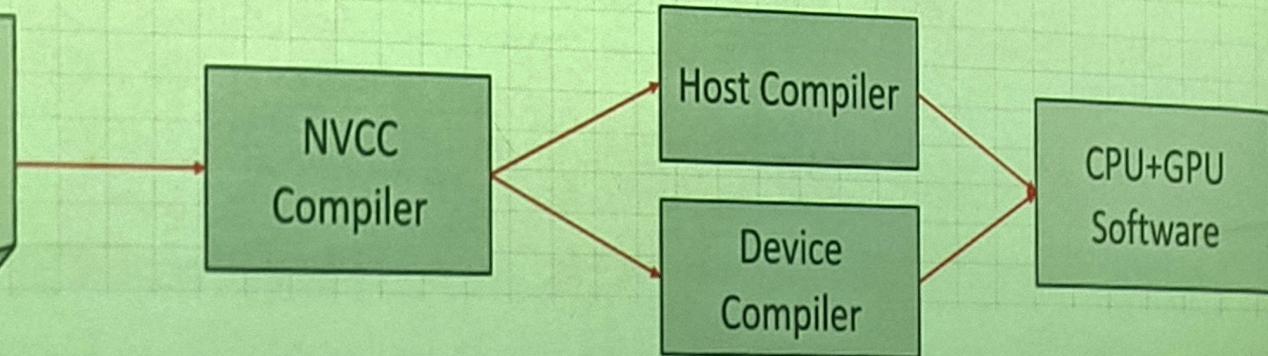
C/C++  
+ CUDA  
Code

NVCC  
Compiler

Host Compiler

Device  
Compiler

CPU+GPU  
Software





# Simple CUDA Example

```
int main()
{
    ...
    // Kernel invocation with N threads
    VecAdd<<<1, N>>>(A, B, C);
    Should wait for kernel to finish
}

// Kernel definition
__global__ void VecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x;
    C[i] = A[i] + B[i];
    Only void allowed
}
```

**1 block**

**N threads per block**

**global**: In GPU, called from host/GPU

**device**: In GPU, called from GPU

**host**: In host, called from host

**N instances of VecAdd spawned in GPU**

**One function can be both**

**Which of N threads am I?**  
See also: blockIdx



# CPU-Only Version

```
void vecAdd(int N, float* A, float* B, float* C) {  
    for (int i = 0; i < N; i++) C[i] = A[i] + B[i];  
}
```

```
int main(int argc, char **argv)  
{  
    int N = 16384; // default vector size  
  
    float *A = (float*)malloc(N * sizeof(float));  
    float *B = (float*)malloc(N * sizeof(float));  
    float *C = (float*)malloc(N * sizeof(float));  
  
    vecAdd(N, A, B, C); // call compute kernel  
  
    free(A); free(B); free(C);  
}
```

# *Adding GPU support*

```
int main(int argc, char **argv)
{
    int N = 16384; // default vector size

    float *A = (float*)malloc(N * sizeof(float));
    float *B = (float*)malloc(N * sizeof(float));
    float *C = (float*)malloc(N * sizeof(float));

    float *devPtrA, *devPtrB, *devPtrC;

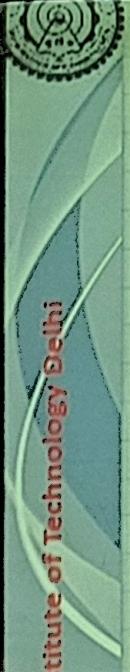
    cudaMalloc((void**)&devPtrA, N * sizeof(float));
    cudaMalloc((void**)&devPtrB, N * sizeof(float));
    cudaMalloc((void**)&devPtrC, N * sizeof(float));

    cudaMemcpy(devPtrA, A, N * sizeof(float), cudaMemcpyHostToDevice);
    cudaMemcpy(devPtrB, B, N * sizeof(float), cudaMemcpyHostToDevice);
```



# Adding GPU support

```
vecAdd<<<N/512, 512>>>(devPtrA, devPtrB, devPtrC);  
  
cudaMemcpy(C, devPtrC, N * sizeof(float), cudaMemcpyDeviceToHost);  
  
cudaFree(devPtrA);  
cudaFree(devPtrB);  
cudaFree(devPtrC);  
  
free(A); free(B); free(C);  
}
```



# GPU Kernel

- CPU version

```
vecAdd<<<N/512>>>(devPtrA, devPtrB, devPtrC) Kernel invocation
void vecAdd(int N, float* A, float* B, float* C)
```

```
{  
    cudaMemcpy(C, devPtrC, N * sizeof(float), cudaMemcpyDeviceToHost);  
    for (int i = 0; i < N; i++)  
        C[i] = A[i] + B[i];  
    cudaFree(devPtrA);  
    cudaFree(devPtrB);  
    cudaFree(devPtrC);  
}
```

Copy results from  
device memory to  
the host memory

free(A); free(B); free(C);

```
__global__ void vecAdd(float* A, float* B, float* C)  
{  
    int i = blockIdx.x * blockDim.x + threadIdx.x;  
    C[i] = A[i] + B[i];  
}
```

Device memory de-  
allocation

```
int main(int argc, char* argv[])
{
    int N = 1024;

    struct timeval t1, t2, ta, tb;
    long msec1, msec2;
    float flop, mflop, gflop;

    float *a = (float *)malloc(N*N*sizeof(float));
    float *b = (float *)malloc(N*N*sizeof(float));
    float *c = (float *)malloc(N*N*sizeof(float));

    minit(a, b, c, N);

    gettimeofday(&t1, NULL);
    mmult(a, b, c, N); // a = b * c
    gettimeofday(&t2, NULL);

    mprint(a, N, 5);

    free(a);
    free(b);
    free(c);

    msec1 = t1.tv_sec * 1000000 + t1.tv_usec;
    msec2 = t2.tv_sec * 1000000 + t2.tv_usec;
    msec2 -= msec1;
    flop = N*N*N*2.0f;
    mflop = flop / msec2;
    gflop = mflop / 1000.0f;
    printf("msec = %10ld  GFLOPS = %.3f\n", msec2, gflop);
}
```

```
// a = b * c
void mmult(float *a, float *b, float *c, int N)
{
    for (int j = 0; j < N; j++)
        for (int k = 0; k < N; k++)
            for (int i = 0; i < N; i++)
                a[i+j*N] += b[i+k*N]*c[k+j*N];
}
```

```
void minit(float *a, float *b, float *c, int N)
{
    for (int j = 0; j < N; j++)
        for (int i = 0; i < N; i++) {
            a[i+N*j] = 0.0f;
            b[i+N*j] = 1.0f;
            c[i+N*j] = 1.0f;
        }
}
```

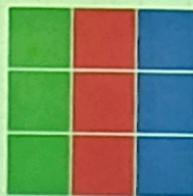
```
void mprint(float *a, int N, int M)
{
    int i, j;

    for (int j = 0; j < M; j++)
    {
        for (int i = 0; i < M; i++)
            printf("%.2f ", a[i+N*j]);
        printf("...\n");
    }
    printf("..\n");
}
```

# *Matrix Representation in Memory*

```
for (i = 0; i < n; ++i)
    for (j = 0; j < n; ++j)
        for (k = 0; k < n; ++k)
            a[i+n*j] += b[i+n*k] * c[k+n*j];
```

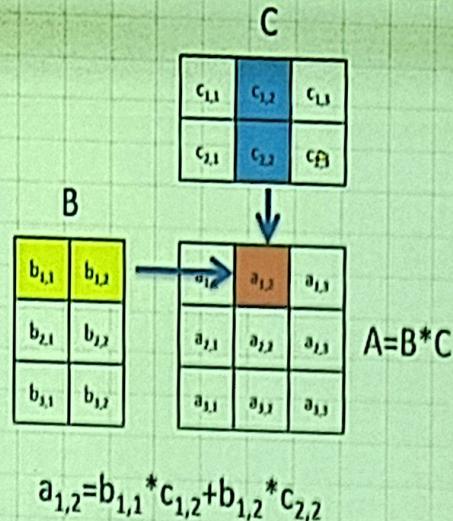
- Matrices are stored in column-major order



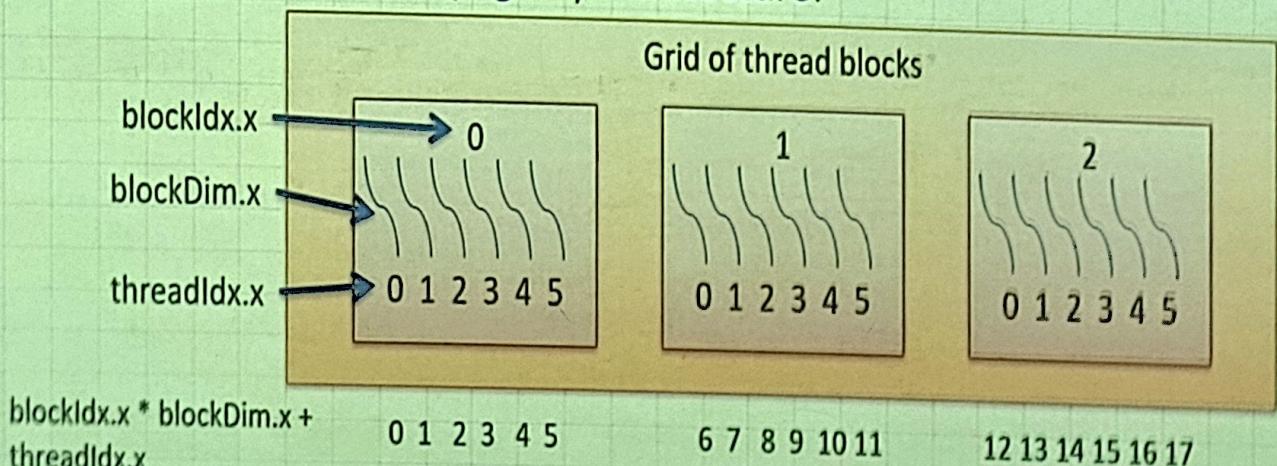
- For reference, jki-ordered version runs at 1.7 GFLOPS on 3 GHz Intel Xeon (single core)

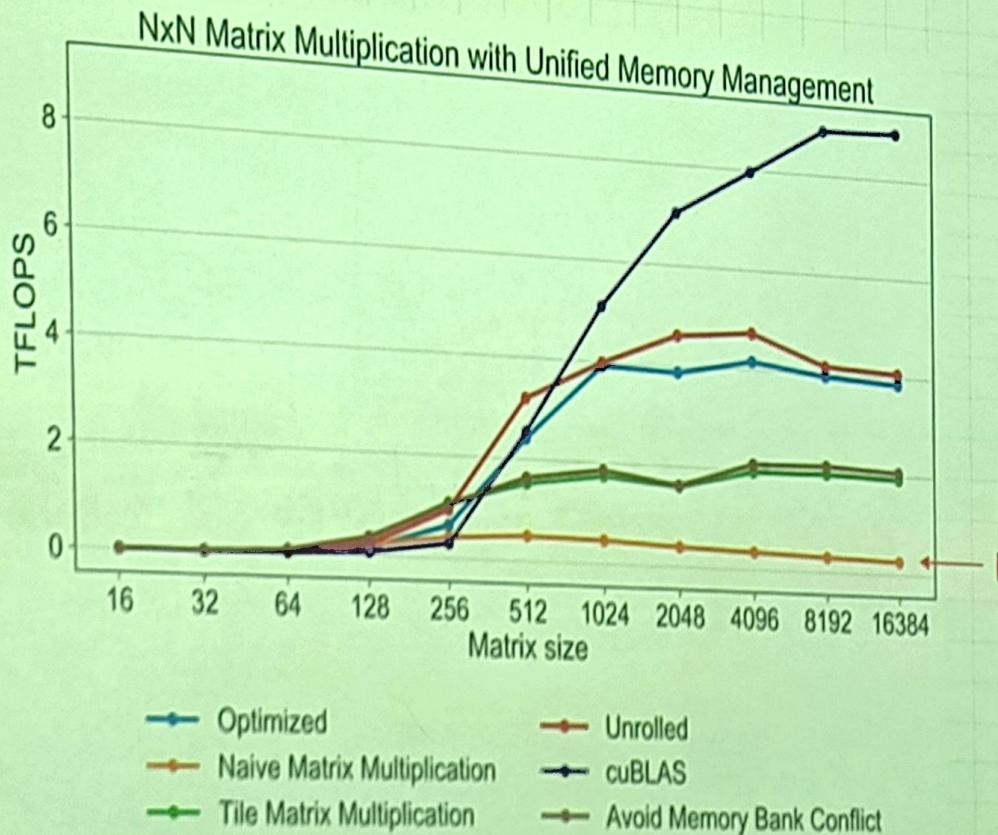
Map this code:

```
for (i = 0; i < n; ++i)
    for (j = 0; j < n; ++j)
        for (k = 0; k < n; ++k)
            a[i+n*j] += b[i+n*k] * c[k+n*j];
```



into this (logical) architecture:





Results from NVIDIA P100



# Kernel

```
void mmult(float *a, float *b, float *c, int N)
{
    for(int i = 0; i < N; i++)
        for(int j = 0; j < N; j++)
            for (int k = 0; k < N; k++)
                a[i+j*N] += b[i+k*N]*c[k+j*N];
}
```

```
__global__
void mmult(float *a, float *b, float *c, int N)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y;
    float sum = 0.0;

    for (int k = 0; k < N; k++)
        sum += b[i+N*k] * c[k+N*j];

    a[i+N*j] = sum;
}
```

```
dim3 dimGrid(32, 1024);
dim3 dimBlock(32);
mmult<<<dimGrid, dimBlock>>>(devPtrA, devPtrB, devPtrC, N);
```