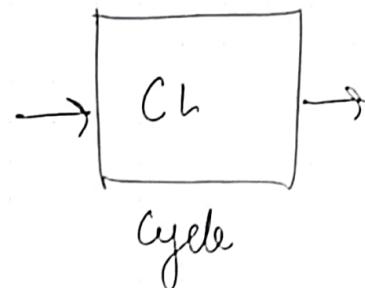




What about frequency?

- What is frequency dependent on ...
- Frequency = $1 / \text{clock period}$
- Clock Period:
 - 1 pipeline stage is expected to take 1 clock cycle
 - Clock period = maximum latency of the pipeline stages
- How to reduce the clock period?
 - Make each stage of the pipeline smaller by increasing the number of pipeline stages
 - Use faster transistors





Limits to increasing frequency - II

- What does it mean to have a very high frequency?
- Before answering, keep these facts in mind:

1

Thumb
Rule

$$P \propto f^3$$

P → power
f → frequency

2

Thermo-
dynamics

$$\Delta T \propto P$$

T → Temperature

3

We need to increase the number of pipeline stages →
more hazards, more forwarding paths



Pipeline Stages vs IPC

- $CPI = \frac{1}{IPC}$

$$CPI = CPI_{ideal} + \text{stall_rate} * \text{stall_penalty}$$

- The stall rate will remain more or less constant per instruction with the number of pipeline stages
- The **stall penalty** (in terms of cycles) will however **increase**
- This will lead to a net increase in CPI and **loss** in IPC

As we increase the number of stages, the IPC goes down.



What is ILP = Instruction level parallelism

- *multiple operations (or instructions) can be executed in parallel, from a single instruction stream*
 - so we are not yet talking about MIMD, multiple instruction streams

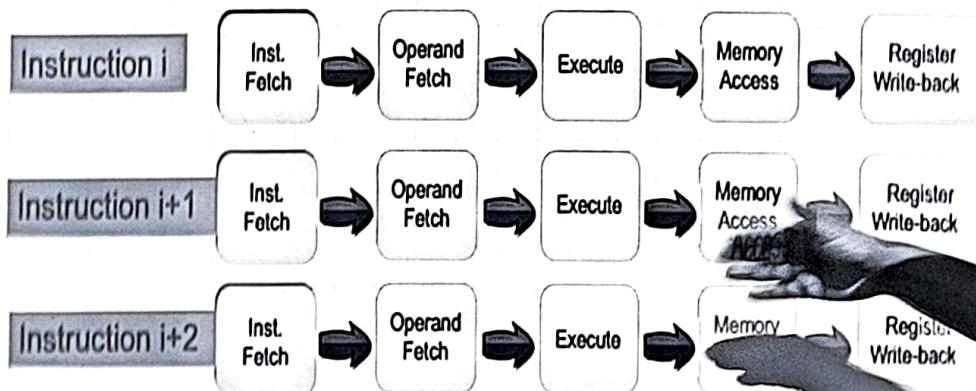
Needed:

- Sufficient (HW) resources
- Parallel scheduling
 - Hardware solution
 - Software solution
- Application should contain sufficient ILP



Since we cannot increase frequency ...

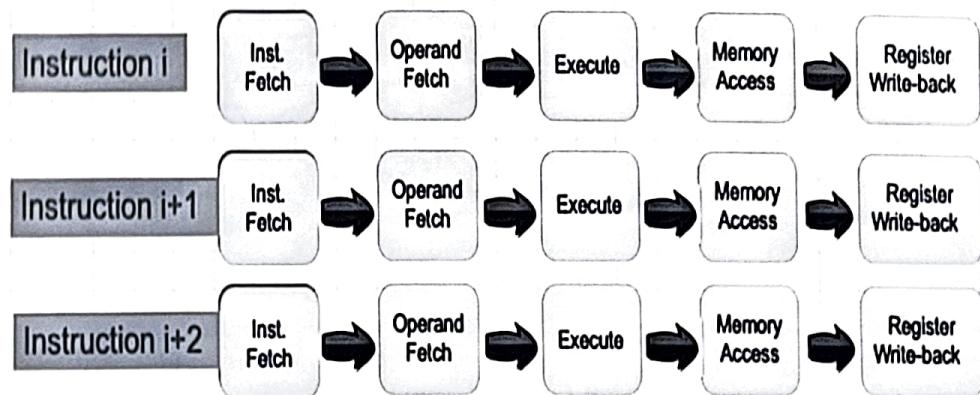
- Increase IPC
 - Issue more instructions per cycle
 - 2, 4, or 8 instructions
- Make it a **superscalar** processor → A processor that can execute multiple instructions per cycle
 - Have multiple in-order pipelines





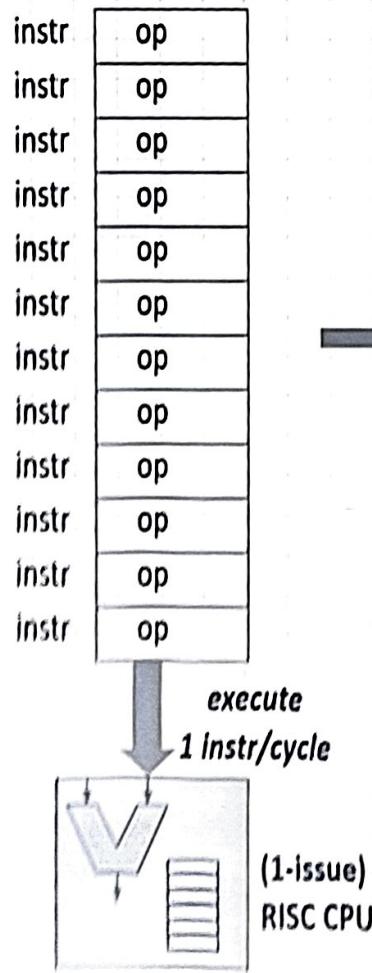
Since we cannot increase frequency ...

- Increase IPC
 - Issue more instructions per cycle
 - 2, 4, or 8 instructions
- Make it a **superscalar** processor → A processor that can execute multiple instructions per cycle
 - Have multiple in-order pipelines





Single Issue RISC vs Superscalar

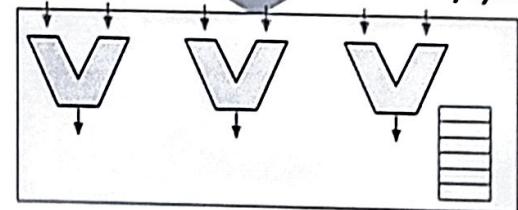


Change HW,

but can use
same code

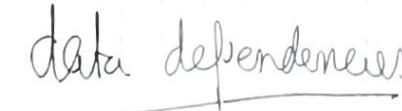
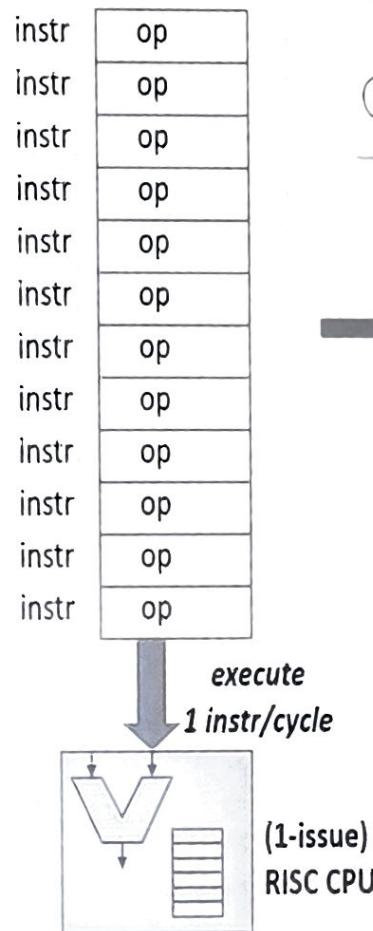
3-issue Superscalar

issue and (try to) execute
3 instr/cycle





Single Issue RISC vs Superscalar

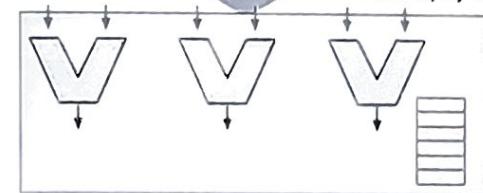


Change HW

but can use
same code

Control (f)
→ by

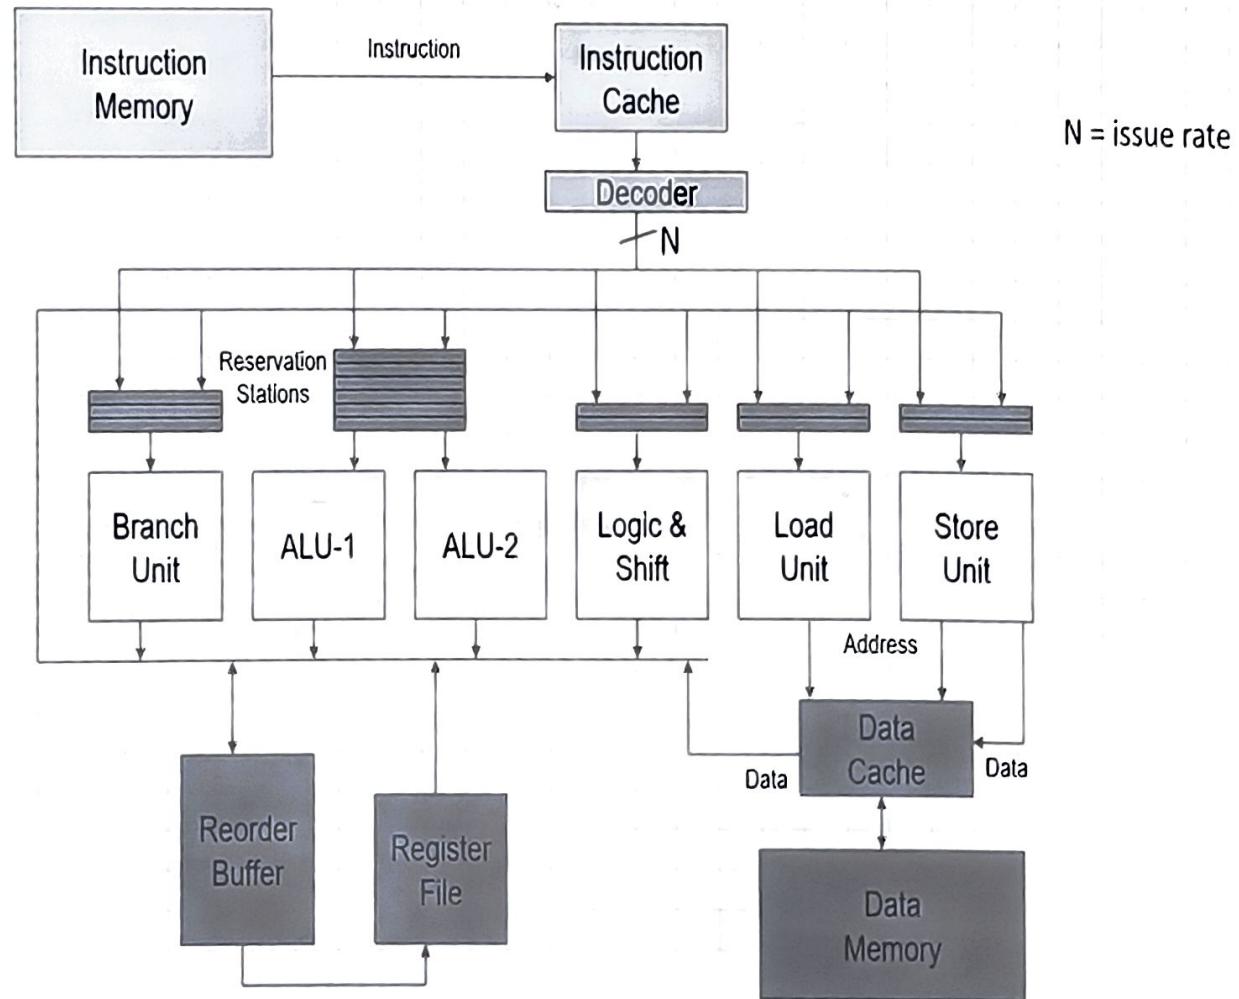
3-issue Superscalar



issue and (try to) execute
3 instr/cycle



Superscalar: General Architecture Concept





Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX				
SUB.D	F8, F2, F6		IF	EX				
DIV.D	F10, F0, F6			IF				
ADD.D	F6, F8, F2			IF				
MUL.D	F12, F2, F4							





Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 Id/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP Id/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX	EX	EX	EX	
SUB.D	F8, F2, F6		IF	EX	EX	WB		
DIV.D	F10, F0, F6			IF				
ADD.D	F6, F8, F2			IF		EX	EX	
MUL.D	F12, F2, F4					IF		

cannot execute
structural hazard



Example of Superscalar Processor Execution

- Superscalar processor organization:
 - simple pipeline: IF, EX, WB
 - fetches 2 instructions each cycle
 - 2 ld/st units, dual-ported memory; 2 FP adders; 1 FP multiplier
 - Instruction window (buffer between IF and EX stage) is of size 2
 - FP ld/st takes 1 cc; FP +/- takes 2 cc; FP * takes 4 cc; FP / takes 8 cc

Cycle		1	2	3	4	5	6	7
L.D	F6, 32 (R2)	IF	EX	WB				
L.D	F2, 48 (R3)	IF	EX	WB				
MUL.D	F0, F2, F4		IF	EX	EX	EX	EX	WB
SUB.D	F8, F2, F6		IF	EX	EX	WB		
DIV.D	F10, F0, F6			IF				EX
ADD.D	F6, F8, F2			IF		EX	EX	WB
MUL.D	F12, F2, F4				IF			?



Hazards

- Three types of hazards
 - Structural
 - Multiple instructions need access to the same hardware at the same time
 - Data dependence
 - There is a dependence between operands (in register or memory) of successive instructions
 - Control dependence
 - Determines the order of the execution of basic blocks
 - When jumping/branching to another address the pipeline has to be (partly) squashed and refilled



Data dependences

4

- **RaW** read after write
 - real or flow dependence
 - can only be avoided by value prediction (i.e. speculating on the outcome of a previous operation)
- **WaR** write after read
- **WaW** write after write
 - WaR and WaW are false or name dependencies
 - Could be avoided by **renaming** (if sufficient registers are available); see later slide



Data dependences

- **RaW** read after write
 - real or flow dependence
 - can only be avoided by value prediction (i.e. speculating on the outcome of a previous operation)
- **WaR** write after read
- **WaW** write after write
 - WaR and WaW are false or name dependencies
 - Could be avoided by renaming (if sufficient registers are available); see later slide

$$\begin{array}{l} c = a + b \\ d = a \cdot c + e \end{array}$$

Notes:

1. data dependences can be **both** between **registers** and **memory** data operations
2. data dependencies are shown in de DDG: Data Dependence Graph

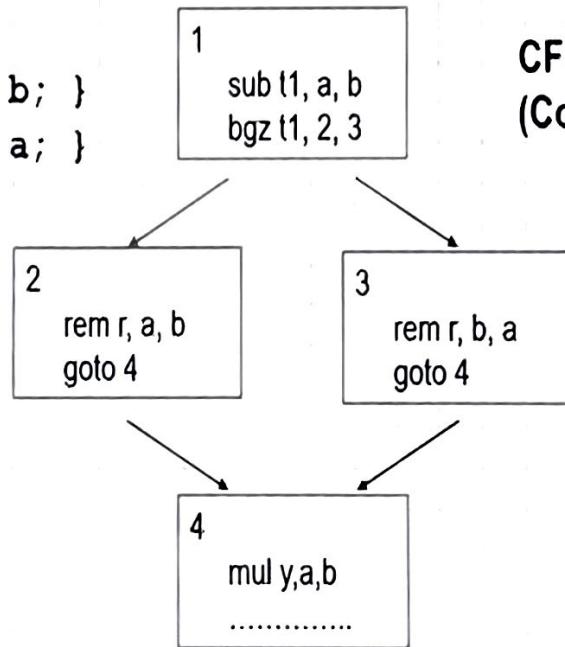


Control Dependences: CFG

C input code:

```
if (a > b) { r = a % b; }
else      { r = b % a; }
y = a*b;
```

CFG
(Control Flow Graph):



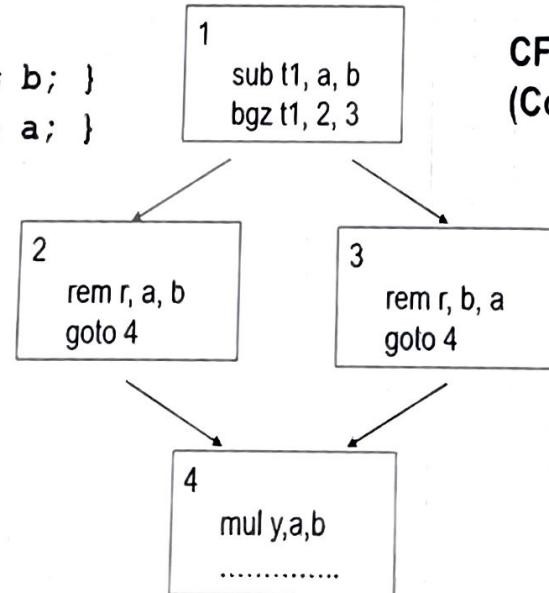


Control Dependences: CFG

C input code:

```
if (a > b) { r = a % b; }
else      { r = b % a; }
y = a*b;
```

CFG
(Control Flow Graph):



Questions:

- How real are control dependences?
- Can 'mul y, a, b' be moved to block 2, 3 or even block 1?
- Can 'rem r, a, b' be moved to block 1 and executed speculatively?



Avoiding pipeline stalls due to Hazards

- **Structural**
 - Buy more hardware
 - Extra units, pipelined units, more ports on RF and data memory (or banked memories), etc.
 - Note: more HW means bigger chip => could increase cycle time t_{cycle}
- **Data dependence**
 - Real (RaW) dependences: add Forwarding (aka Bypassing) logic
 - Compiler optimizations
 - False (WaR & WaW) dependences: use renaming (either in HW or in SW)



Avoiding pipeline stalls due to Hazards

- **Structural**

- Buy more hardware
 - Extra units, pipelined units, more ports on RF and data memory (or banked memories), etc.
 - Note: more HW means bigger chip => could increase cycle time t_{cycle}

- **Data dependence**

- Real (RaW) dependences: add Forwarding (aka Bypassing) logic
 - Compiler optimizations
 - False (WaR & WaW) dependences: use renaming (either in HW or in SW)

- **Control dependence**

- Adding extra pipeline HW to reduce the number of Branch delay slots
 - Branch prediction
 - Avoiding Branches

FP Loop: Where are the Hazards?

```
for (i=1000; i>0; i=i-1)  
    x[i] = x[i] + s;
```

Instruction producing result	Instruction using result	Latency in cycles	stalls between in cycles
FP ALU op	Another FP ALU op	4	3
FP ALU op	Store double	3	2
Load double	FP ALU op	1	1
Load double	Store double	1	0
Integer op	Integer op	1	0

FP Loop: Where are the Hazards?

```
for (i=1000; i>0; i=i-1)  
    x[i] = x[i] + s;
```

Instruction producing result	Instruction using result	Latency in cycles	stalls between in cycles
FP ALU op	Another FP ALU op	4	3
FP ALU op	Store double	3	2
Load double	FP ALU op	1	1
Load double	Store double	1	0
Integer op	Integer op	1	0

```
Loop: L.D      F0,0(R1) ;F0=vector element  
      ADD.D    F4,F0,F2 ;add scalar from F2  
      S.D      0(R1),F4 ;store result  
      DADDUI   R1,R1,-8 ;decrement pointer 8B (DW)  
      BNEZ    R1,Loop ;branch R1!=zero
```

```
1 Loop: L.D      F0,0(R1) ;F0=vector element  
2          stall  
3          ADD.D    F4,F0,F2 ;add scalar in F2  
4          stall  
5          stall  
6          S.D      0(R1),F4 ;store result  
7          DADDUI   R1,R1,-8 ;decrement pointer 8B (DW)  
8          stall      ;assumes can't forward to branch  
9          BNEZ    R1,Loop ;branch R1!=zero
```

1 L.D F0,0(R1) ;F0=vector element
2 ADD.D F4,F0,F2;add scalar from F2
3 S.D 0(R1),F4;store result
4 DADDUI R1,R1,-8;decrement pointer 8B
5 BNEZ R1,Loop ;branch R1!=zero

1 Loop: L.D F0,0(R1)
2 DADDUI R1,R1,-8
3 ADD.D F4,F0,F2
4 stall
5 stall
6 S.D 8(R1),F4
7 BNEZ R1,Loop

<pre> Integer op 1 0 Integer op 1 0 </pre> <p>P: L.D F0,0(R1) ;F0=vector element ADD.D F4,F0,F2 ;add scalar from F2 S.D 0(R1),F4;store result DADDUI R1,R1,-8 ;decrement pointer 8B BNEZ R1,Loop ;branch R1!=zero</p>	<pre> 5 stall 6 S.D 0(R1),F4 ;store result 7 DADDUI R1,R1,-8 ;decrement pointer 8B (DW) 8 stall ;assumes can't forward to branch 9 BNEZ R1,Loop ;branch R1!=zero </pre>
---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

<pre> Loop: L.D F0,0(R1) DADDUI R1,R1,-8 ADD.D F4,F0,F2 stall stall S.D 8(R1),F4 BNEZ R1,Loop </pre>	<pre> 1 Loop:L.D F0,0(R1) 3 ADD.D F4,F0,F2 6 S.D 0(R1),F4 7 L.D F6,-8(R1) ;drop DSUBUI & BNEZ 9 ADD.D F8,F6,F2 12 S.D -8(R1),F8 ;drop DSUBUI & BNEZ 13 L.D F10,-16(R1) 15 ADD.D F12,F10,F2 18 S.D -16(R1),F12 ;drop DSUBUI & BNEZ 19 L.D F14,-24(R1) 21 ADD.D F16,F14,F2 24 S.D -24(R1),F16 25 DADDUI R1,R1,#-32 ;alter to 4*8 26 BNEZ R1,LOOP </pre>
------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

27 clock cycles, or 6.75 per iteration
(Assumes R1 is multiple of 4)

Double	Store double	1	1	5	stall
Integer op	Integer op	1	0	6	S.D 0(R1),F4 ;store result
				7	DADDUI R1,R1,-8 ;decrement pointer 8B (DW)
				8	stall ;assumes can't forward to branch
				9	BNEZ R1,Loop ;branch R1!=zero

Loop: L.D F0,0(R1) ;F0=vector element
 ADD.D F4,F0,F2 ;add scalar from F2
 S.D 0(R1),F4 ;store result
 DADDUI R1,R1,-8 ;decrement pointer 8B
 BNEZ R1,Loop ;branch R1!=zero

1	Loop: L.D	F0,0(R1)	1	Loop:L.D	F0,0(R1)
2	DADDUI	R1,R1,-8	2	L.D	F6,-8(R1)
3	ADD.D	F4,F0,F2	3	L.D	F10,-16(R1)
4	stall		4	L.D	F14,-24(R1)
5	stall		5	ADD.D	F4,F0,F2
6	S.D	8(R1),F4	6	ADD.D	F8,F6,F2
7	BNEZ	R1,Loop	7	ADD.D	F12,F10,F2

14 clock cycles, or 3.5 per iteration

27 clock cycles, or 6.75 per iteration
 (Assumes R1 is multiple of 4)

1	Loop:L.D	F0,0(R1)	1	Loop:L.D	F0,0(R1)																																																																														
3	ADD.D	F4,F0,F2	3	S.D	0(R1),F4	4	ADD.D	F8,F6,F2	5	L.D	F6,-8(R1)	6	S.D	-8(R1),F8	7	L.D	F10,-16(R1)	8	ADD.D	F12,F10,F2	9	S.D	-16(R1),F12	10	L.D	F14,-24(R1)	11	ADD.D	F16,F14,F2	12	S.D	-24(R1),F16	13	DADDUI	R1,R1,#-32	14	BNEZ	R1,LOOP	15	ADD.D	F16,F14,F2	16	S.D	-32(R1),F32	17	ADD.D	F16,F14,F2	18	S.D	-40(R1),F40	19	L.D	F14,-24(R1)	20	ADD.D	F16,F14,F2	21	S.D	-48(R1),F48	22	DADDUI	R1,R1,#-32	23	BNEZ	R1,LOOP	24	ADD.D	F16,F14,F2	25	S.D	-56(R1),F56	26	ADD.D	F16,F14,F2	27	S.D	-64(R1),F64	28	DADDUI	R1,R1,#-32	29	BNEZ	R1,LOOP



Dynamic Scheduling Principle

- What we examined so far is *static scheduling*
 - Compiler reorders instructions so as to avoid hazards and reduce stalls
- *Dynamic scheduling*:
hardware rearranges instruction execution to reduce stalls
- Example:

```
DIV.D F0,F2,F4      ; takes 24 cycles and  
                      ; is not pipelined
```

```
ADD.D F10,F0,F8
```

```
SUB.D F12,F8,F14
```

Dynamic Scheduling Principle

- What we examined so far is *static scheduling*
 - Compiler reorders instructions so as to avoid hazards and reduce stalls
- *Dynamic scheduling:*
hardware rearranges instruction execution to reduce stalls
- Example:

```
DIV.D F0,F2,F4      ; takes 24 cycles and  
RaW; real dependence ; is not pipelined
```

```
ADD.D F10,F0,F8
```

```
SUB.D F12,F8,F14
```

*This instruction cannot continue
even though it does not depend
on previous Div and Add*



Register Renaming: General Idea

- Example, look at F6:

DIV.D F0, F2, F4

ADD.D F6, F0, F8

S.D F6, O(R1)

SUB.D F8, F10, F14

MUL.D F6, F10, F8

F6: RaW

F6: WaR

F6: WaW



Register Renaming: General Idea

- Example, look at F6:

DIV.D F0, F2, F4

ADD.D F6, F0, F8

S.D F6, O(R1)

SUB.D F8, F10, F14

MUL.D F6, F10, F8

F6: RaW

F6: WaW

- False (aka "name") dependences with F6
 - anti: WaR and
 - output: WaW in this example



Register Renaming Technique

- Eliminate name/false dependencies:
 - anti- (WaR) and output (WaW)) dependencies
- Can be implemented
 - by the **compiler**
 - advantage: low cost
 - disadvantage: “old” codes perform poorly
 - by **hardware**
 - advantage: binary compatibility
 - disadvantage: extra hardware needed



Register Renaming by HW

- Same example after renaming:

DIV.D R, F2, F4

ADD.D S, R, F8

S.D S, 0(R1)

SUB.D T, F10, F14

MUL.D U, F10, T

Original code:

DIV.D F0, F2, F4

ADD.D F6, F0, F8

S.D F6, 0(R1)

SUB.D F8, F10, F14

MUL.D F6, F10, F8

- Each destination gets a new (physical) register assigned
- Now only RaW hazards remain, which can be strictly ordered
- We will see several HW implementations of Register Renaming
 1. use ReOrder Buffer (ROB) & Reservation Stations, or
 2. use large register file with mapping table



Speculation (Hardware based)

- Execute instructions along predicted execution paths but *only commit the results if the prediction was correct*



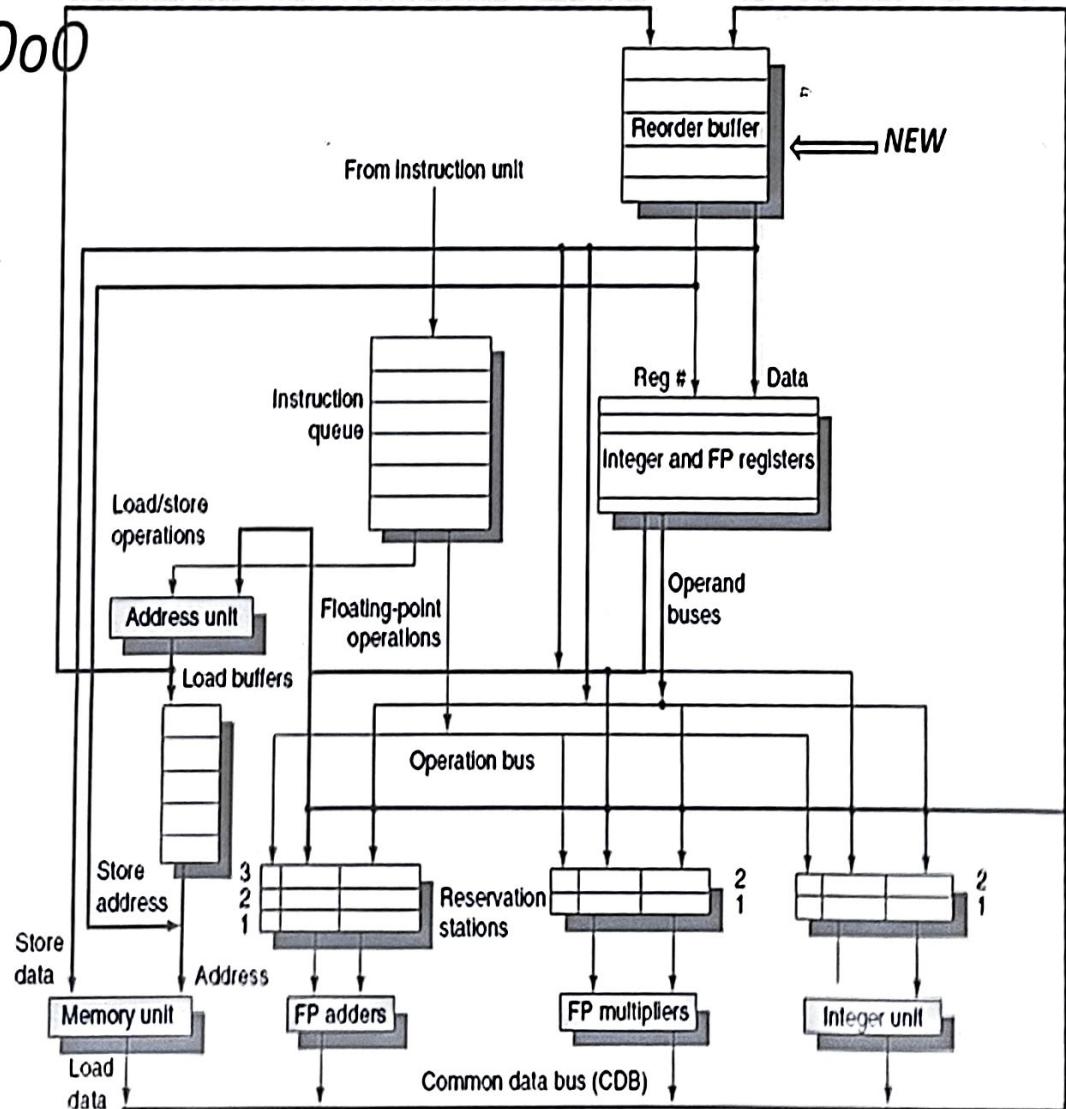
Speculation (Hardware based)

- Execute instructions along predicted execution paths but ***only commit the results if the prediction was correct***
- Instruction commit: *allowing an instruction to only update the register file when instruction is no longer speculative*
- Need an additional piece of hardware to prevent any irrevocable action until an instruction commits:
 - Reorder buffer, or Large renaming register file
 - why? think about it?



Speculative OoO

execution with
speculation
using RoB





Reorder Buffer (RoB)

- Register values and memory values are not written until an instruction commits
- RoB effectively renames the destination registers
 - every destination gets a new entry in the RoB
- On misprediction:
 - Speculated entries in RoB are cleared
- Exceptions:
 - Not recognized/taken until it is ready to commit
 - Precise exceptions require that ‘later’ entries in RoB are cleared