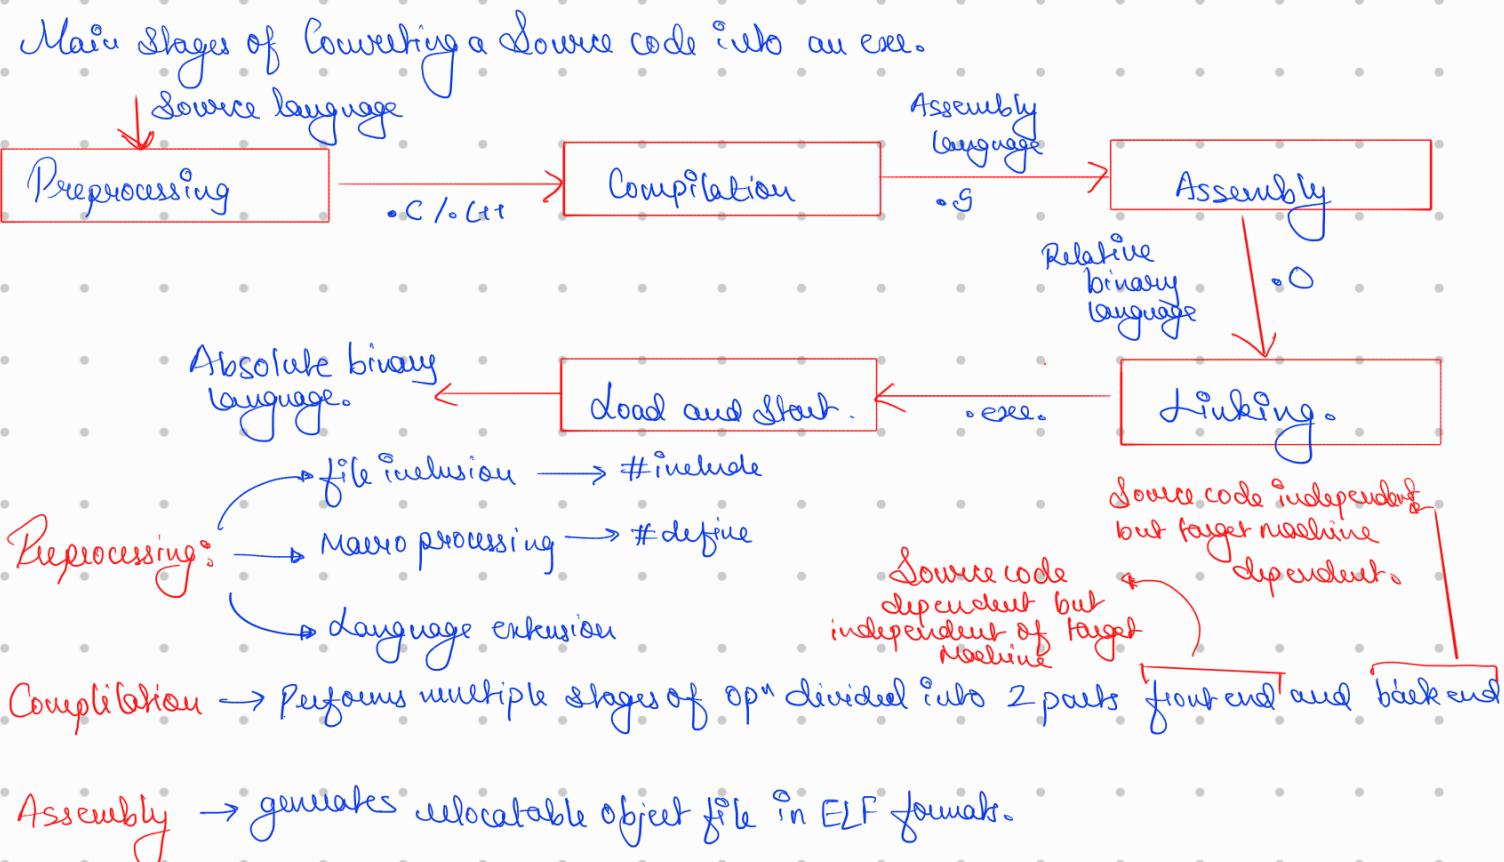


Q1: Roles — Compiler, Assembler, Loader, Linker

- Compiler: Translates high-level source (e.g., C) into target machine form through front-end analysis (lexing, parsing, semantic checks), optimization, and code generation, typically producing assembly or intermediate representation that follows the platform's calling conventions and ABI; it is part of a larger compilation system that also includes preprocessing and later stages (assembly and linking). [1] [2]
- Assembler: Converts assembly language emitted by the compiler into relocatable object files (ELF relocatable objects on Unix-like systems), encoding instructions, data, symbol tables, and relocation entries for later linking. [2] [1]
- Linker: Resolves symbol references across one or more object files and libraries, applies relocation to fix up addresses, and produces an executable (or shared object) with proper sections and segments ready for loading. It supports static and dynamic linking, symbol resolution rules, GOT/PLT for PIC, and library interpositioning. [1] [2]
- Loader: Part of the OS program loader that maps the executable's segments into memory (e.g., ELF program headers), initializes .bss, resolves dynamic dependencies (via the dynamic linker), sets up the initial stack, and transfers control to the program entry point. [2] [1]

Resources: CS:APP Chapter 1 §1.2–1.4 and Chapter 7; Lecture-22-08 (ELF, loader). [1] [2]



Linking: Resolves external symbols and applies relocations to produce an exe.
↳ Static linking:
↳ Dynamic linking:

loader: loads the program into memory / basically allocates actual memory.

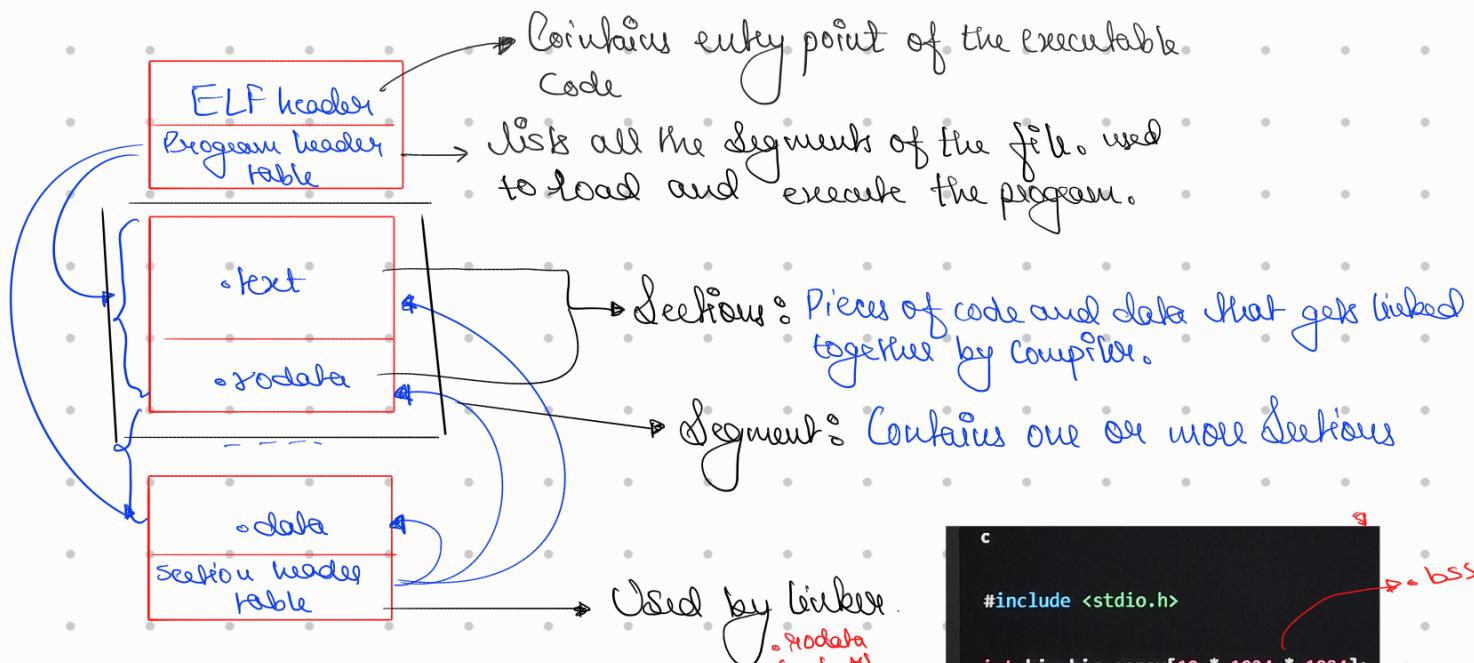
↳ places segments into memory and expands .bss file.

ELF format (Executable and Linkable)

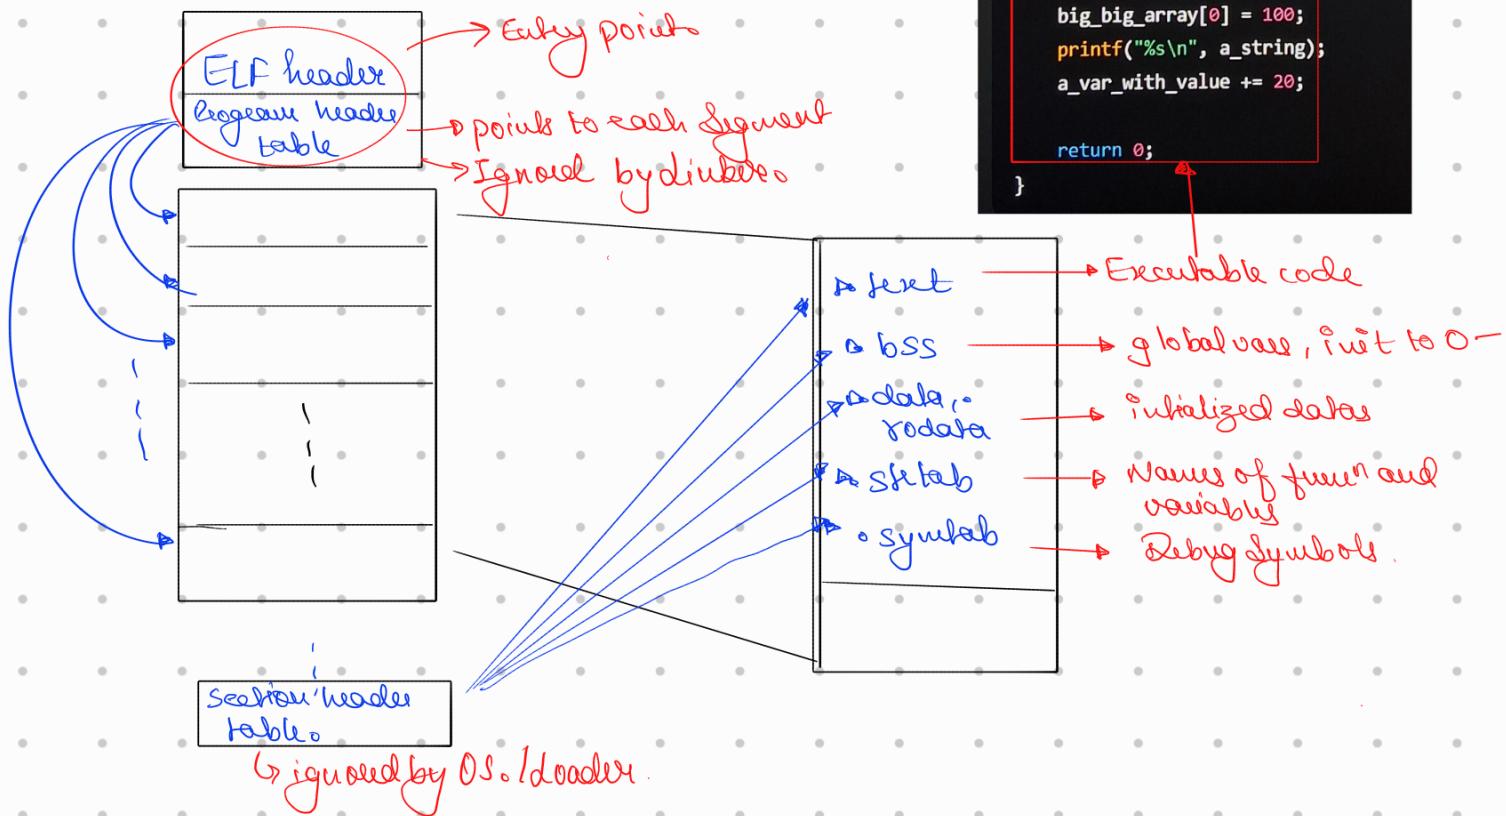
↳ It is basically a standard Linux/unix binary format that stores code and data in a structured way. So that assembler, linker and loader can function properly.

↳ Using the same header/section/segment conventions it can describe

- ① relocatable objects (.o)
- ② shared objects (.so)
- ③ executables (.exe)



we can basically define it as.



```
#include <stdio.h>

int big_big_array[10 * 1024 * 1024];
char *a_string = "Hello, World!";
int a_var_with_value = 0x100;

int main(void) {
    big_big_array[0] = 100;
    printf("%s\n", a_string);
    a_var_with_value += 20;

    return 0;
}
```

Maps contiguous file sections to run-time memory segments

Describes object file sections

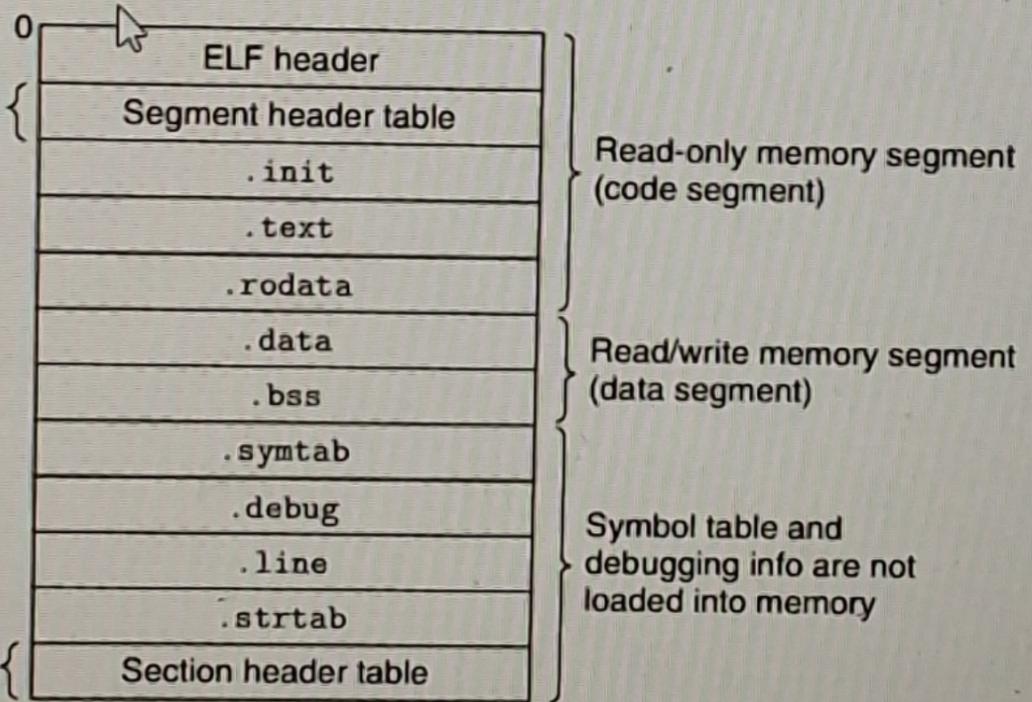


Figure 7.13 Typical ELF executable object file.

Q2: Compile-time vs Runtime Errors (with examples)

- **Compile-time errors:** Detected by the compiler or assembler before execution; include syntax errors, type mismatches, undefined identifiers, or invalid inline assembly syntax.
Example: Missing semicolon or calling a function with wrong parameter types causes compilation to fail.^[1]
 - **Runtime errors:** Occur during execution; include segmentation faults, divide-by-zero traps, illegal instruction, null-pointer dereference, out-of-bounds access, and assertion failures.
Example: Dereferencing a freed pointer triggers a fault handled as an exception/signal on Unix (SIGSEGV).^[1]
- Resources: CS:APP Chapter 1 (translation pipeline), Chapter 8 (exceptions, signals), Chapter 3 §3.10 (buffer overflows).^[1]

Compile time error.

↳ found during parsing, syntax, semantic analysis or code optimization phase.

↳ binary cannot be produced until corrected.

↳ ex: Syntax error, type mismatch etc.

↳ categories

① Syntax error: missing semicolon etc

② Type/semantic error: incompatible types

③ Name/symbol issue at run time: undefined reference to a function

④ Configuration/ABI mismatch: wrong object format detected by linker/loader.

```
c                                     Copy code

#include <stdio.h>

void print_number(int x) {
    printf("Number: %d\n", x);
}

int main(void) {
    char *str = "Hello";
    print_number(str); // ✗ Passing char* instead of int
    return 0;
}
```

- At compile-time, the compiler will emit an error/warning like:

pgsql

Copy code

warning: passing argument 1 of 'print_number' makes integer from pointer without a cast

Runtime error

↳ found only when the program runs with specific inputs or env.

↳ the process starts but then

↳ may fail via traps/exceptions

↳ give incorrect results.

↳ categories (typical)

① Memory faults : Dereferencing Null or bad pointers

② Arithmetic faults : division by zero

③ Resource / OS errors : failed system calls

④ Stack errors : stack overflow.

c

```
#include <stdio.h>
#include <stdlib.h>

int main(void) {
    int *ptr = malloc(0); // Could return NULL depending on system
    if (ptr == NULL) {
        printf("malloc failed, but we'll dereference anyway...\n");
    }
    *ptr = 42; // ❌ Dereferencing NULL → runtime segfault
    free(ptr);
    return 0;
}
```

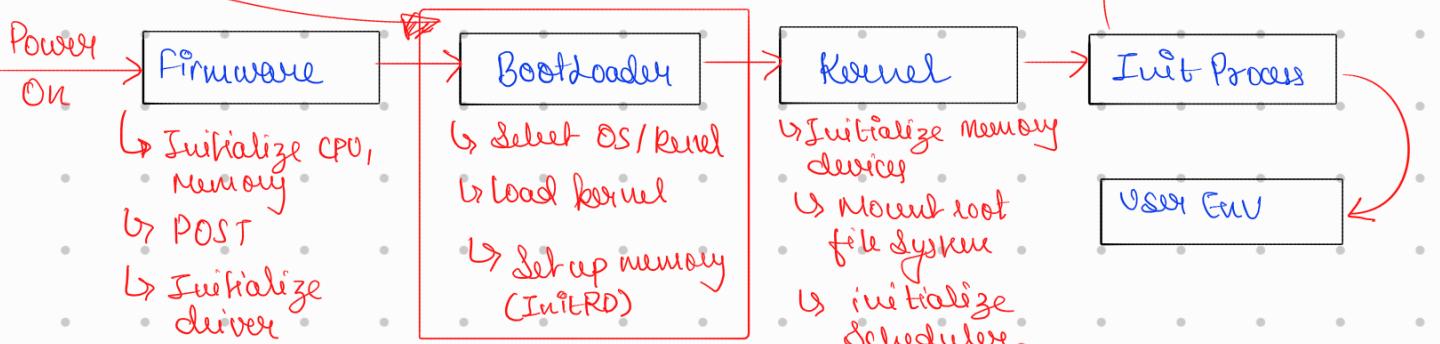
Q3: Role of a Bootloader

- A bootloader initializes minimal hardware, selects a bootable device, and loads the OS image and initramfs from disk into memory. In multi-stage designs (e.g., GRUB), stage-1 fits in MBR constraints (~512 bytes), stage-1.5 adds filesystem awareness, and stage-2 provides menu/config parsing; the bootloader then places the kernel and initrd, sets kernel parameters, and transfers control to the kernel's entry code. [2]
- It bridges firmware (BIOS/UEFI) and OS by setting up CPU mode, memory maps, and passing command line or device tree. [2]

Resources: Lecture-22-08 (Boot Sequence, GRUB stages, kernel placement). [2]

Boot loader.

- Its job is to bridge the firmware and the OS
- It is the first software loaded by the firmware.
- It selects and loads the kernel, Initrd, prepares memory and CPU state and then transfers control to kernel entry point



What bootloader sets up

- Load kernel image to a suitable physical address and places the initial Ram Disk (InitRD) in memory.

Boot Loader Vs program loader.

- Runs before any OS exists
- maps ELF segments
- Expands .bss
- Locates
- Locals Shared Libraries.

what is an initial RAM disc

An initial RAM disk (often called initrd) is a temporary root filesystem that is loaded into memory (RAM) by the bootloader and used by the Linux kernel during the very early stages of booting.

Why it exists

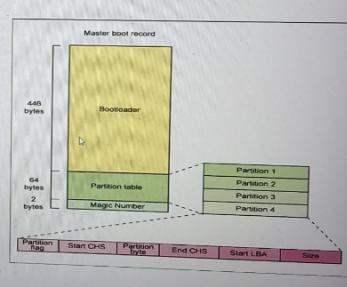
- When the Linux kernel first starts, it doesn't know how to talk to your actual root filesystem (for example, it may be on a SATA drive, NVMe, RAID, encrypted partition, or network).
- The kernel needs some drivers, modules, and setup scripts before it can mount the "real" root filesystem.
- That's where initrd (or initramfs in modern Linux) comes in.

How it works

- Bootloader (like GRUB):**
 - Loads the Linux kernel image.
 - Loads the initrd image into memory.
- Kernel startup:**
 - Kernel unpacks the initrd into a RAM-based filesystem.
 - Kernel mounts initrd as its temporary root filesystem.
- Init scripts inside initrd:**
 - Load necessary kernel modules (e.g., disk drivers, filesystem support, RAID/LVM drivers, encryption).
 - Locate and mount the real root filesystem (e.g., /dev/sda2).
 - Switch root from initrd → real root filesystem.
- System continues booting from the real root.**

Boot Loaders - Linux

- A multi-stage program which eventually loads the kernel image and initial RAM Disk(initrd)
 - Stage-1 Boot Loader is less than 512 bytes (why?)
 - Just does enough to load next stage
- Next stage can reside in boot sector or the partition or area in the disk which is hardcoded in MBR
- Stage-1.5 is a crucial feature
 - Makes grub file-system aware



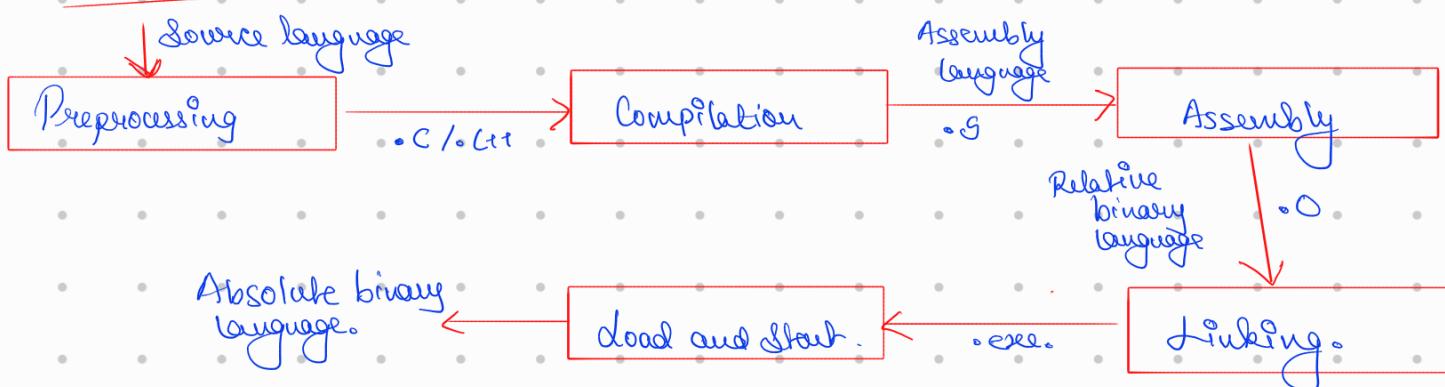
Q4: Interpreted vs Compiled Language (one key difference)

- Key difference: Execution model and timing of translation. Compiled languages (like C) are translated ahead-of-time into machine code and executed natively; interpreted languages (like Python) translate/execute at runtime via a VM/interpreter (and often a bytecode step), enabling portability and dynamic features at the cost of overhead.^[1]
- Consequence: Compiled programs typically have faster steady-state performance; interpreted environments enable rapid development and introspection but rely on a runtime to execute code.^[1]

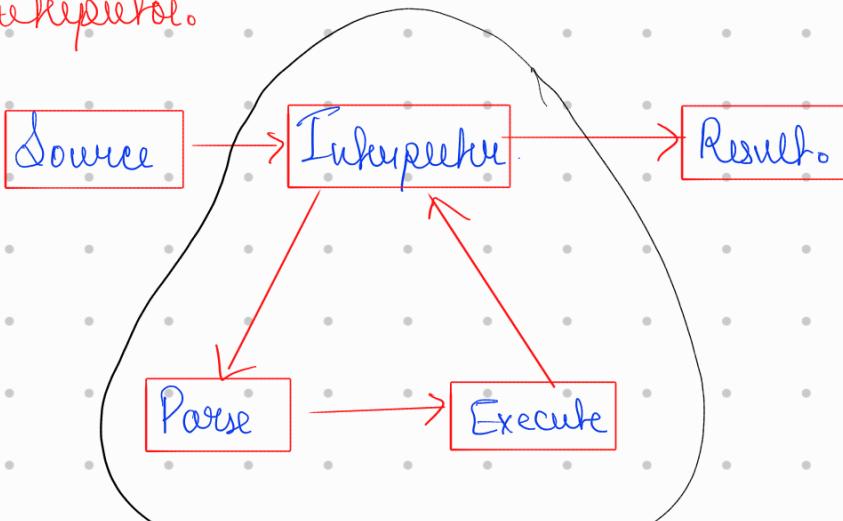
Resources: CS:APP Chapter 1 §1.2 (program translation forms) and §1.4 (execution).^[1]

↳ flow of execution.

Compiler.



Interpreter.



↳ executed for every statement of the code.

Typical tradeoffs.

① Performance

Compiler > interpreter.
↳ less translation overhead
↳ instructions are native

② Portability and flexibility

interpreter > Compiler.
↳ dynamic

Given:

```
int g = 10; // global
int main() {
    int x = 5;
    char *p = malloc(100);
}
```

- g: Stored in the data segment (specifically .data since it has a non-zero initializer). [2] [1]
- x: Automatic local on the stack (lives in the current stack frame of main). [1]
- p: The pointer variable p itself is on the stack; the dynamically allocated 100 bytes are on the heap (managed by the allocator via brk/mmap). [1]

Resources: CS:APP Chapter 9 §9.8–9.9 (memory mapping, malloc), Lecture-22-08 (sections .text/.data/.bss, process memory layout). [2] [1]

Where can variable reside

① Register (register keyword)

↳ nowadays register keyword is not used as CPU optimally uses registers to store largely used data.

② Stack (automatic storage) (auto keyword)

- ↳ local non static variables
- ↳ scope is current block
- ↳ lifetime: current block.

③ Static Storage (static keyword)

- ↳ initialized global / static → .data (data)
 - ↳ zero initialized " " → .bss
 - ↳ constants → .rodata.
- ↳ lifetime: Program.

④ Heap

- ↳ dynamically allocated memory / memory allocated at runtime
- ↳ generally malloc / calloc
- ↳ lifetime is till they are freed.

⑤ Shared libraries global.

- ↳ DSO / data or DSO / bss.

Storage classes, qualifiers, and effects

- **auto:** Default for block-scope variables; automatic storage on the stack; scope is the block; lifetime is the block's execution; rarely written explicitly in modern C. [attached_file:2 p.275]
- **register:** Hint to keep an automatic variable in a CPU register; compilers may ignore; same scope/lifetime as automatic; taking its address is disallowed by the language if truly in a register. [attached_file:2 p.287]
- **static (at file scope):** Gives the variable internal linkage (visible only in that translation unit) and static storage duration (program-long lifetime); placed in .data/.bss/.rodata depending on initialization/constness. [attached_file:4 p.6][attached_file:2 p.46]
- **static (at block scope):** Changes an otherwise local variable to static storage duration; a single instance persists across calls (initialized once); scope remains the block/function; typically in .data/.bss. [attached_file:2 p.46]
- **extern:** Declares but does not define; refers to a definition elsewhere; storage duration is static; linkage is external; placement per defining translation unit (.data/.bss/.rodata). [attached_file:2 p.715]
- **const:** Qualifier that forbids modification through that identifier; does not by itself change storage duration; many compilers place file-scope const objects in read-only segments (.rodata); lifetime follows storage class. [attached_file:4 p.7]
- **volatile:** Prevents optimization that would remove or reorder accesses; required for MMIO, signal-shared flags; does not change storage duration or placement. [attached_file:1 p.84]
- **restrict (C99):** Alias qualifier for pointers, enabling optimizations; no effect on placement, scope, or lifetime. [attached_file:2 p.538]

Q6: Static vs Dynamic Linking (with scenario)

- Static linking: The linker copies library code into the executable at link time, producing a self-contained binary; larger file size, no runtime dependency on shared libs, but no benefit from system library updates and potential duplication across processes.^[1]
- Dynamic linking: The executable contains references resolved at load/run time by the dynamic linker; code is shared in memory across processes, smaller executables, and security/bug fixes in shared libraries apply without rebuilding the app. Uses PLT/GOT and PIC.^[1]
- Prefer dynamic linking: For system tools using libc and other common libraries to reduce memory footprint and receive security updates transparently.^[1]

Resources: CS:APP Chapter 7 §7.10–7.13 (dynamic linking, PIC, interpositioning).^[1]

* Static linking → CSAPP (chapter 7 § 7.2, Page 708)

- ↳ The linker pulls the required code from the static archives and writes it into the executable itself.
- ↳ Symbol resolution and relocation are done at link time.
- ↳ Resulting files are larger but independent from external files at run time.
- ↳ Causes duplication.
- ↳ Slightly faster calls and startup.

* Dynamic linking → CSAPP (chapter 7 § 7.10, Page 784)

- ↳ Executable contains relocation records and references.
- ↳ maps the shared objects (so) at the time of loading the program.
- ↳ Resolves the symbols lazily on first call.
- ↳ has a smaller exe. Supports hot fixes.
- ↳ Requires loader.
- ↳ greater startup overhead.
- ↳ ABI mismatches can break if libraries change incompatibly.

Dynamic linking is preferred when system libraries frequently update or
Code sharing matters

Static linking is preferred in minimal, single binary environment
Where reproducible behaviour without external libraries is required.

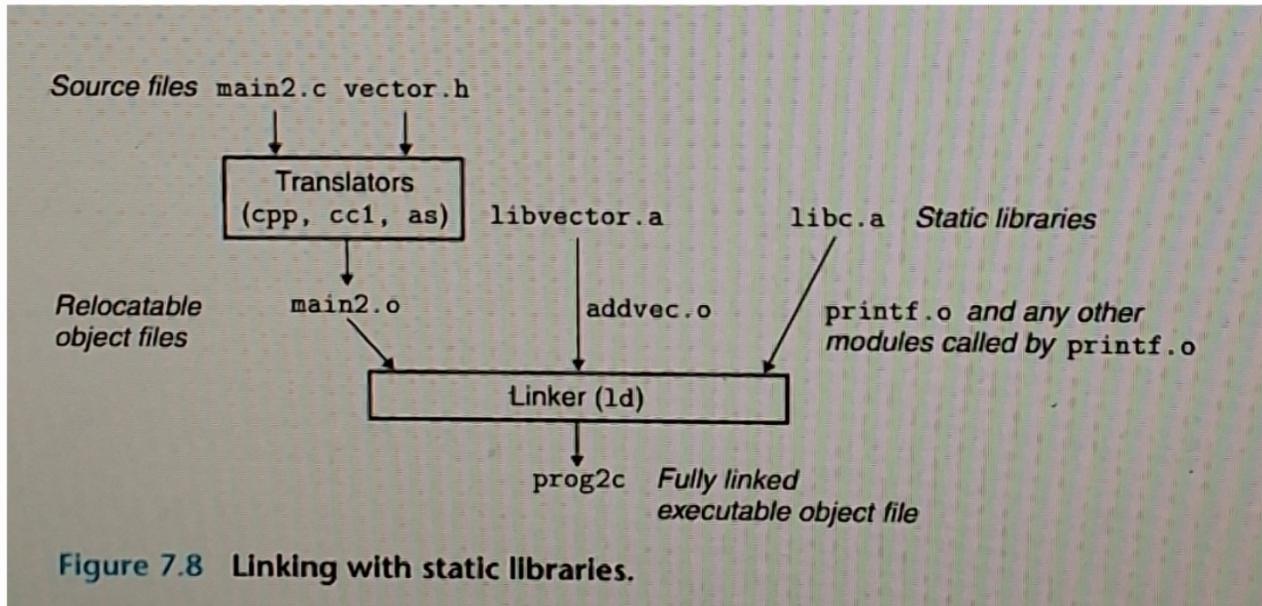
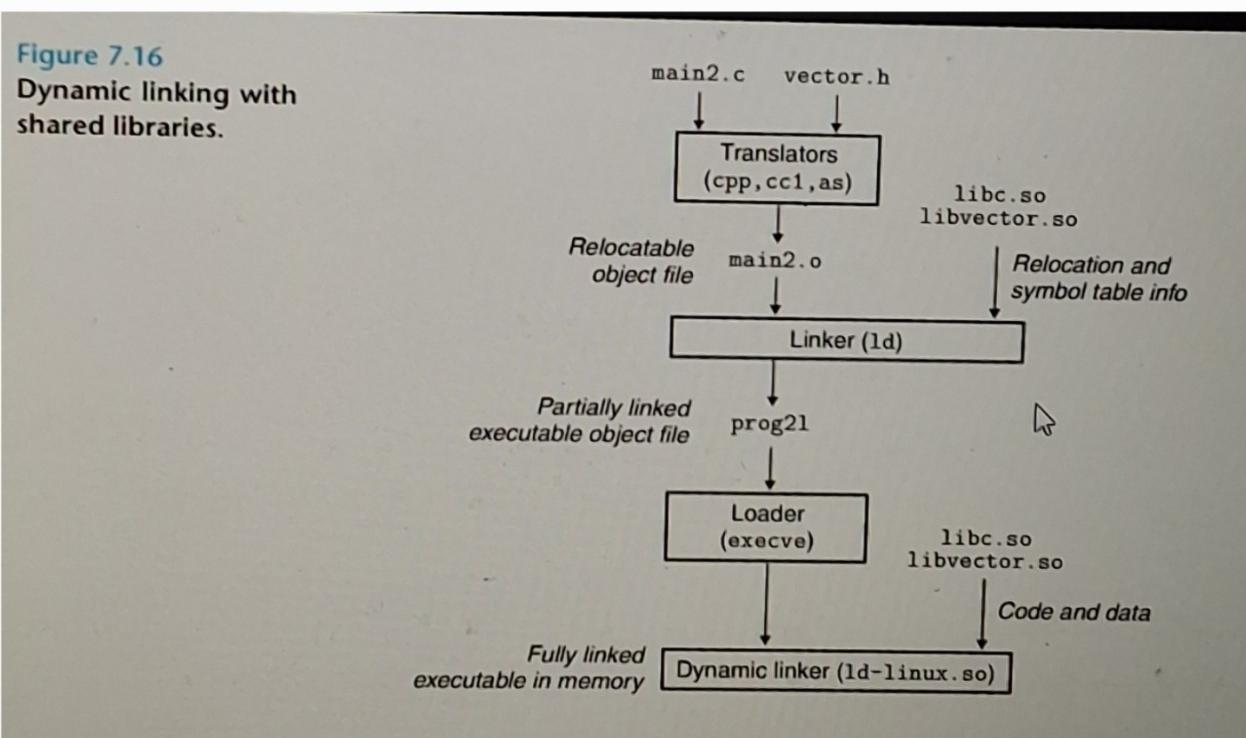


Figure 7.8 Linking with static libraries.



Real life examples of where what type of linking is preferable
Static

- ↳ Embedded Systems.
- ↳ Firmware running on MC
- ↳ Bootloader or kernels.
- ↳ where a certain version of a library is required.

- Dynamic
- ↳ A large GUI framework.
 - ↳ Instead of embedding large code every file share the code.
 - ↳ where multiple program may share a memory.

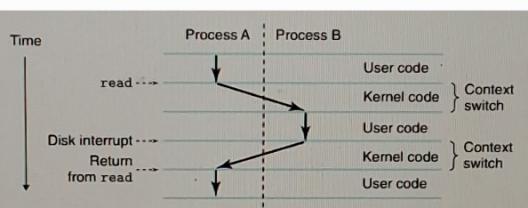
Q7: System Call Interface and Privileged Operations

- System call interface: Provides a controlled gateway from user mode to kernel mode for privileged services (I/O, process control, memory management, networking). User programs issue syscalls (e.g., via syscall instruction) passing arguments; the kernel validates, executes in privileged mode, and returns a result or errno.^[1]
- Why not direct privileged I/O: Hardware protection rings (user vs kernel) prevent untrusted code from manipulating devices or memory mappings directly, ensuring isolation, safety, and fairness; exceptional control flow and mode bits enforce this boundary.^[1]

Resources: CS:APP Chapter 8 (exceptions, processes, user/kernel modes), Chapter 10 (I/O).^[1]

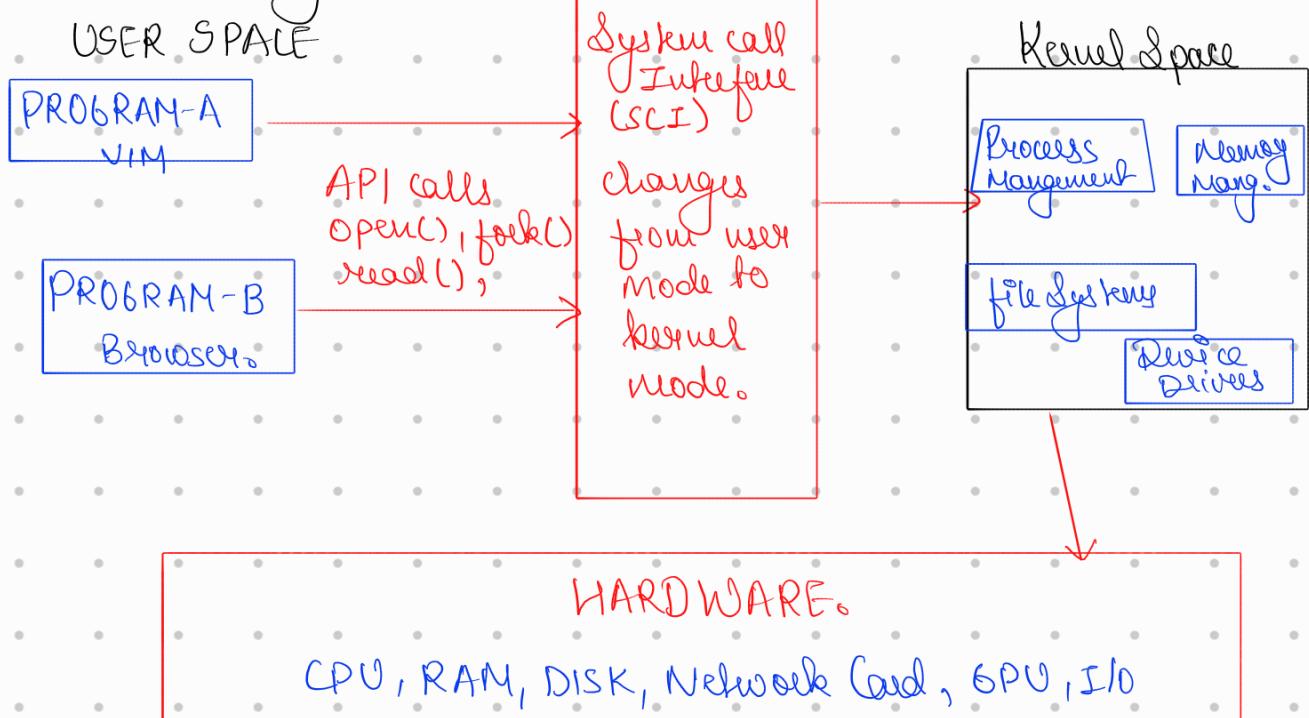
System call Interface (SCI)

- ↳ A system call is the boundary b/w user space (where our code runs) and kernel space (where OS core runs).
- ↳ Applications cannot directly access hardware as that is unsafe, instead they use these services by making a System call to the kernel.



- This is an interface b/w the software and hardware
- ↳ Its responsibilities are
 - ① Resource management
 - ② hardware abstraction (for apps that don't care you are intel or ARM)
 - ③ Security and Isolation

System call visualization



▶ Example: Writing to a File

- 1. User program calls:

```
c

write(fd, buf, count);
```

Copy code

2. The C library (glibc) translates that into a system call number (e.g., `SYS_WRITE`).
3. CPU switches from user mode → kernel mode via a trap instruction (like `int 0x80` or `syscall`).
4. Kernel executes the write handler:
 - Finds the file in its table.
 - Writes data to disk (through driver).
5. Control returns back to user mode with success/failure code.

So in essence:

- Kernel = the OS core that manages resources and hardware.
- System Call Interface = the controlled door that user programs must pass through to ask the kernel for services.

- ↳ Privileged Instructions are only executable in Kernel mode. If a user tries them directly the CPU raises an exception and prevents the execution.
- ↳ This ensures hardware protection against malicious code. How?
- ↳ If a system can directly perform I/O then following this are possible
 - ① One program can write another program's memory
 - ② Disk sectors can be corrupted
 - ③ whole OS can be deleted ultimately killing everything.
 - ④ bypassing of sweet/important files.
- ↳ no application requires hardware details like Sectors, NIC registers and hence kernel provides uniform System Calls (API) and performs such tasks internally

Q8: Typical Process Address Space Layout (diagram description)

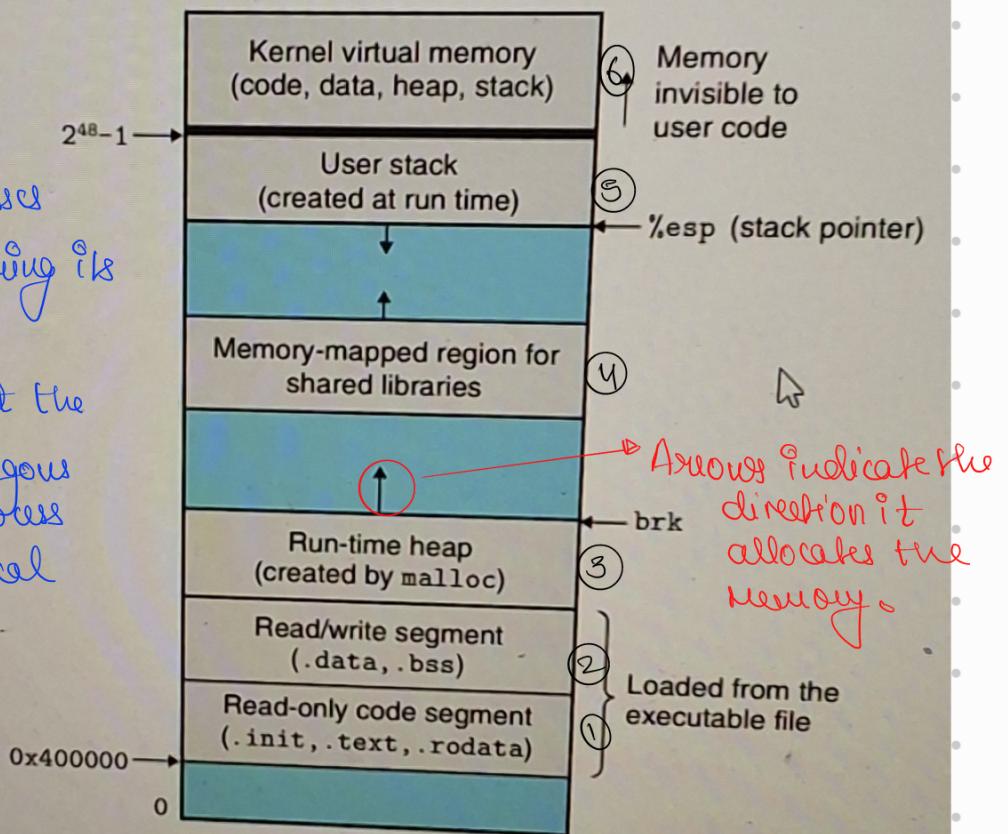
A typical 64-bit Unix process (high to low addresses):

- Kernel space (inaccessible in user mode)
- User space top
- Stack: grows down; holds frames, return addresses, locals
- Memory-mapped region: shared libs, mmap'd files, anonymous mappings
- Heap: grows up from end of .bss/.data; managed by malloc
- BSS: zero-initialized globals/static
- Data: initialized globals/static
- Text (.text): executable code (read/execute) [2] [1]

Resources: CS:APP Chapter 9 §9.8 (memory mapping, address spaces), Lecture-22-08 (ELF layout, loader, stack/heap). [2] [1]

Figure 8.13
Process address space.

- ↳ What is this?
- ↳ Set of memory addresses that process can use during its execution.
- ↳ It gives the illusion that the process has its own contiguous memory while every process shares the same physical memory space.



Components of PAs

① Read only code segment.

- ↳ Contains machine instr^us loaded from executable.
- ↳ no modification allowed during execution.
- ↳ Sub Comp
 - ① .init → Initialization Code.
 - ② .text → actual programme instr^us.
 - ③ .rodata → read only constants. (Constant tables, string literals.)

② Read write segments.

- ↳ Contains global / static variables.
- ↳ These are updateable at runtime
- ↳ Sub Comp
 - ① .data → initialized globals / statics (int x=5)
 - ② .bss → Uninitialized or zero initialized globals (int x,)

③ Runtime heap.

- ↳ Region of Memory for dynamic allocation. (malloc, calloc, new)
- ↳ grows upwards towards higher memory address.
- ↳ brk (program break) defines current end of heap

④ Memory mapped Region for shared libraries.

- ↳ used to load shared files / memory mapped files.
- ↳ grows upward.
- ↳ Enables loading of shared libraries to avoid duplication.

⑤ User Space → Decreasing address.

- ↳ created at runtime when the process starts
- ↳ used for function calls and flow control.
- ↳ stack pointer (esp) is used to track top of stack.
- ↳ Each thread in a process has its own stack.

⑥ Kernel Virtual memory.

- ↳ Contains kernel code, data, heap etc.
- ↳ Invisible to user code.
- ↳ used for system calls, interrupt handling and device driver execution.

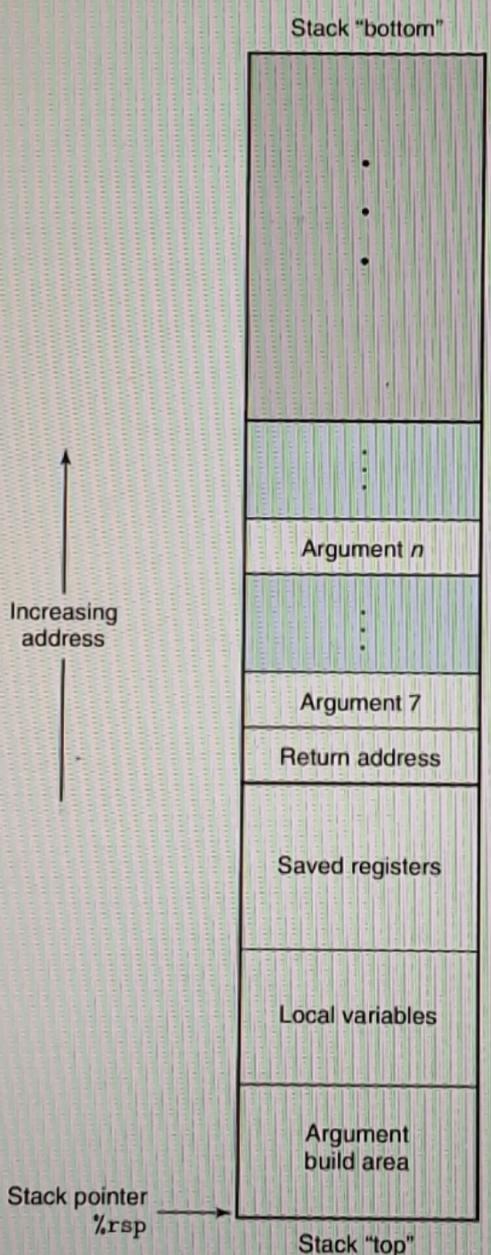
Q9: Recursion Exceeding Stack Size — What Error and Why?

- Error: Stack overflow leads to a protection fault (page fault) when the stack exceeds its guard pages/allocated region, typically delivered as a segmentation fault (SIGSEGV) to the process. The fault arises when the program attempts to access unmapped memory while growing the stack with deeper frames.^[1]
- Root cause: Each recursive call consumes additional frame space; without base case or with large frames, the stack crosses guard boundaries.^[1]

Resources: CS:APP Chapter 3 §3.7 (run-time stack), Chapter 8 (exceptions/signals), Chapter 9 (VM and page faults).^[1]

Figure 3.25

General stack frame structure. The stack can be used for passing arguments, for storing return information, for saving registers, and for local storage. Portions may be omitted when not needed.



↳ Each new recursive call pushes the Stack Pointer downwards

↳ If the stack crosses the the guard boundaries then it will give a Segmentation fault.

Frame for calling function P

Frame for executing function Q

stack frame. Figure 3.25 shows the overall structure of the run-time stack, including its partitioning into stack frames, in its most general form. The frame for the currently executing procedure is always at the top of the stack. When procedure P calls procedure Q, it will push the *return address* onto the stack, indicating where

Q10: Steps to Add a New Linux System Call (high level)

- Kernel changes:
 - Implement the syscall handler function in the appropriate kernel subsystem directory, following kernel coding conventions and argument checking.^[1]
 - Assign a syscall number by updating the syscall table for the target architecture (e.g., arch/x86/entry/syscalls/syscall_64.tbl) and add the prototype to the corresponding header (e.g., include/linux/syscalls.h).^[1]
 - Wire up any required capability checks and documentation (man pages in section 2 usually come later).^[1]
- Build and install:
 - Rebuild the kernel, install modules, update bootloader if necessary, and reboot into the new kernel.^[2]
- User space:
 - Add a libc wrapper (optional until glibc adopts it). In the interim, use syscall(2) with the new number from user programs for testing.^[1]

Resources: CS:APP Chapter 8 (syscall path, user/kernel boundary), Lecture-22-08 (boot/loader context for new kernels).^[2] ^[1]

Steps to add a
New Linux System Call.

1. Write the system call function

- Implement your new system call in the kernel source tree.
- Typically placed in an appropriate file under `kernel/`, `fs/`, or `arch/<arch>/kernel/`.
- Follow the convention:

```
c

asmlinkage long sys_mycall(int arg1, const char __user *arg2) {
    // kernel code here
    return 0;
}
```

[Copy code](#)

2. Add a prototype

- Add the function prototype to the global header:
 - `include/linux/syscalls.h`

3. Assign a syscall number

- System calls are indexed in the `syscall` table.
- Update the `syscall` table file:
 - On x86_64 → `arch/x86/entry/syscalls/syscall_64.tbl`
 - On x86 (32-bit) → `arch/x86/entry/syscalls/syscall_32.tbl`
- Add a line with a new unique `syscall` number, ABI, and name:

```
449 common mycall sys_mycall
```

[Copy code](#)

4. Update user-space interface (optional)

- Define a wrapper function (e.g., in glibc) so user programs can call your syscall via a library API.
- Or call it directly using `syscall(SYS_mycall, ...)`.

5. Recompile the kernel

- Rebuild the kernel and modules:

```
bash  
  
make -j$(nproc)  
make modules_install  
make install
```

 Copy code

- Update the bootloader if necessary and reboot into the new kernel.

6. Test the new system call

- From user space:

```
c  
  
#include <unistd.h>  
#include <sys/syscall.h>  
#include <stdio.h>  
  
#define __NR_mycall 449  
  
int main() {  
    long ret = syscall(__NR_mycall, 42, "Hello");  
    printf("Return value: %ld\n", ret);  
    return 0;  
}
```

 Copy code